

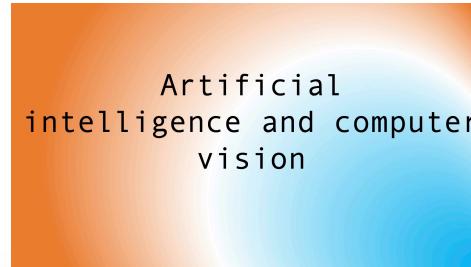
# OPTICAL CHARACTER RECOGNITION (OCR) IN PYTHON



# COURSE CONTENT

- OCR with Tesseract
- Techniques for pre-processing images
- OCR with EAST
- Training an OCR from scratch
  - Artificial Neural Networks and Convolutional Neural Networks
- EasyOCR for natural scenarios
- OCR in videos
- Projects
  - Searching for specific terms
  - Scanner + OCR
  - License plate reading

# TEXT DETECTION – CONTROLLED SCENARIOS



Source: <https://www.pyimagesearch.com/2017/07/17/credit-card-ocr-with-opencv-and-python/>

## TEXT DETECTION – CONTROLLED SCENARIOS



# TEXT DETECTION – NATURAL SCENARIOS



Source: [Mancas-Thillou e Gosselin](#)

- Lighting conditions
- Blur
- Resolution
- Viewing angle
- Non-planar objects
- Objects that are not paper
- Noise in the image
- Unknown layout
- Slanted text
- Different letters

What to do?



Source: <https://www.rsipvision.com/real-time-ocr/>

# TESSERACT

- It is an OCR engine (*Optical Character Recognition*)
- Originally started as a PhD project at Hewlett-Packard (HP) laboratories – scanner
- After gaining popularity, it was developed by HP in 1984. The team worked until 1994
- It was ported to Windows in 1996
- In 2005 it was released to the community as an open source project
- In 2006 its development began to be sponsored by Google
- Since 2006, it is considered to be one of the best and most popular OCR tools
- Official repository: [github.com/tesseract-ocr/tesseract](https://github.com/tesseract-ocr/tesseract)

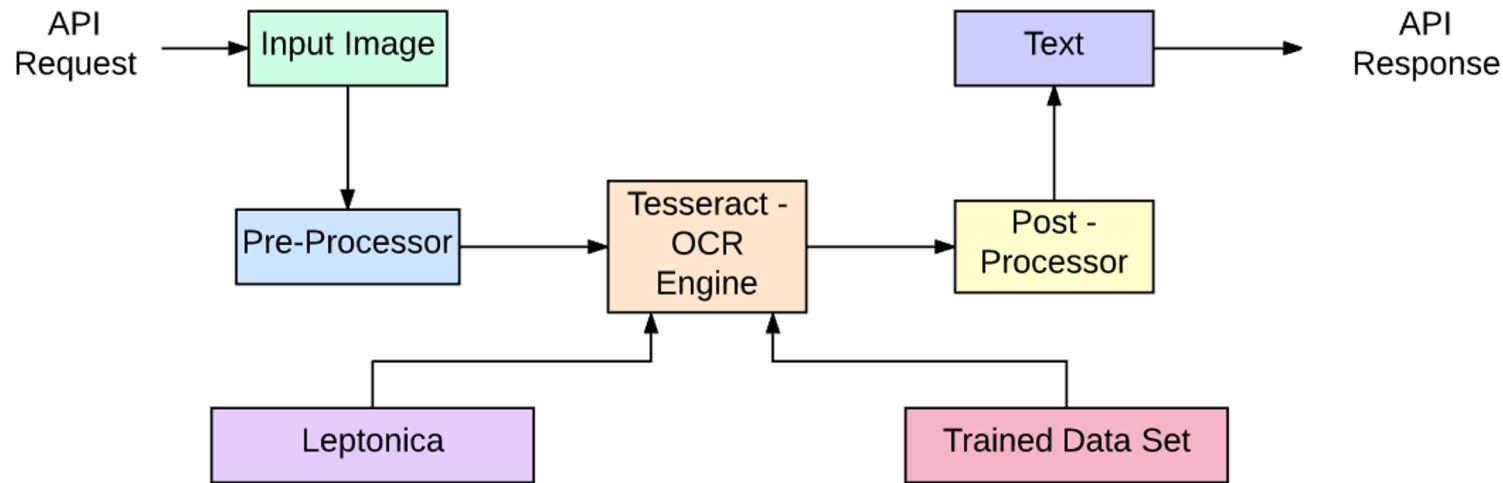


Tesseract OCR

# TESSERACT

- The **first** version of Tesseract only supported English
- The **second** version supported Brazilian Portuguese, French, Italian, German and Dutch
- The **third** version dramatically expanded support to include ideographic (symbolic) languages such as Japanese and Chinese, as well as right-to-left writing languages such as Arabic and Hebrew.
- The **fourth** (June 2021) supports over 100 languages for characters and symbols
- Pytesseract: <https://pypi.org/project/pytesseract/>

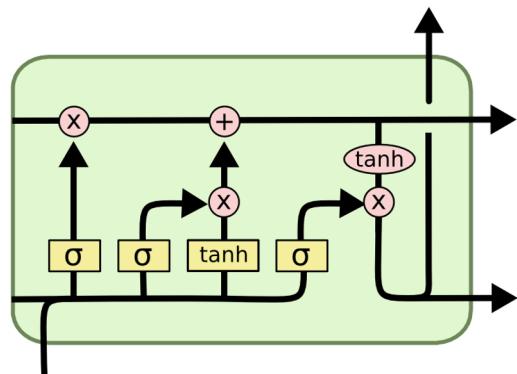
# TESSERACT



Source: [Balaaji Parthasarathy](#)

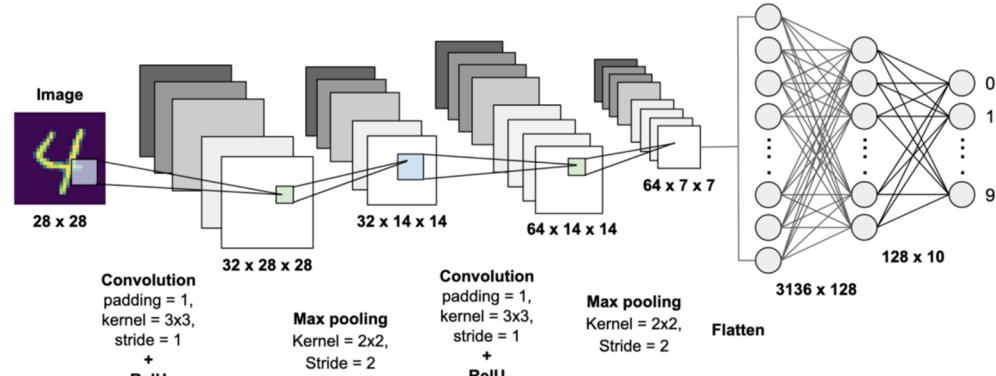
# TESSERACT

To recognize an image containing a single character, we normally use a Convolutional Neural Network (CNN)



LSTM architecture

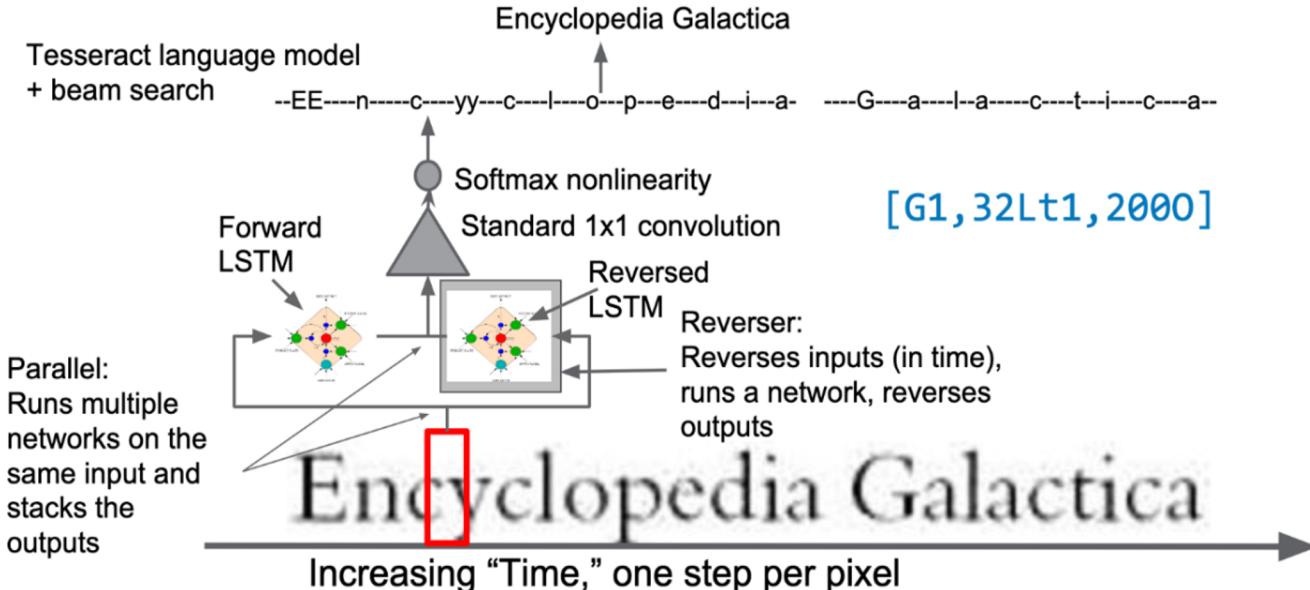
Source: [Colah's blog](#)



Source: Krut Patel

A text of arbitrary length is a sequence of characters and it is more interesting to use RNNs (Recurrent Neural Networks). LSTM (Long short-term memory) is a popular form of RNN.

# TESSERACT



Google

Tesseract Blends Old and New OCR Technology - DAS2016 Tutorial - Santorini - Greece

Source: [Tesseract 3 OCR process paper](#)

# TESSERACT

- Tesseract has a function to detect the orientation of the text in the image, as well as the language it is written in
- This option is called OSD - Orientation and script detection
- Detect if the text in the image is rotated
- If it's rotated, we can apply some kind of preprocessing

# THRESHOLDING

- Thresholding (binarization) is the simplest method of image segmentation
- It consists of separating an image into regions of interest and non-interest by choosing a cut-off point (called threshold)



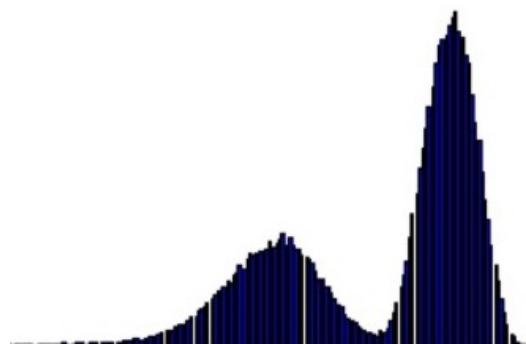
# SIMPLE THRESHOLDING

The value of the new color that the pixel will have is calculated according to the cutoff point (threshold)

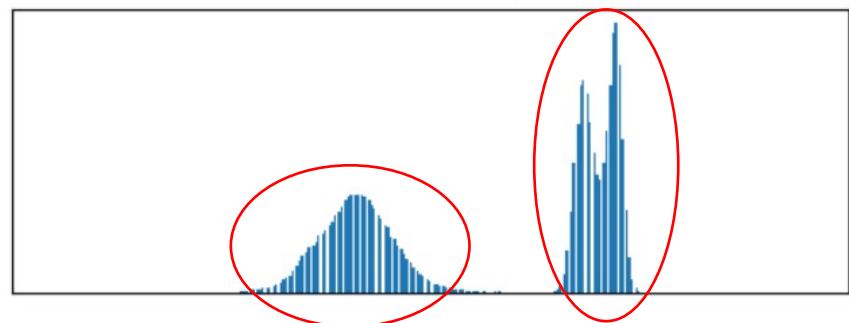
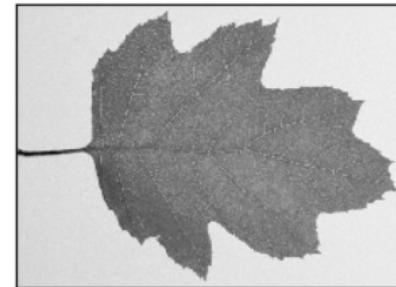
Any pixel with intensity less than or equal to the threshold becomes black. If the pixel has an intensity greater than the threshold, it becomes white.



# OTSU METHOD



Example of a histogram of a bimodal image



# ADAPTIVE THRESHOLDING (GAUSSIAN)

- The threshold is calculated for small regions of the image, different thresholds are obtained for small regions of the image (using the average)
- Gaussian: also uses standard deviation (considering the variation in the pixels) – uses convolution

.....

**Troca de moedas**

A primeira figura mostra 6 moedas de prata, A, C, E, G, I e K e 6 moedas de ouro, B, D, F, H, J e L. A sua tarefa é mover as moedas para a disposição mostrada na segunda figura. Cada movimento deve consistir em uma troca entre uma moeda de prata e uma moeda de ouro adjacente; duas moedas são adjacentes se estiverem unidas por uma linha reta. Sabe-se que o menor número de movimentos que resolve este quebra-cabeça é 17. Você consegue encontrar uma solução em 17 jogadas?

Mova as moedas da primeira posição para a segunda.

.....

**Troca de moedas**

A primeira figura mostra 6 moedas de prata, A, C, E, G, I e K e 6 moedas de ouro, B, D, F, H, J e L. A sua tarefa é mover as moedas para a disposição mostrada na segunda figura. Cada movimento deve consistir em uma troca entre uma moeda de prata e uma moeda de ouro adjacente; duas moedas são adjacentes se estiverem unidas por uma linha reta. Sabe-se que o menor número de movimentos que resolve este quebra-cabeça é 17. Você consegue encontrar uma solução em 17 jogadas?

Mova as moedas da primeira posição para a segunda.

# RESIZING

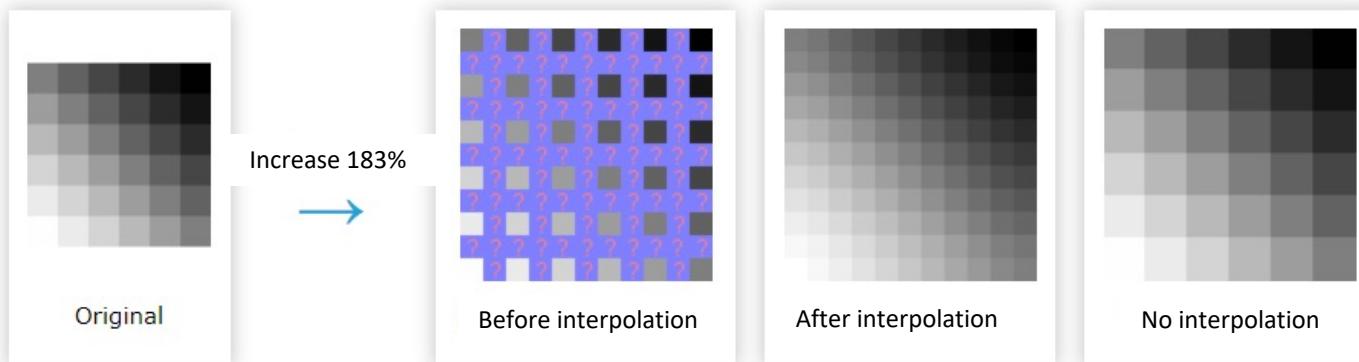
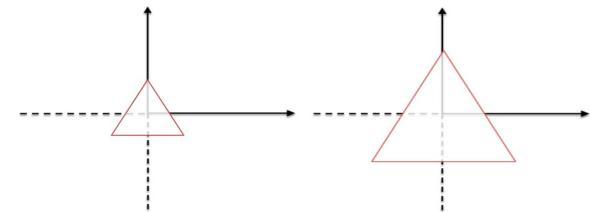
$$\begin{bmatrix} x_T \\ y_T \end{bmatrix} = \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} \cdot \begin{bmatrix} x \\ y \end{bmatrix}$$

Scale factor:

>1 increase

<1 decrease

$s_x = s_y$  uniform scale factor (no distortion)



# RESIZING

OpenCV options

- **INTER\_NEAREST** - a nearest neighbor interpolation. It is widely used because it is the fastest
- **INTER\_LINEAR** - a bilinear interpolation (it's used by default), generally good for zooming in and out of images
- **INTER\_AREA** - uses the pixel area ratio. May be a preferred method for image reduction as it provides good results
- **INTER\_CUBIC** - bicubic (4x4 neighboring pixels). It has better results
- **INTER\_LANCZOS4** - Lanczos interpolation (8x8 neighboring pixels). Among these algorithms, it is the one with the best quality results

# RESIZING

## Comparison between interpolation algorithms – increasing the size

original (50x50)



Region-based segmentation

With this downsizing method of the image we can see that the boundaries of the objects are smoother than in the original. There is no noise in the image, but it is also not very sharp. The image is not very good in terms of contrast.

© 2013-2014, Tanakuchi

area (400x400)



Region-based segmentation

With this downsizing method of the image we can see that the boundaries of the objects are smoother than in the original. There is no noise in the image, but it is also not very sharp. The image is not very good in terms of contrast.

© 2013-2014, Tanakuchi

nearest (400x400)



Region-based segmentation

With this downsizing method of the image we can see that the boundaries of the objects are smoother than in the original. There is no noise in the image, but it is also not very sharp. The image is not very good in terms of contrast.

© 2013-2014, Tanakuchi

linear (400x400)



Region-based segmentation

With this downsizing method of the image we can see that the boundaries of the objects are smoother than in the original. There is no noise in the image, but it is also not very sharp. The image is not very good in terms of contrast.

© 2013-2014, Tanakuchi

cubic (400x400)



Region-based segmentation

With this downsizing method of the image we can see that the boundaries of the objects are smoother than in the original. There is no noise in the image, but it is also not very sharp. The image is not very good in terms of contrast.

© 2013-2014, Tanakuchi

lanczos4 (400x400)



Region-based segmentation

With this downsizing method of the image we can see that the boundaries of the objects are smoother than in the original. There is no noise in the image, but it is also not very sharp. The image is not very good in terms of contrast.

© 2013-2014, Tanakuchi

<http://tanbakuchi.com/posts/comparison-of-opencv-interpolation-algorithms/>

# RESIZING

## Comparison between interpolation algorithms – decreasing the size

original (400x400)



Region-based segmentation

Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:

```
>>> markers = np.zeros_like(coins)
```

area (50x50)

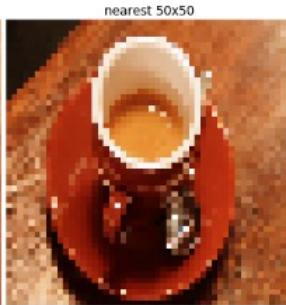


Region-based segmentation

Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:

```
>>> markers = np.zeros_like(coins)
```

nearest (50x50)



Region-based segmentation

In this step, we use the markers of the coins and the background to determine which pixels in the image are to be assigned to the coins and which are to be assigned to the background. This is done by applying a threshold to the markers. All pixels with a value above the threshold are assigned to the coins, and all pixels with a value below the threshold are assigned to the background.

```
>>> markers = np.zeros_like(coins)
```

linear (50x50)



Region-based segmentation

In this step, we use the markers of the coins and the background to determine which pixels in the image are to be assigned to the coins and which are to be assigned to the background. This is done by applying a threshold to the markers. All pixels with a value above the threshold are assigned to the coins, and all pixels with a value below the threshold are assigned to the background.

```
>>> markers = np.zeros_like(coins)
```

cubic (50x50)



Region-based segmentation

In this step, we use the markers of the coins and the background to determine which pixels in the image are to be assigned to the coins and which are to be assigned to the background. This is done by applying a threshold to the markers. All pixels with a value above the threshold are assigned to the coins, and all pixels with a value below the threshold are assigned to the background.

```
>>> markers = np.zeros_like(coins)
```

lanczos4 (50x50)



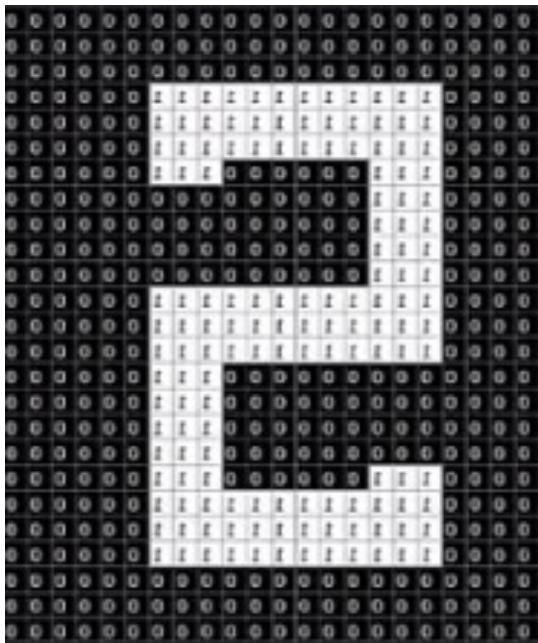
Region-based segmentation

In this step, we use the markers of the coins and the background to determine which pixels in the image are to be assigned to the coins and which are to be assigned to the background. This is done by applying a threshold to the markers. All pixels with a value above the threshold are assigned to the coins, and all pixels with a value below the threshold are assigned to the background.

```
>>> markers = np.zeros_like(coins)
```

<http://tanbakuchi.com/posts/comparison-of-opencv-interpolation-algorithms/>

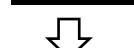
# MORPHOLOGICAL OPERATIONS – EROSION AND DILATION



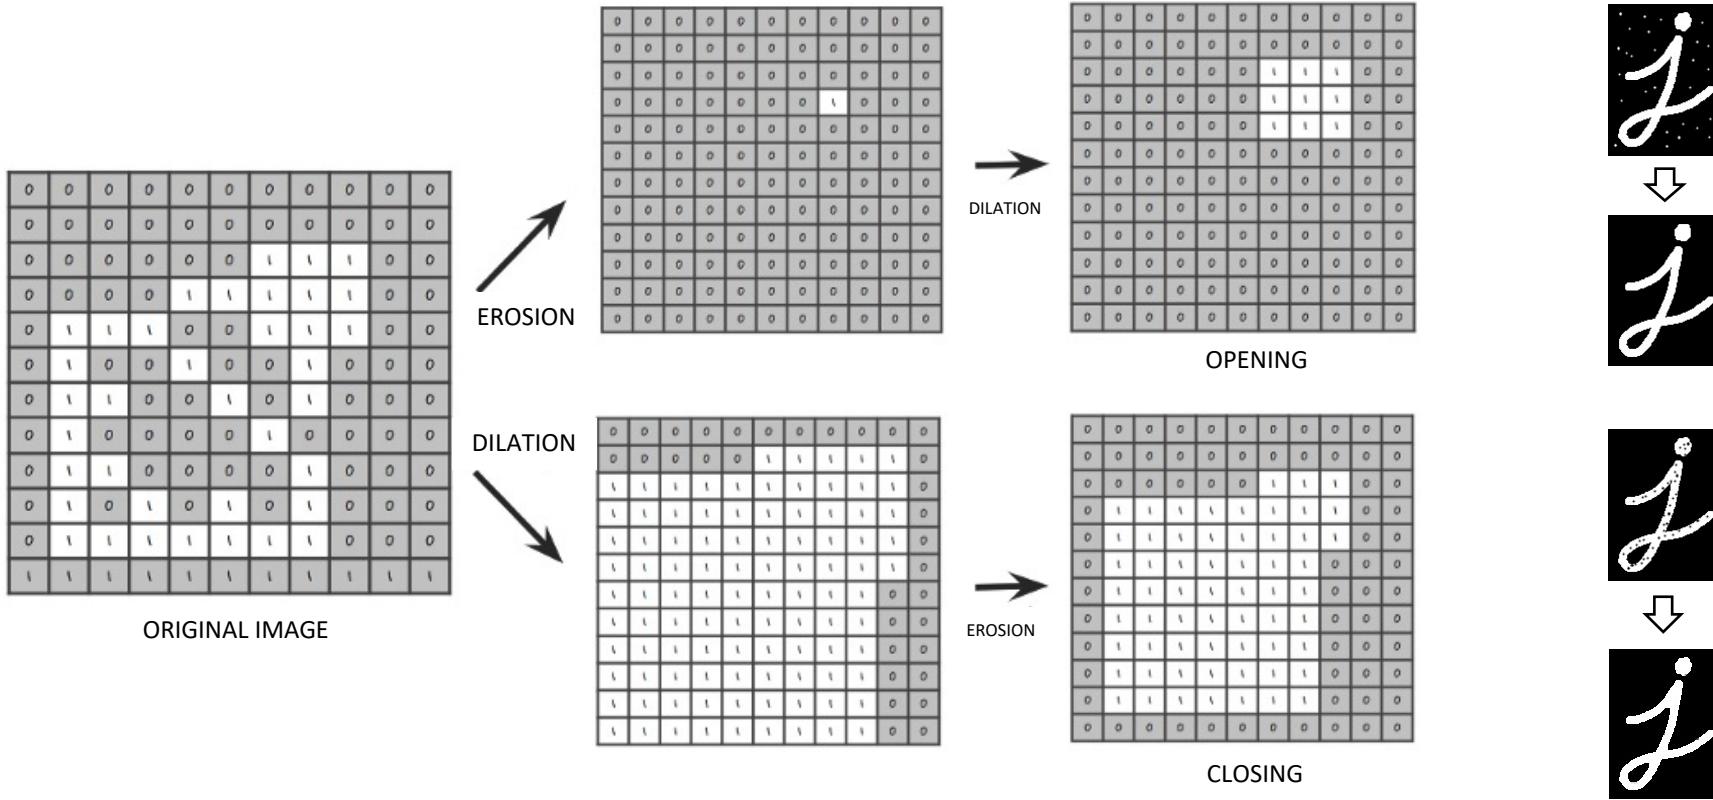
0	0	0	0	0	0	0	0	0	0
0	0	1	0	0	0	0	0	0	0
0	0	1	1	1	0	0	0	0	0
0	1	1	1	1	0	0	0	0	0
0	1	1	1	1	0	0	0	0	0
0	1	1	1	1	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

EROSION      DILATION

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	1	0	0	0	0	0
0	0	1	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0



# MORPHOLOGICAL OPERATIONS – OPENING AND CLOSING



# NOISE REMOVAL WITH BLUR

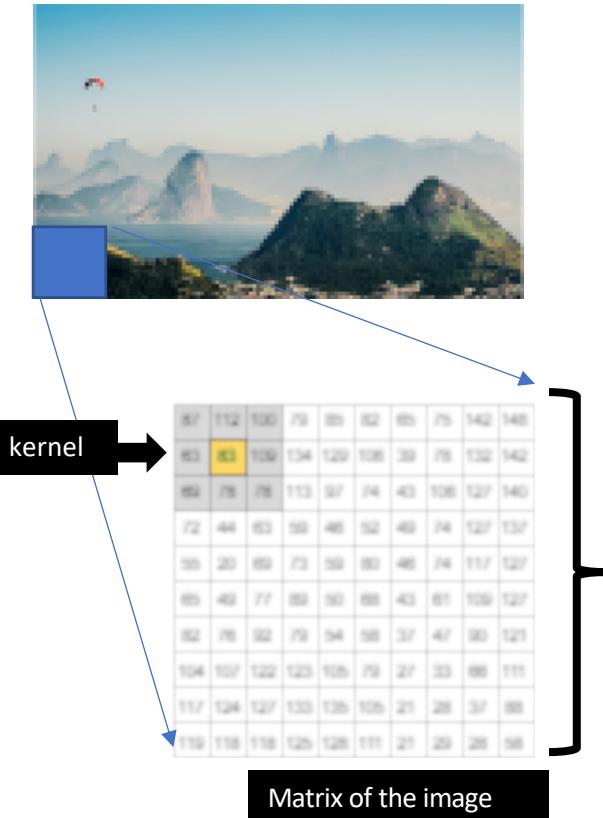
Spatial filters are implemented through kernels (masks/arrays), with odd dimensions. Filters can be::

- Low-pass – used for blur
- High-pass – used to sharpen



# CONVOLUTION

- Convolution is a mathematical operation performed on two matrices, which produces a third matrix that is the result of the operation
- The primary matrix is the image to be treated, and the treatment of this matrix of the original image is done by another matrix called “kernel”, or mask
- Depending on the kernel values it is possible to obtain filters of different types, such as blur and sharpen



# BLUR WITH AVERAGE

161	137	244	254	255	254
154	75	200	249	255	255
109	96	143	223	255	255
69	107	196	236	255	255
51	45	134	218	251	255
58	56	75	224	255	255

Original image

x

$$1 + 1 + 1 + 1 + 1 + 1 + 1 + 1 + 1 = 9$$

Kernel = average

161	137	244
154	75	200
109	96	143

=

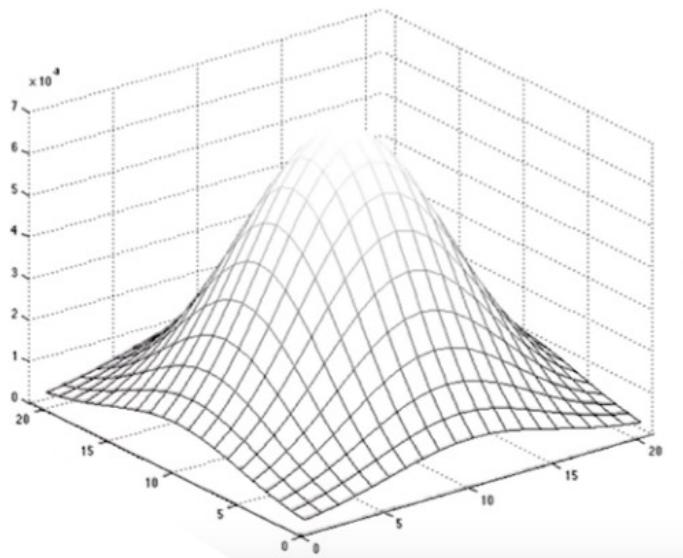
1	1	1
1	1	1
1	1	1

Result  
 $1319 / 9$   
= 147

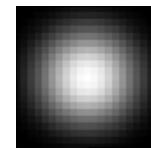
$$\begin{aligned}161 + 137 + 244 \\154 + 75 + 200 \\109 + 96 + 143 \\= 1319\end{aligned}$$

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

# GAUSSIAN BLUR

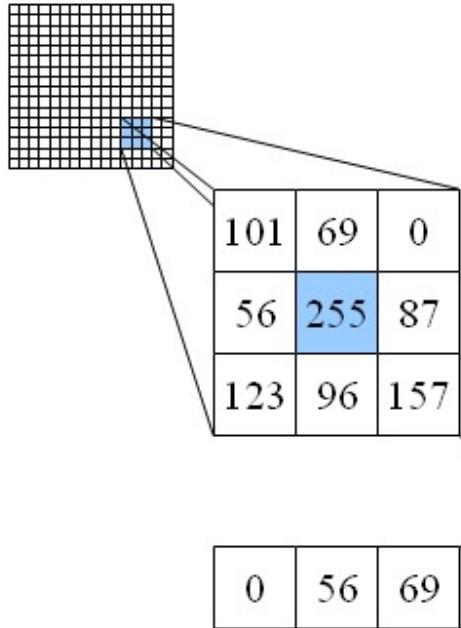


$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$



Graphic representation of  
a 21x21 Gaussian filter

# BLUR WITH MEDIAN



2, 2, 3, **7**, 8, 9, 9

Median = **7**

1, 4, 4, **5**, **6**, 7, 7, 7

Median =  $(5+6) \div 2$   
= **5.5**

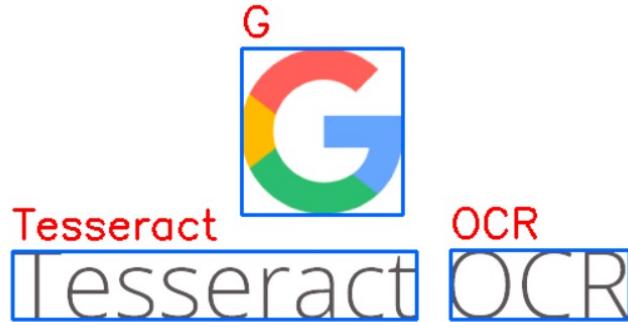
# BILATERAL FILTER



# TEXT DETECTION

Text **detection** before **recognition** is an essential step

Example of images where there would be no need for prior detection



```
'conf': ['-1', '-1', '-1', '-1', 90, '-1', 74, 66],  
'height': [400, 236, 236, 92, 92, 87, 76, 87],  
'left': [0, 38, 38, 38, 38, 102, 102, 307],  
'text': ['', '', '', '', 'TESTANDO', '', '0', 'OCR...'],  
'top': [0, 79, 79, 79, 79, 228, 233, 228],  
'width': [700, 607, 607, 607, 607, 532, 77, 327],
```



# EAST – TEXT DETECTOR

**EAST** (*Efficient Accurate Scene Text detector*) is a deep learning model, officially published in 2017 by Zhou et al.

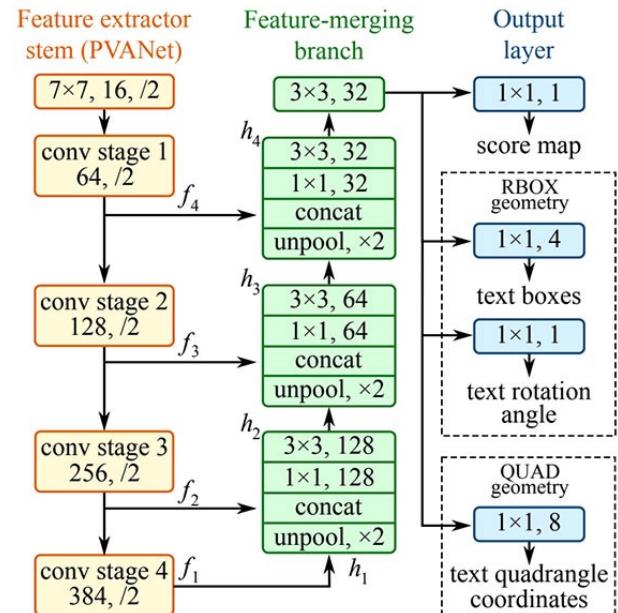
- Uses convolutional layers to extract features from images and thus detect the existence of texts
- It only identifies the **location of the text**. It is not an OCR that will convert it to characters
- It is one of the most accurate text detection techniques
- It is also one of the best known and gained popularity after the release of version 3.4.2 of OpenCV, which made it possible to implement it more easily through the DNN module

# EAST – TEXT DETECTOR

Features at different processing levels are treated in a matrix, leading to the identification of bounding boxes where the text appears in the image

- It has only two stages: FCN (Fully Convolutional Network) and NMS (non-maximum suppression) to suppress weak and overlapping bounding boxes

## EAST architecture (Fully-Convolutional Network)



Source: [Zhou et al.](#)

# ADITIONAL READING

EAST Paper

<https://arxiv.org/abs/1704.03155v2>

Mancas-Thillou e Gosselin

[https://www.tcts.fpms.ac.be/publications/regpapers/2007/VS\\_cmtbg2007.pdf](https://www.tcts.fpms.ac.be/publications/regpapers/2007/VS_cmtbg2007.pdf)

Tesseract OCR

<https://github.com/tesseract-ocr/tesseract>

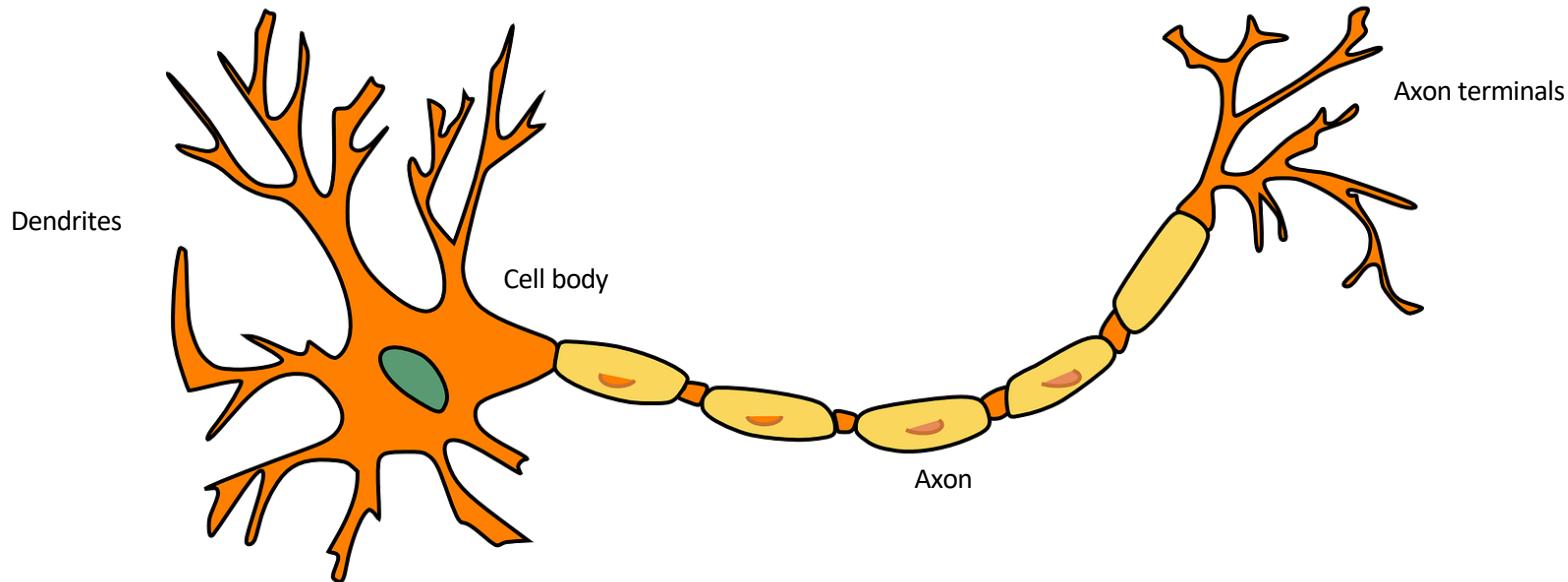
EAST geometry

<https://stackoverflow.com/questions/55583306/decoding-geometry-output-of-east-text-detection>

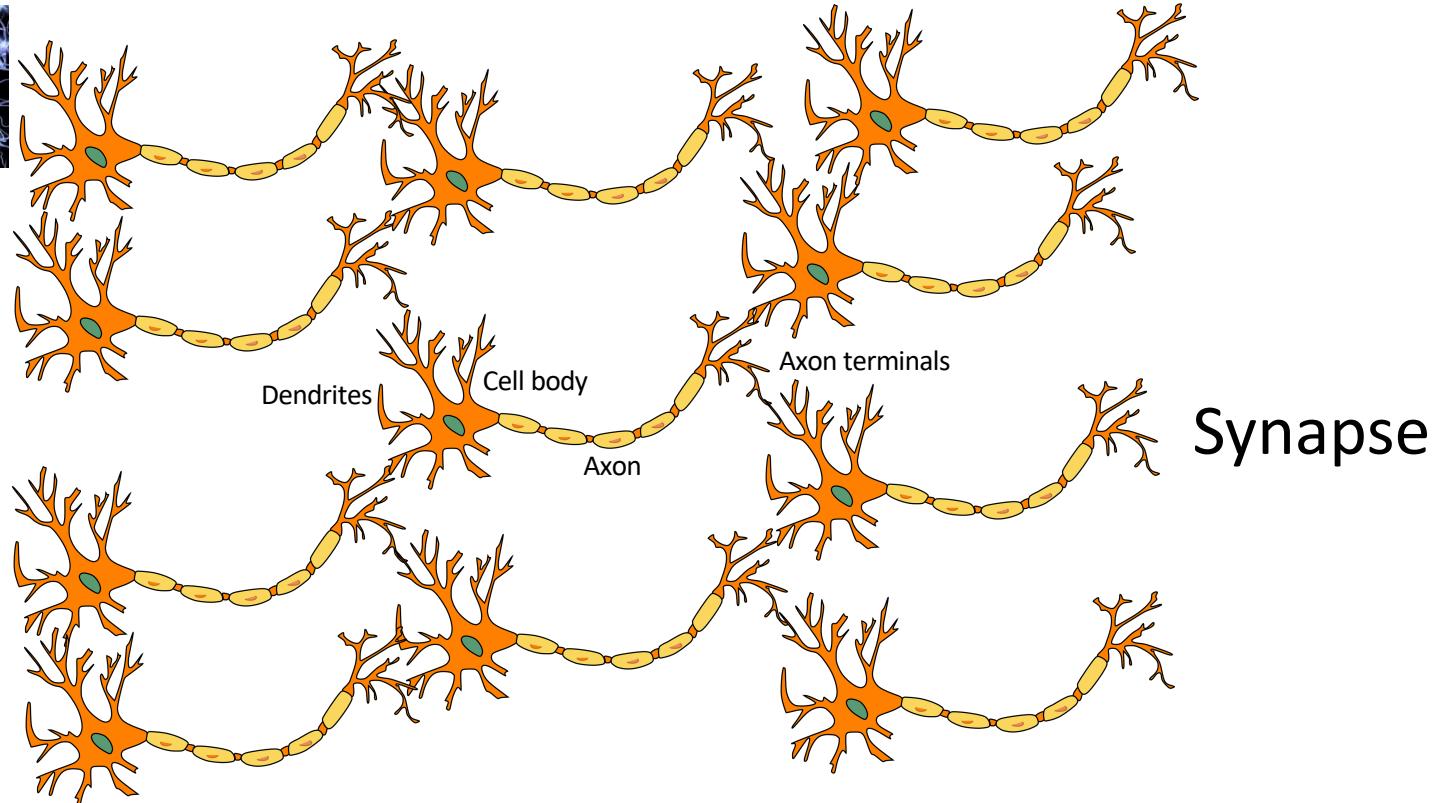
# ARTIFICIAL NEURAL NETWORKS



# ARTIFICIAL NEURAL NETWORKS



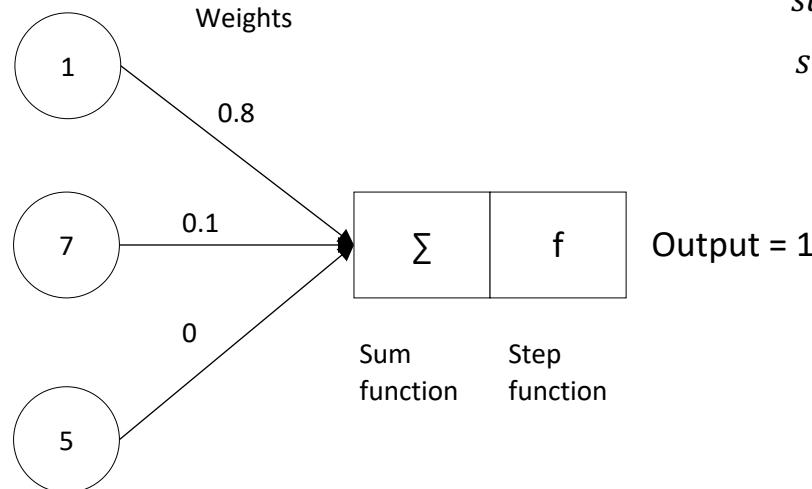
# ARTIFICIAL NEURAL NETWORKS



# ARTIFICIAL NEURON

$$sum = \sum_{i=1}^n x_i * w_i$$

Inputs



$$sum = (1 * 0.8) + (7 * 0.1) + (5 * 0)$$

$$sum = 0.8 + 0.7 + 0$$

$$sum = 1.5$$

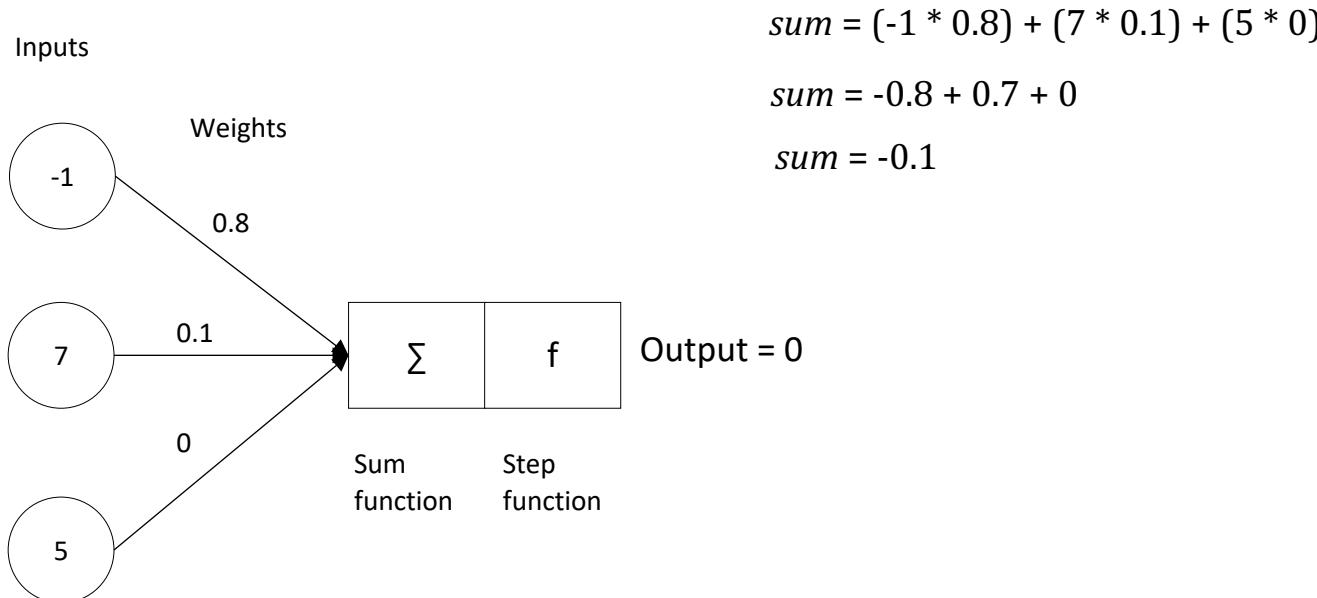
Step function

Greater or equal to 1 = 1

Otherwise = 0

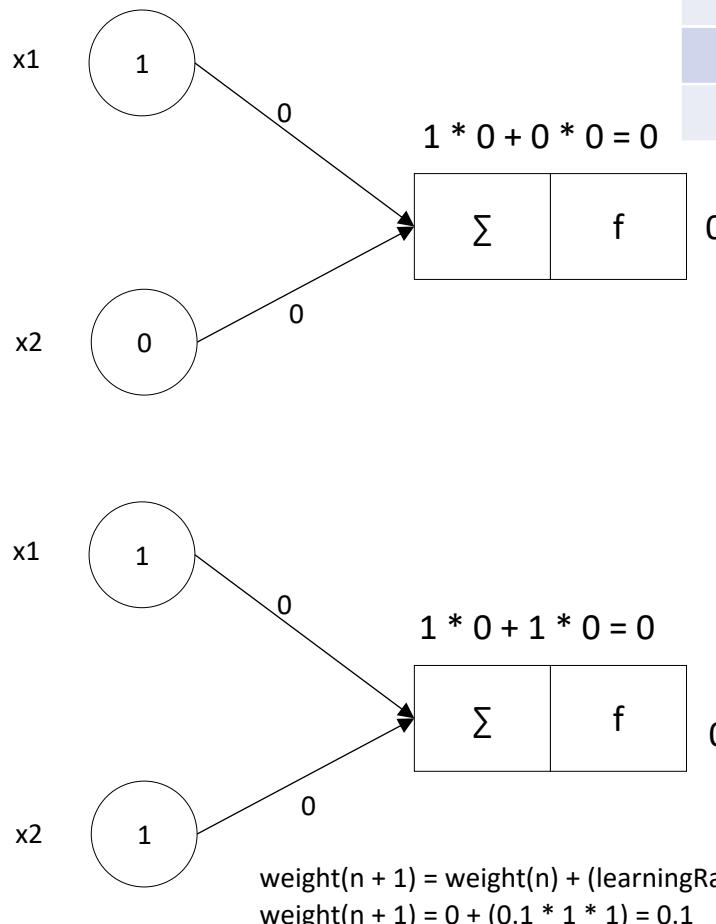
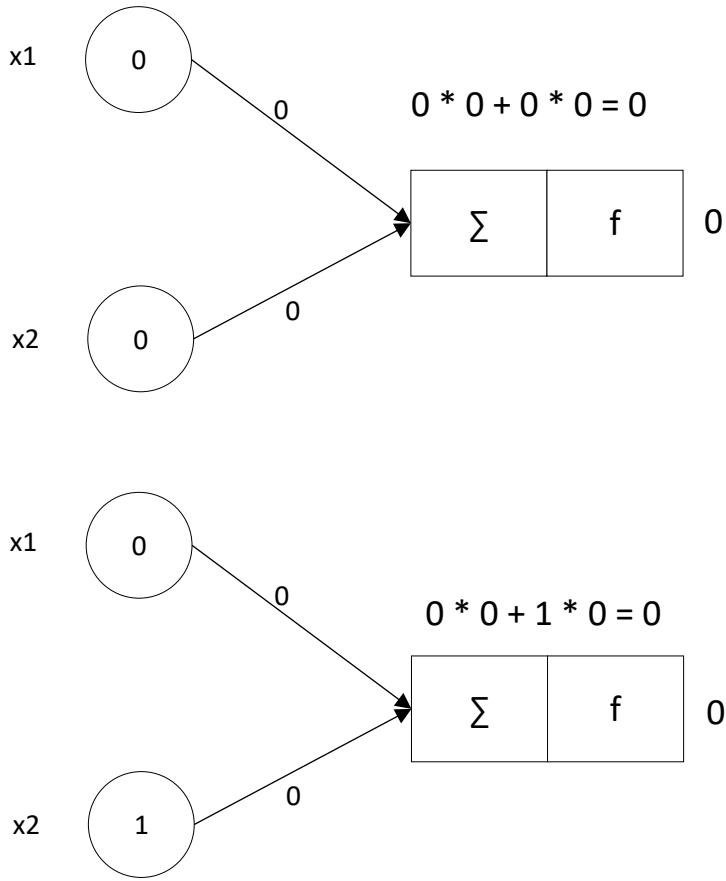
# ARTIFICIAL NEURON

$$sum = \sum_{i=1}^n x_i * w_i$$



# “AND” OPERATOR

X1	X2	Class
0	0	0
0	1	0
1	0	0
1	1	1



x1	x2	Class
0	0	0
0	1	0
1	0	0
1	1	1

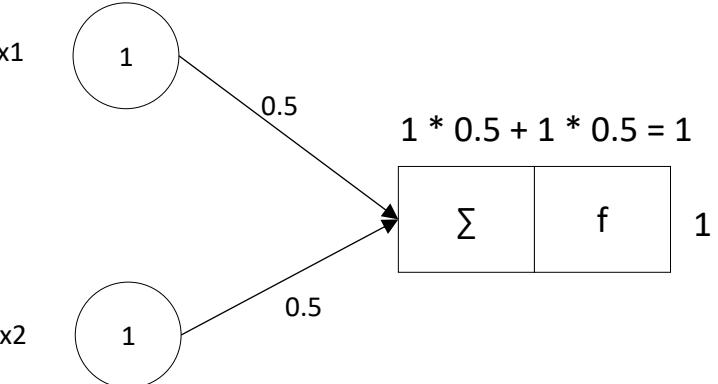
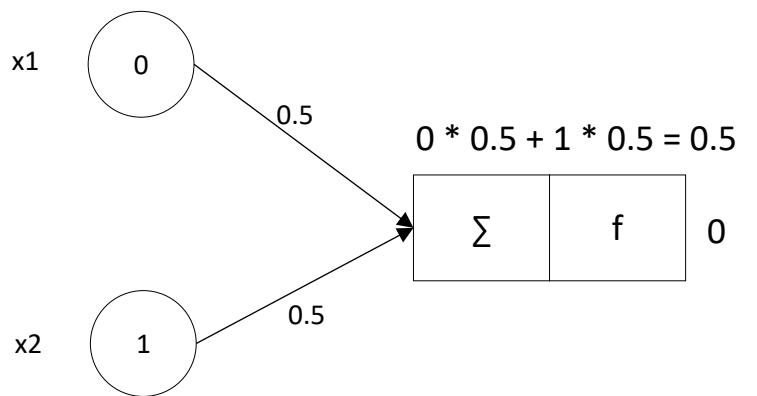
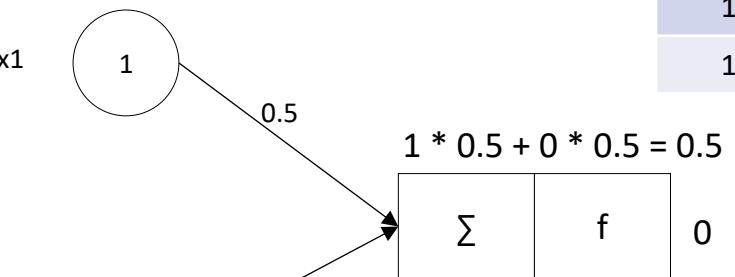
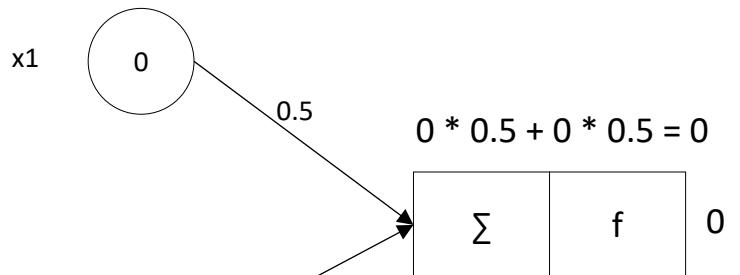
$$0 - 0 = 0$$

$$0 - 0 = 0$$

$$0 - 0 = 0$$

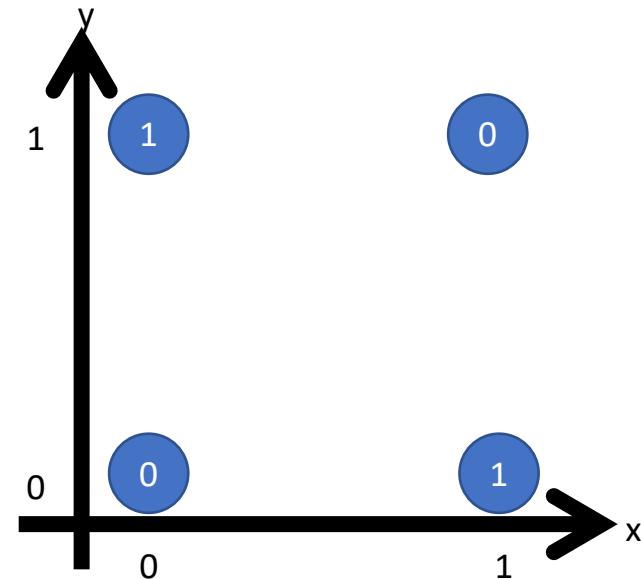
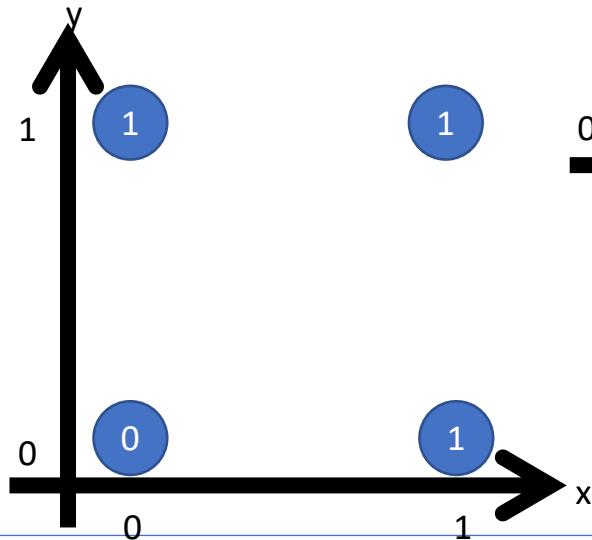
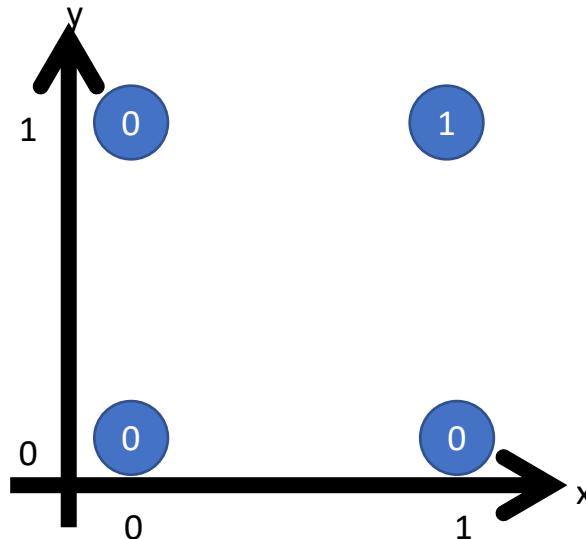
$$1 - 0 = 1$$

x1	x2	Class
0	0	0
0	1	0
1	0	0
1	1	1

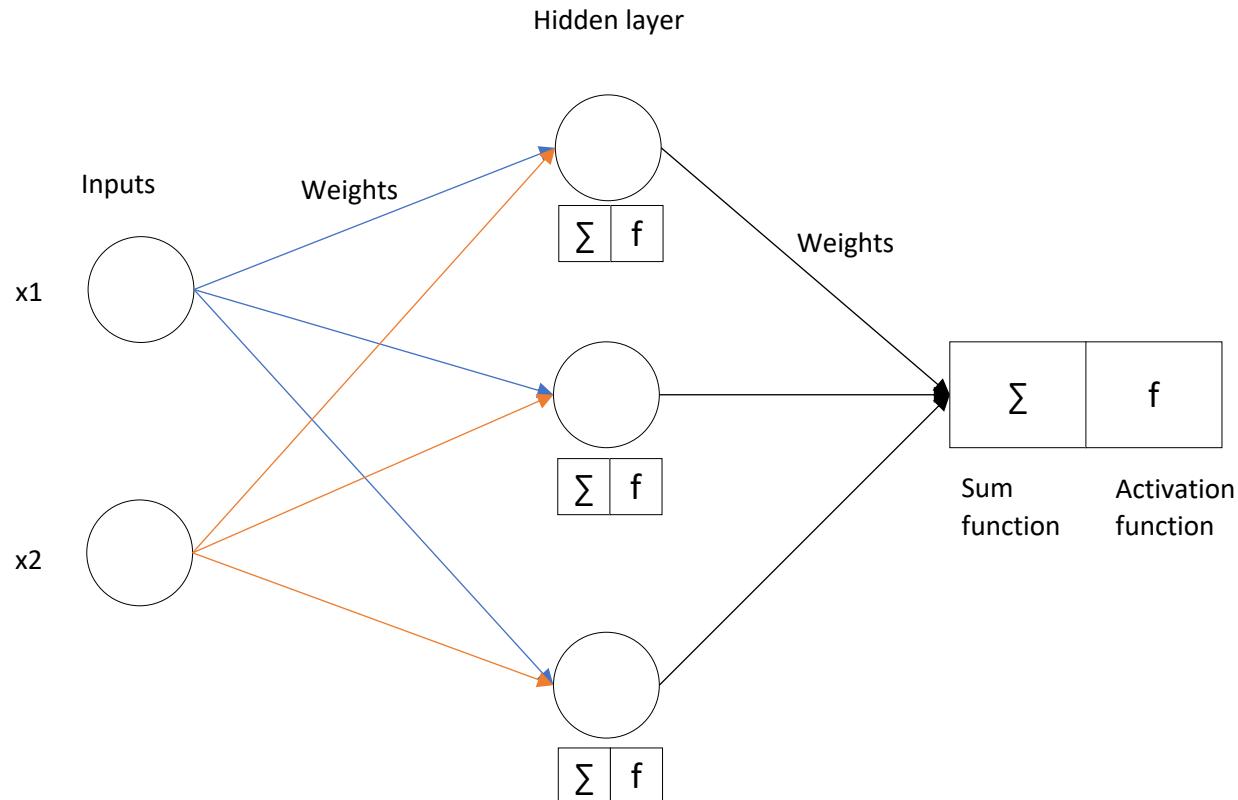


$$\begin{aligned} 0 - 0 &= 0 \\ 0 - 0 &= 0 \\ 0 - 0 &= 0 \\ 1 - 1 &= 0 \end{aligned}$$

# SINGLE LAYER PERCEPTRON

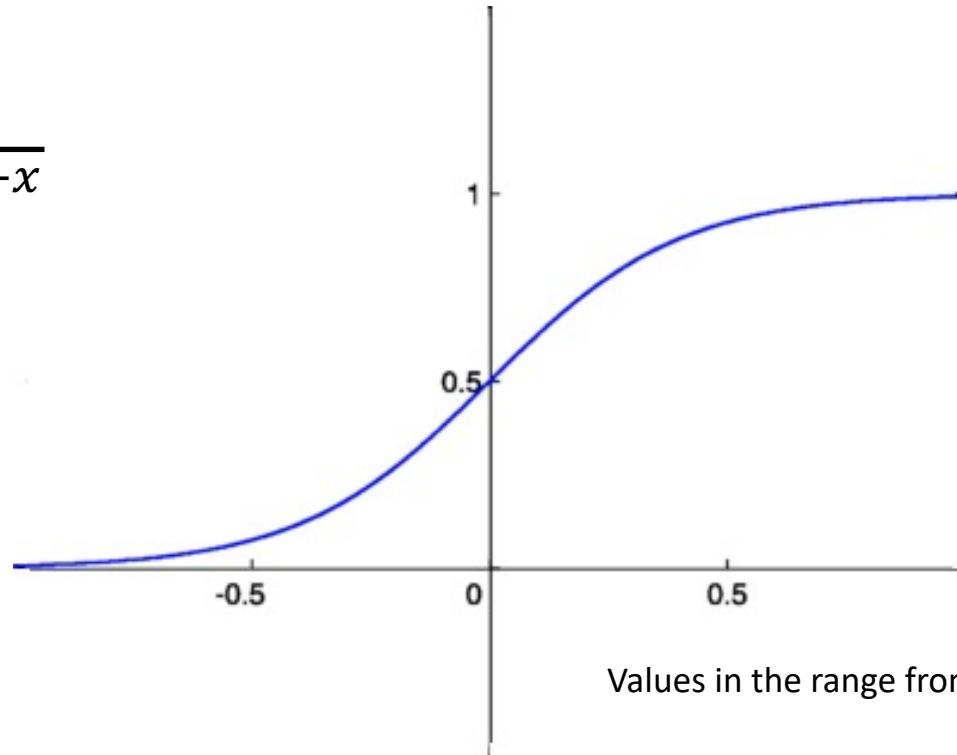


# MULTILAYER PERCEPTRON



# SIGMOID FUNCTION

$$y = \frac{1}{1 + e^{-x}}$$



Values in the range from 0 to 1

If X is high the result will be approximately 1

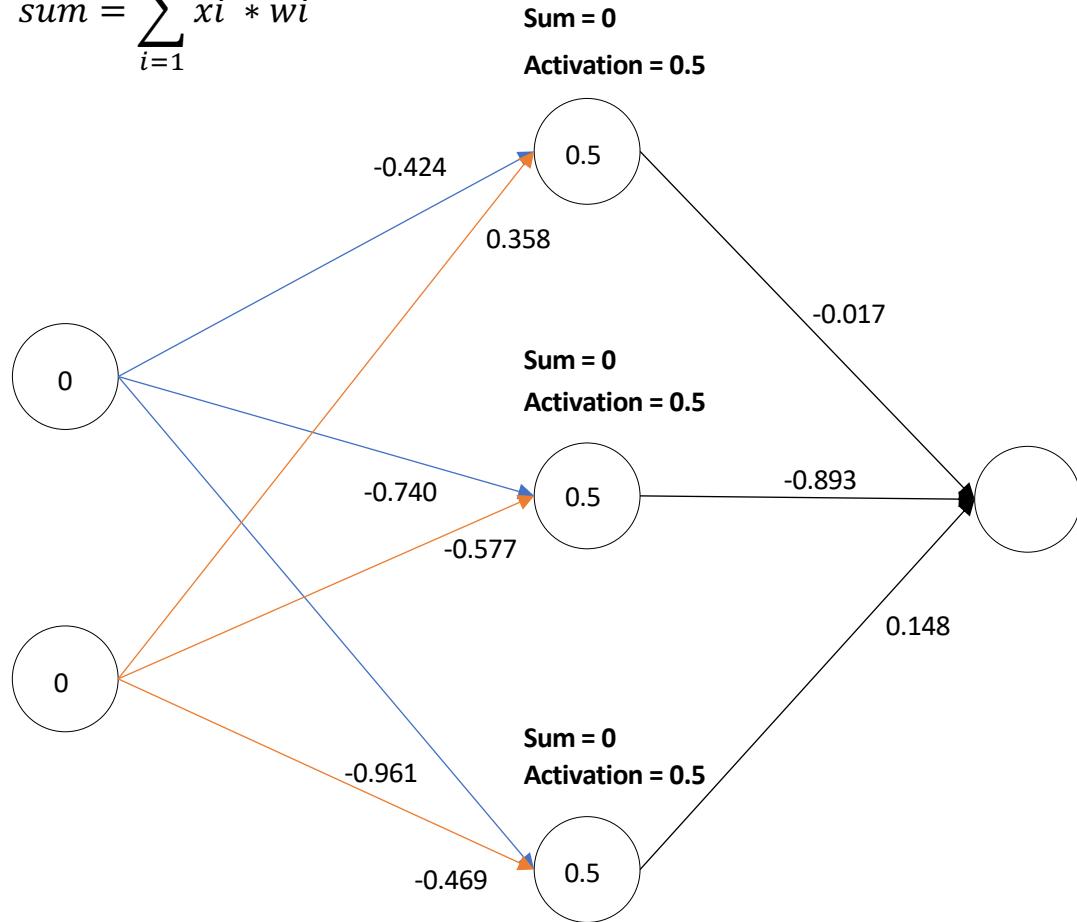
If X is low the result will be approximately 0

No negative number is returned

# “XOR” OPERATOR

X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0

$$sum = \sum_{i=1}^n x_i * w_i$$



$$y = \frac{1}{1 + e^{-x}}$$

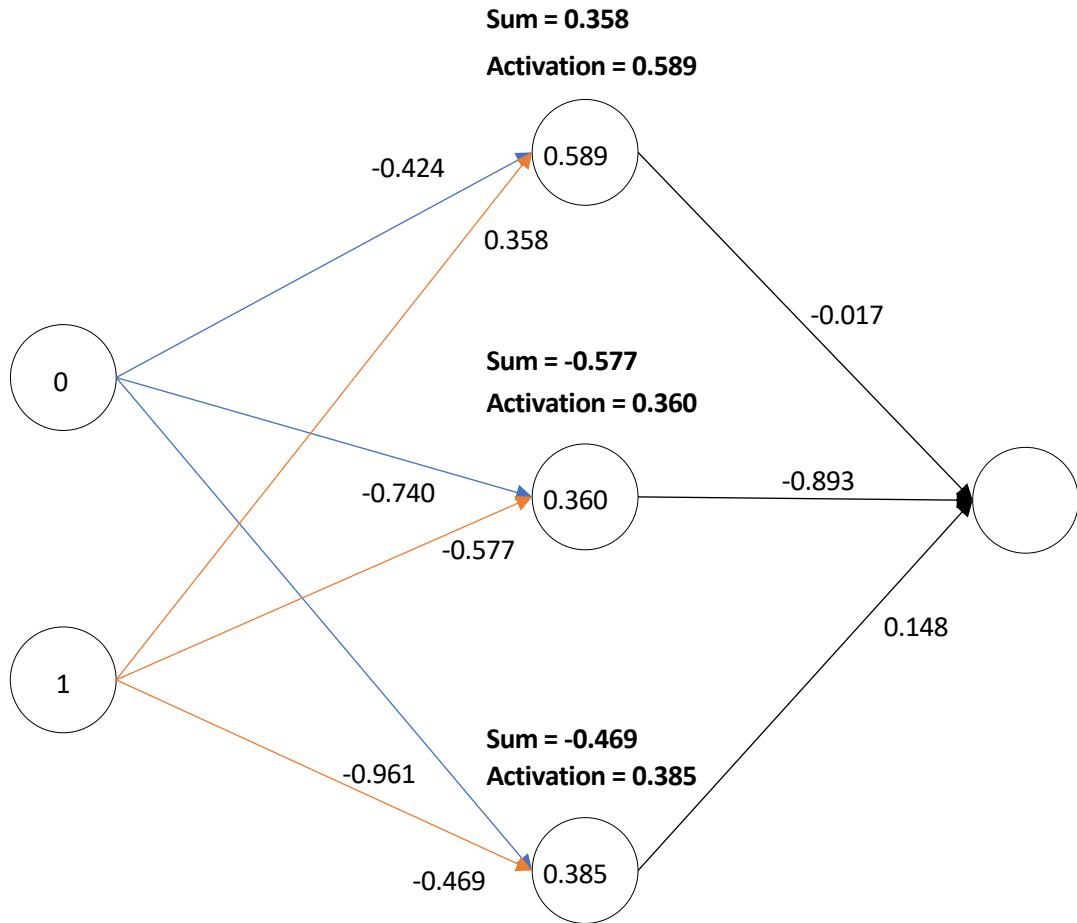
x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0

$$0 * (-0.424) + 0 * 0.358 = 0$$

$$0 * (-0.740) + 0 * (-0.577) = 0$$

$$0 * (-0.961) + 0 * (-0.469) = 0$$

x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0

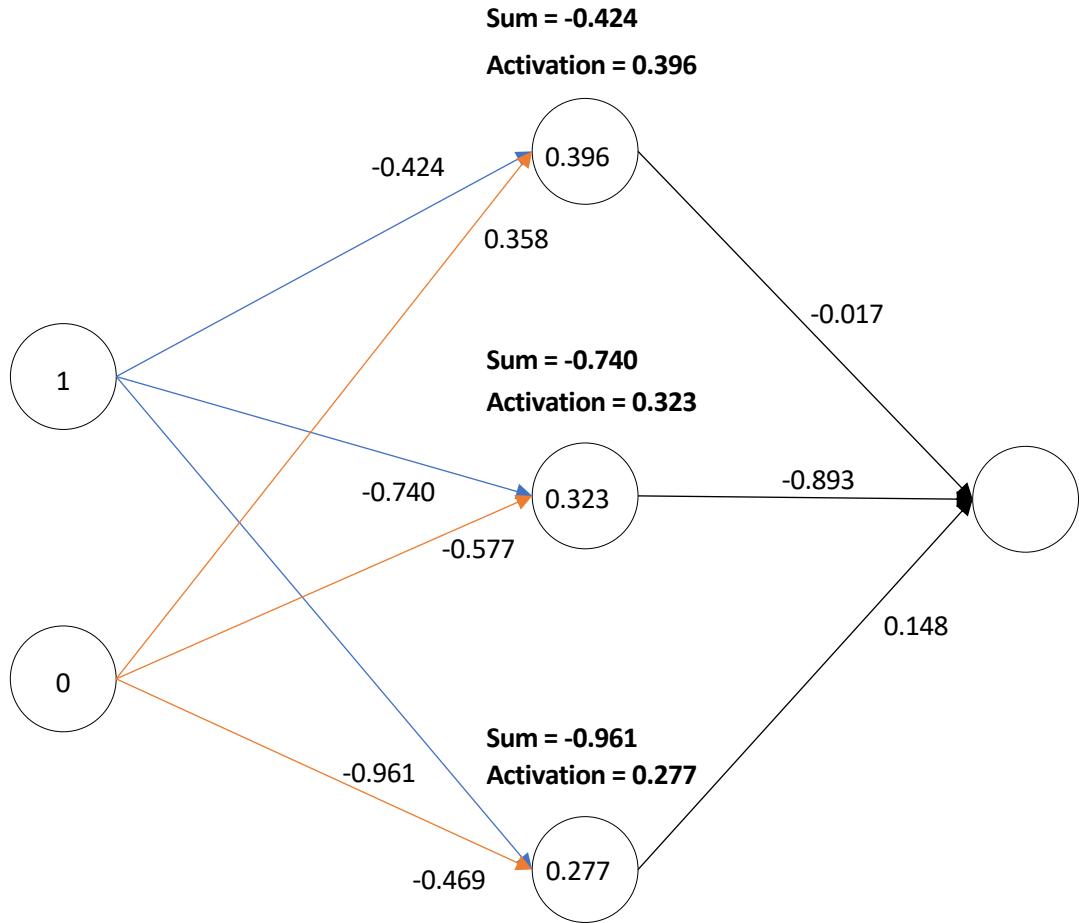


$$0 * (-0.424) + 1 * 0.358 = 0.358$$

$$0 * (-0.740) + 1 * (-0.577) = -0.577$$

$$0 * (-0.961) + 1 * (-0.469) = -0.469$$

x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0

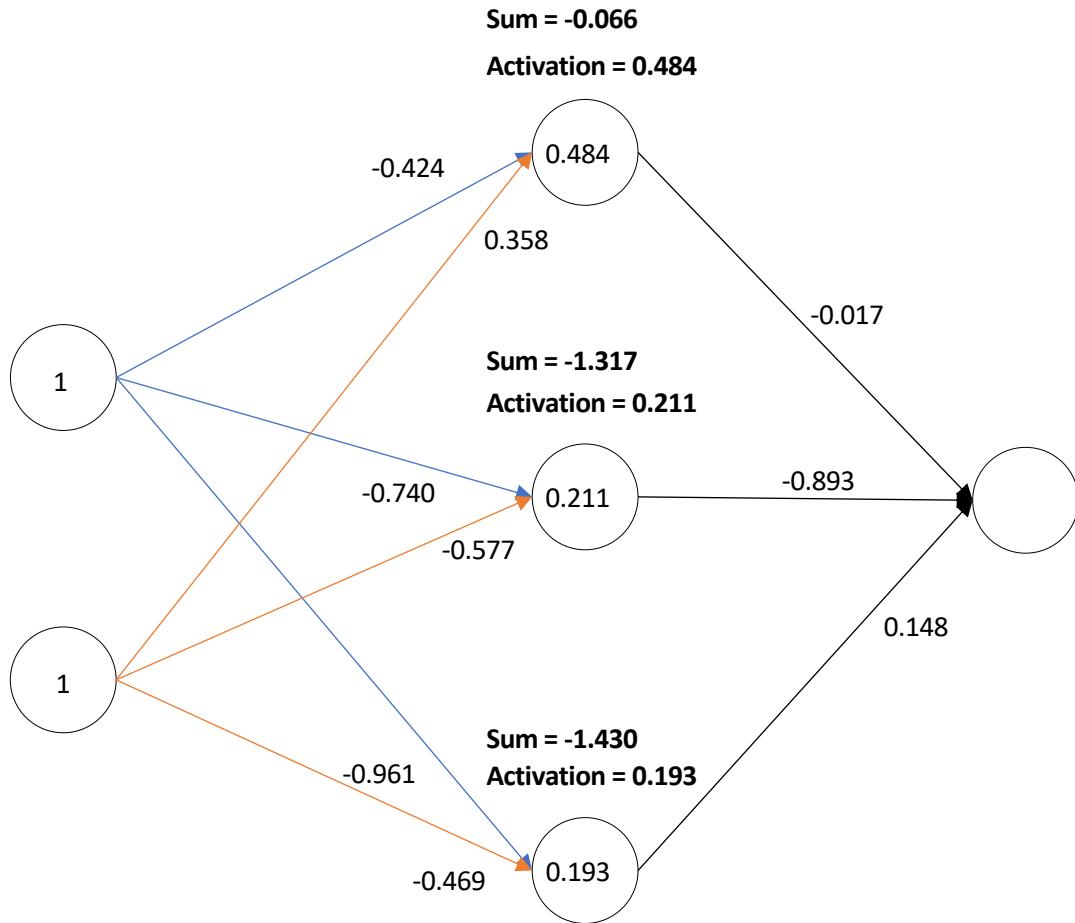


$$1 * (-0.424) + 0 * 0.358 = -0.424$$

$$1 * (-0.740) + 0 * (-0.577) = -0.740$$

$$1 * (-0.961) + 0 * (-0.469) = -0.961$$

x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0



$$1 * (-0.424) + 1 * 0.358 = -0.066$$

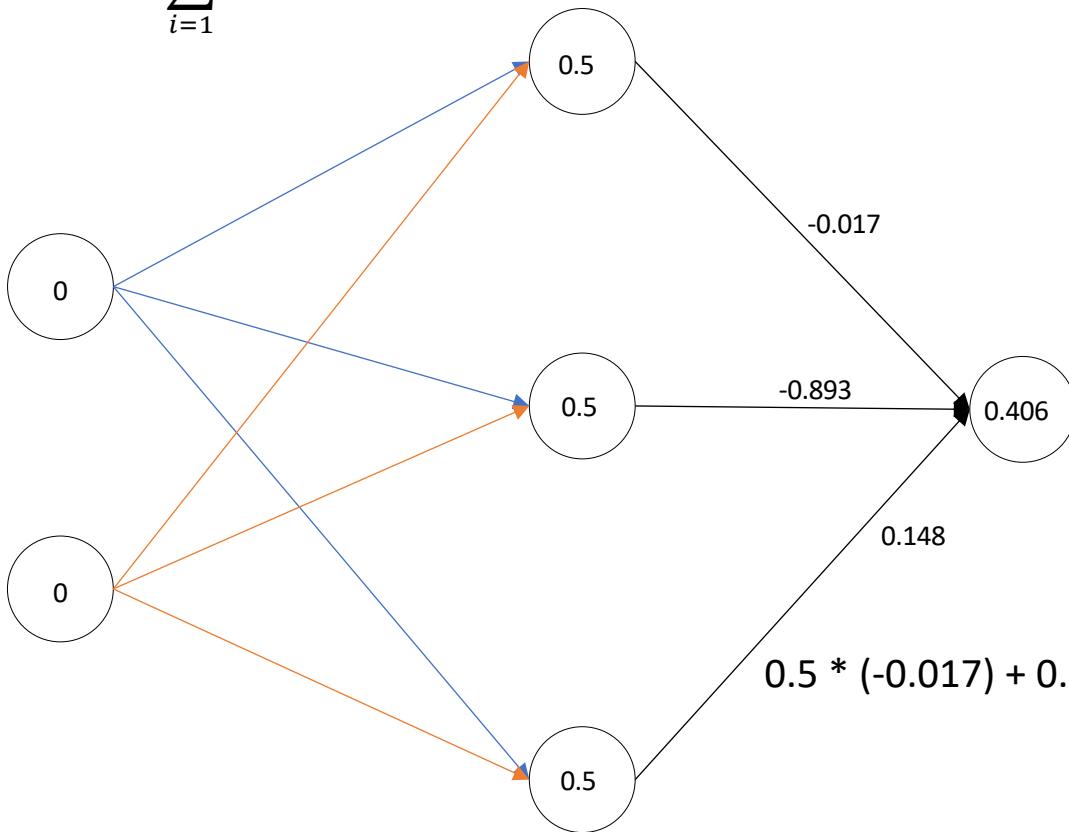
$$1 * (-0.740) + 1 * (-0.577) = -1.317$$

$$1 * (-0.961) + 1 * (-0.469) = -1.430$$

$$sum = \sum_{i=1}^n x_i * w_i$$

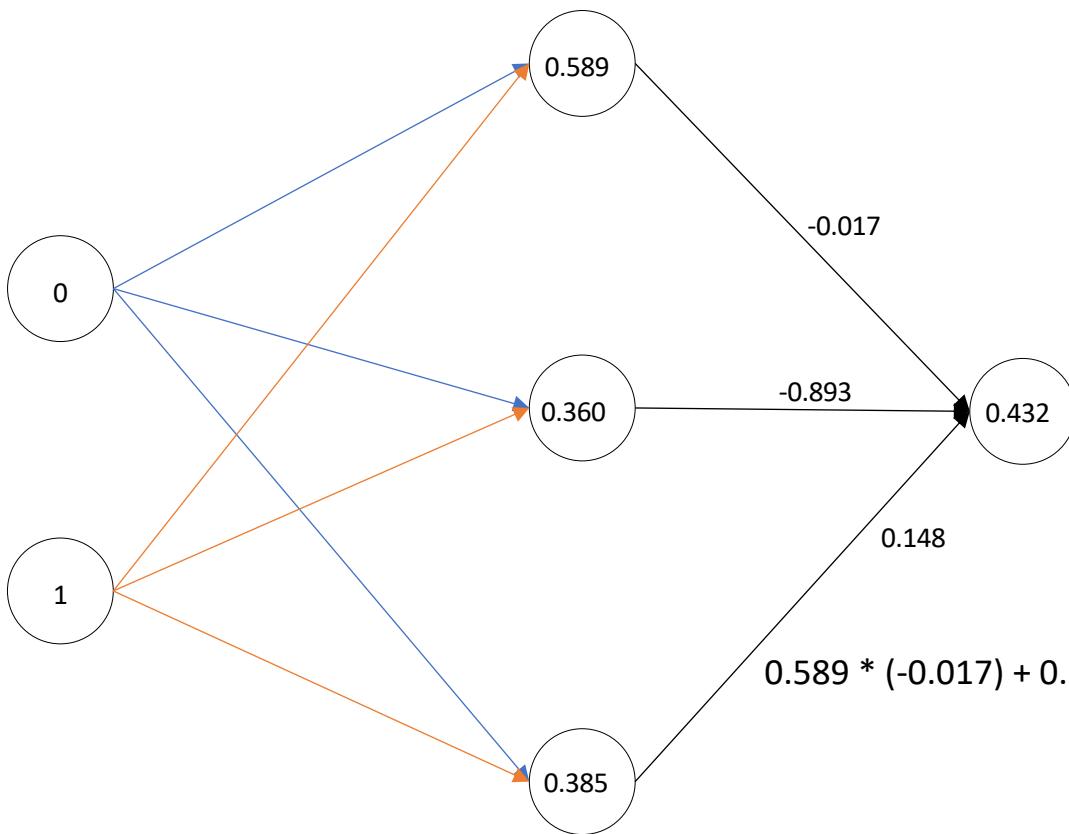
$$y = \frac{1}{1 + e^{-x}}$$

x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0



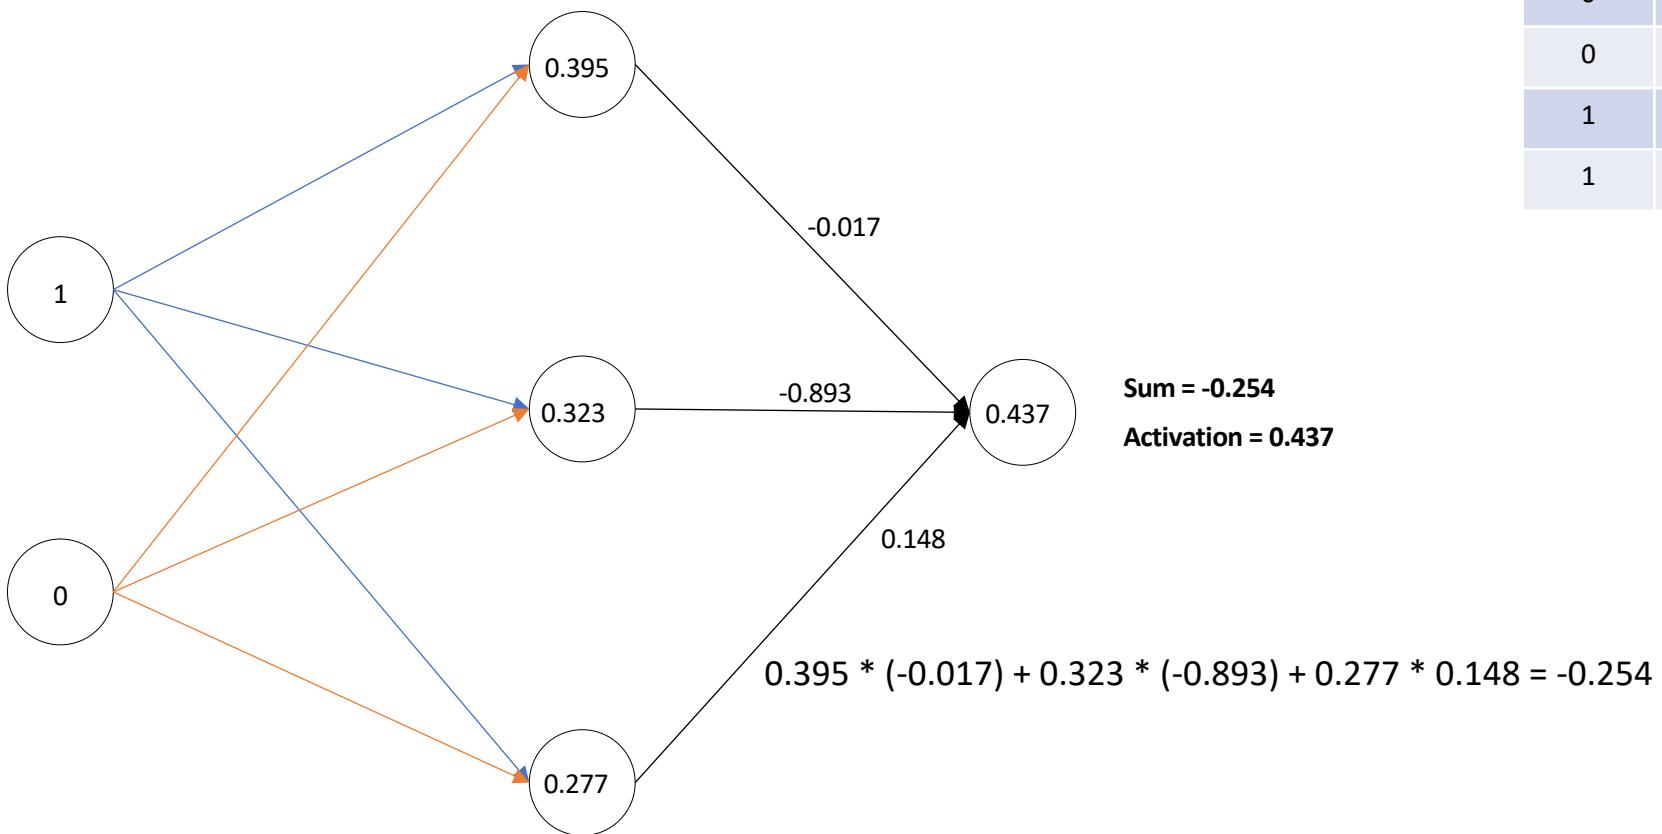
Sum = -0.381

Activation = 0.406

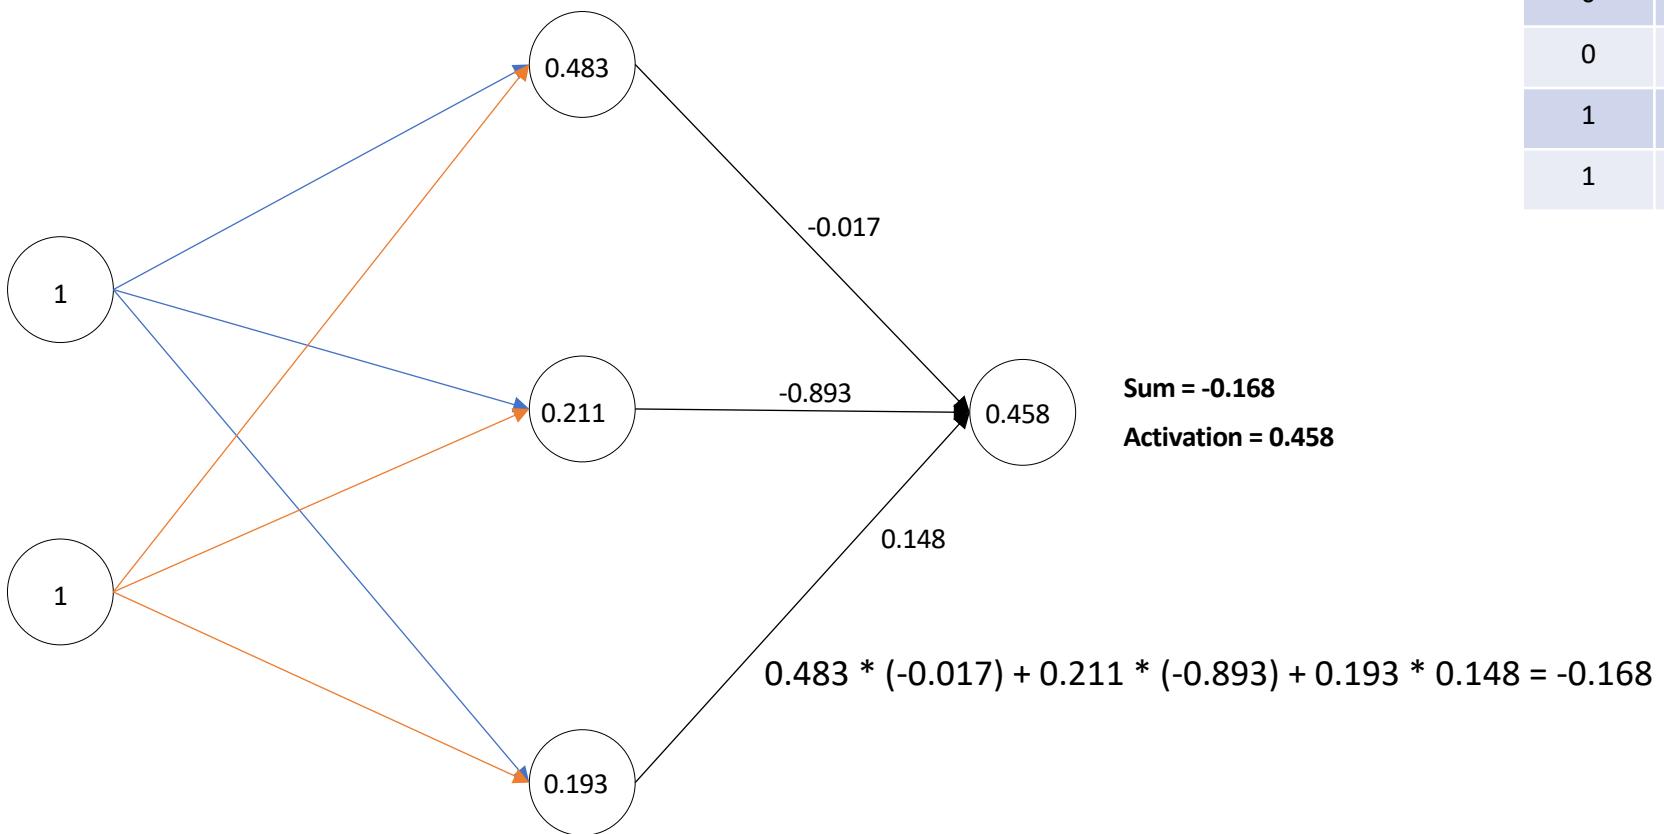


x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0

x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0



x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0



# ERROR (LOSS)

- Simplest algorithm
  - error = expectedOutput – prediction

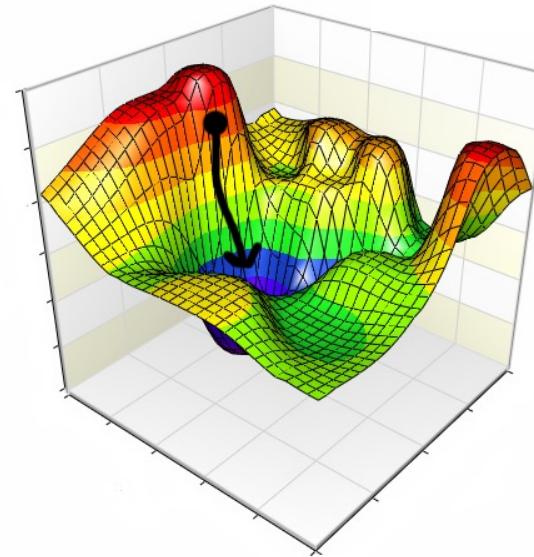
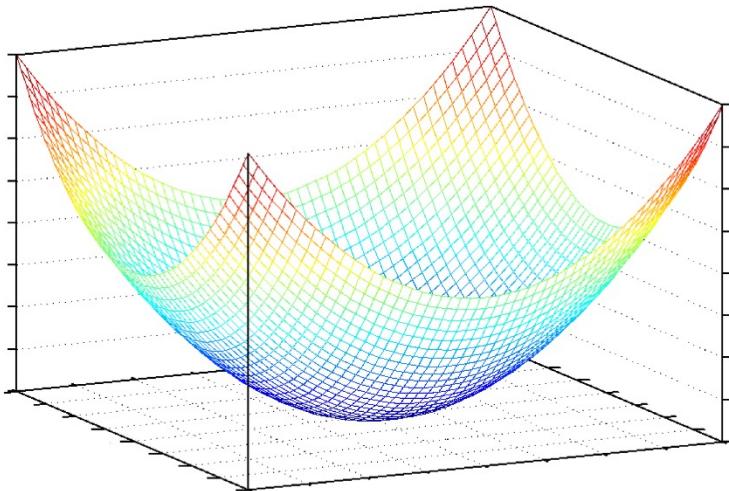
x1	x2	Class	Prediction	Error
0	0	0	0.406	-0.406
0	1	1	0.432	0.568
1	0	1	0.437	0.563
1	1	0	0.458	-0.458

Absolute average = 0.49

# GRADIENT

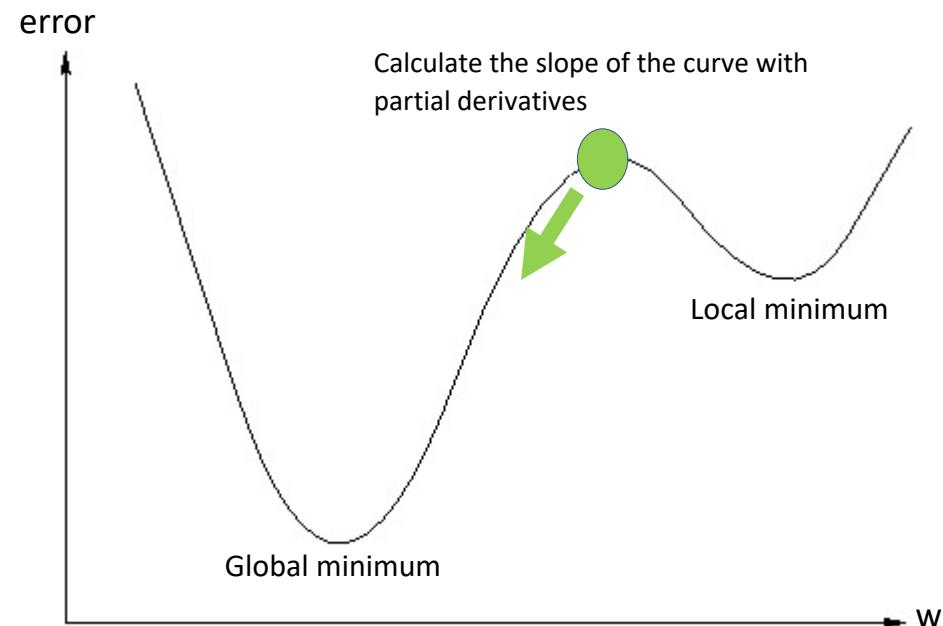
$$\min C(w_1, w_2 \dots w_n)$$

Calculate partial derivative to move to the gradient direction



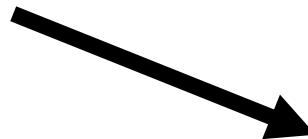
# GRADIENT

- Find the combination of weights where the error is as small as possible
- Gradient is calculated to know how much to adjust the weights



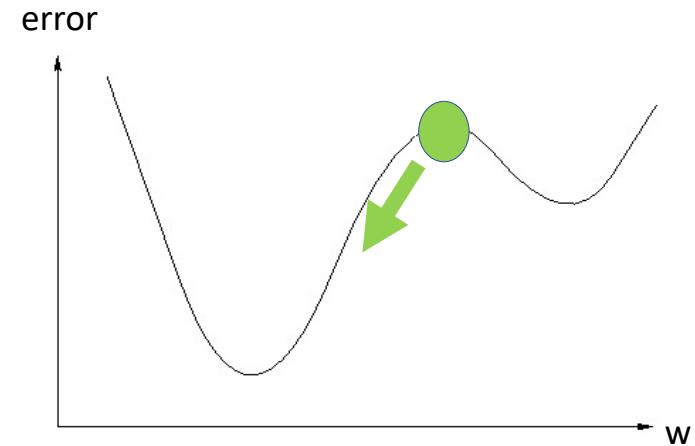
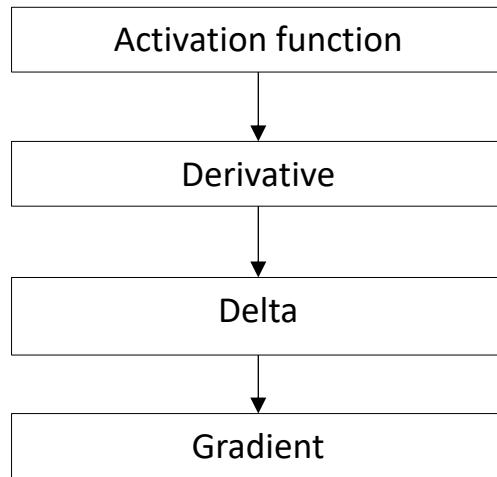
# GRADIENT (DERIVATIVE)

$$y = \frac{1}{1 + e^{-x}}$$

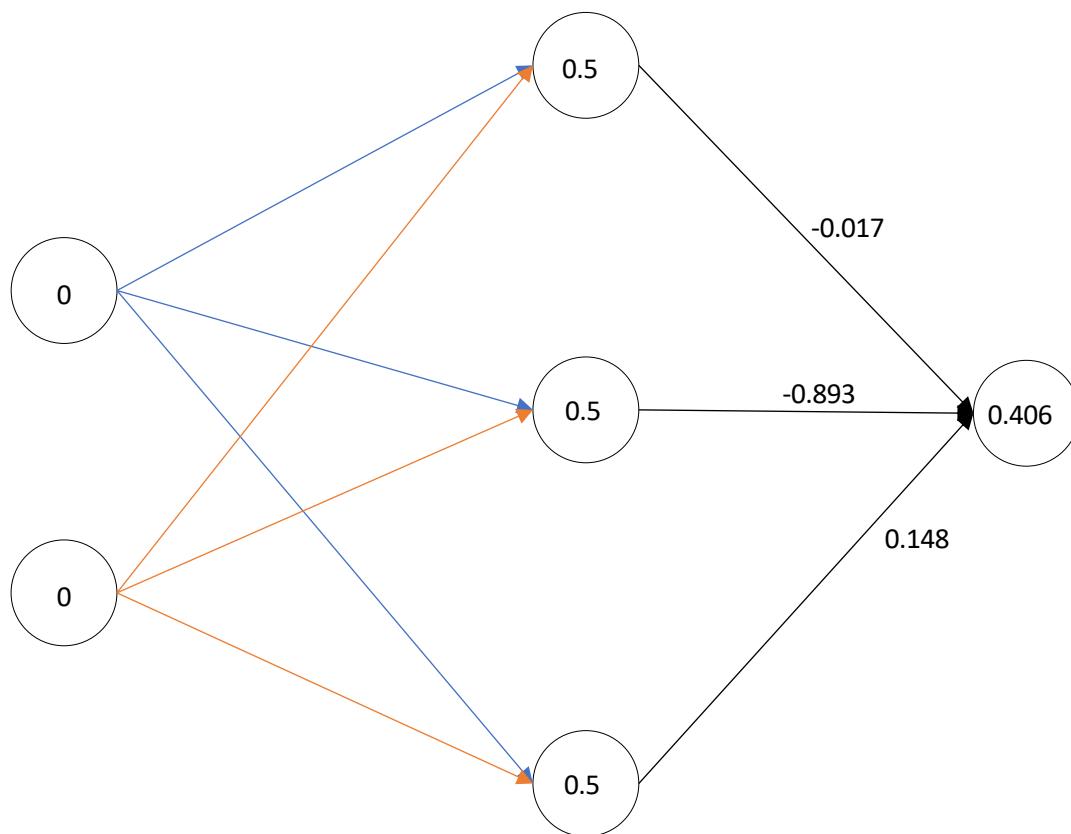


$$d = y * (1 - y)$$

# DELTA PARAMETER



$$\text{DeltaOutput} = \text{Error} * \text{SigmoidDerivative}$$



**Sum = -0.381**

**Activation = 0.406**

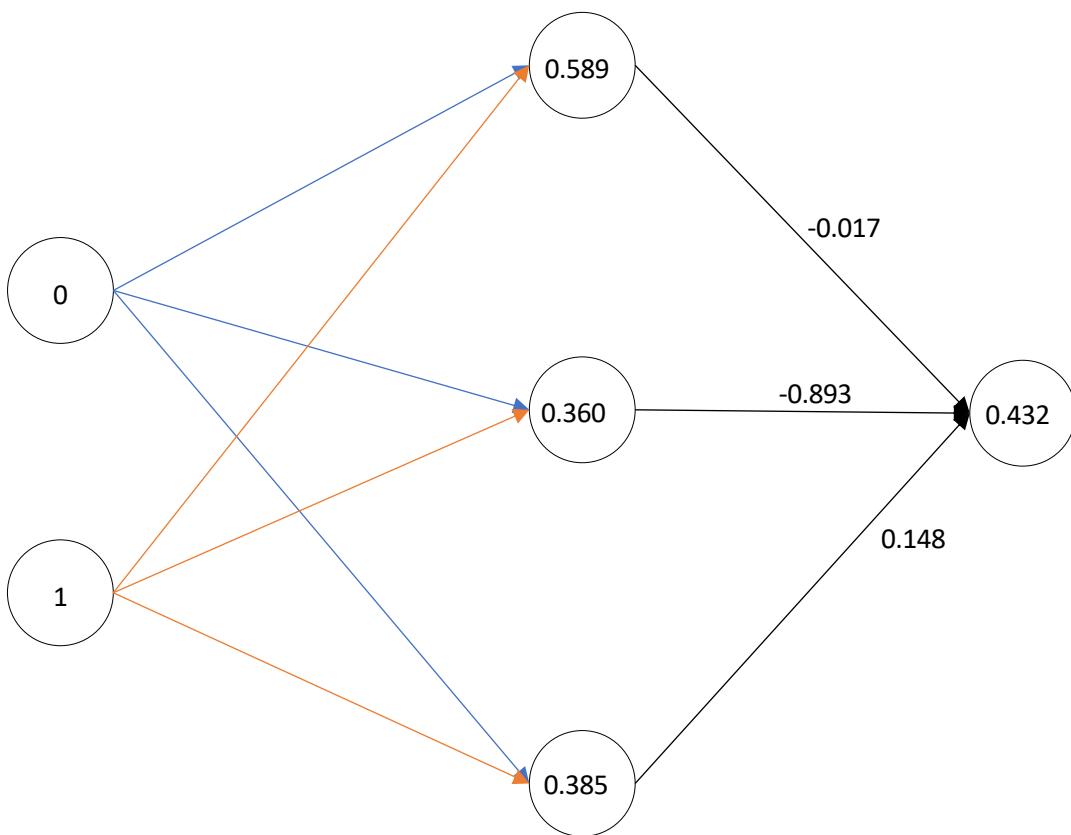
**Error = 0 – 0.406 = -0.406**

**Derivative (sigmoid) = 0.241**

**DeltaOutput = -0.406 \* 0.241 = -0.098**

x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0

$$\text{DeltaOutput} = \text{Error} * \text{SigmoidDerivative}$$



$$\text{Sum} = -0.274$$

$$\text{Activation} = 0.432$$

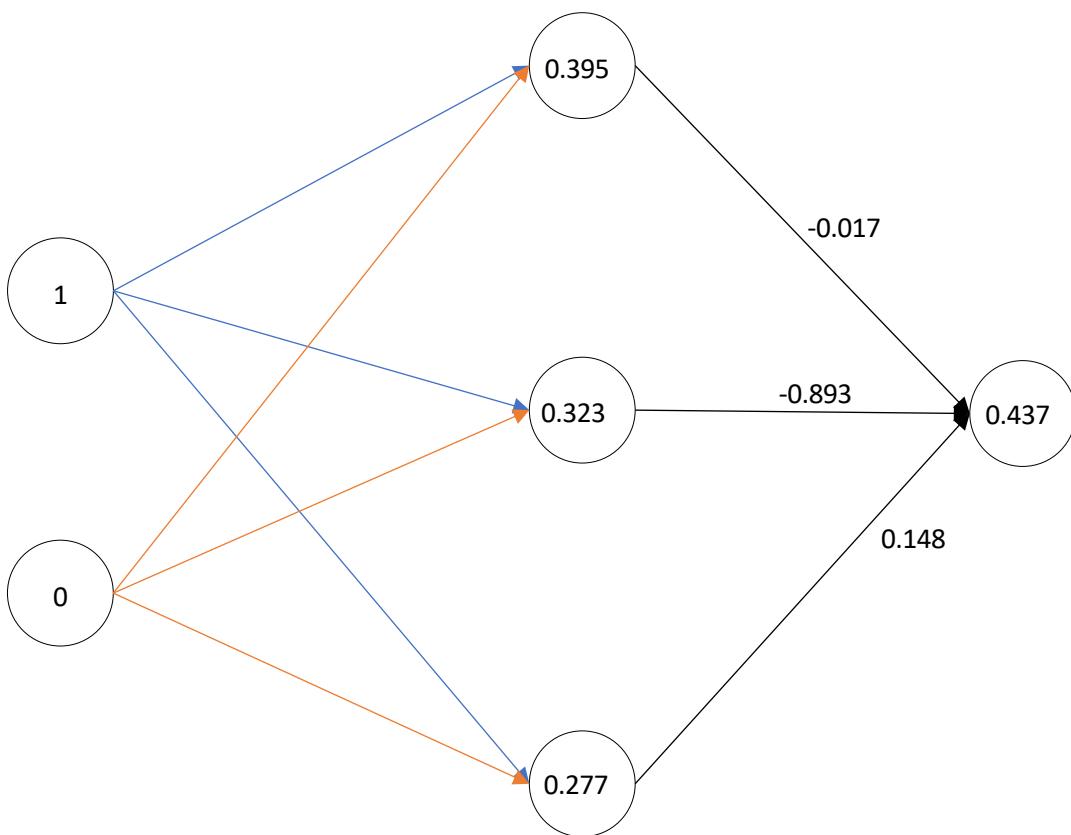
$$\text{Error} = 1 - 0.432 = 0.568$$

$$\text{Derivative (sigmoid)} = 0.245$$

$$\text{DeltaOutput} = 0.568 * 0.245 = 0.139$$

x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0

$$\text{DeltaOutput} = \text{Error} * \text{SigmoidDerivative}$$



Sum = -0.254

Activation = 0.437

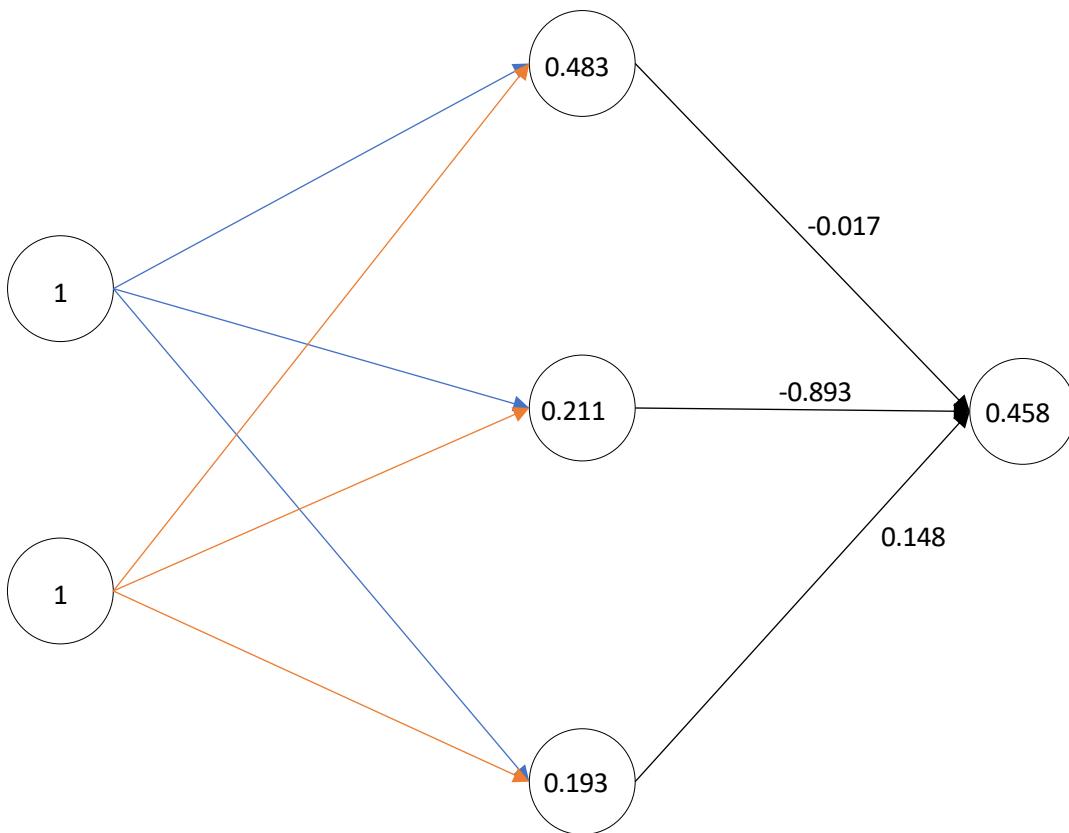
Error = 1 - 0.437 = 0.563

Derivative (sigmoid) = 0.246

DeltaOutput = 0.563 \* 0.246 = 0.139

x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0

$$\text{DeltaOutput} = \text{Error} * \text{SigmoidDerivative}$$



$$\text{Sum} = -0.168$$

$$\text{Activation} = 0.458$$

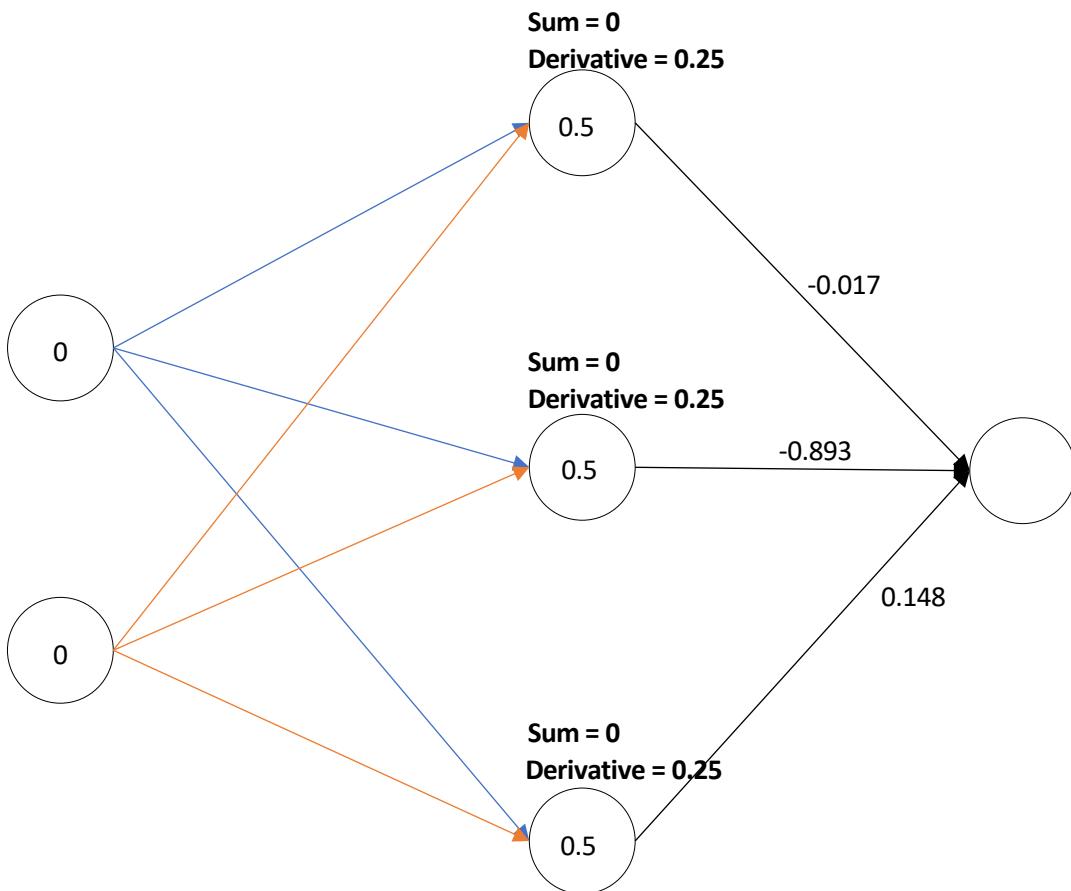
$$\text{Error} = 0 - 0.458 = -0.458$$

$$\text{Derivative (sigmoid)} = 0.248$$

$$\text{DeltaOutput} = -0.458 * 0.248 = -0.114$$

x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0

$$\Delta_{\text{Hidden}} = \text{DerivativeSigmoid} * \text{weight} * \Delta_{\text{Output}}$$



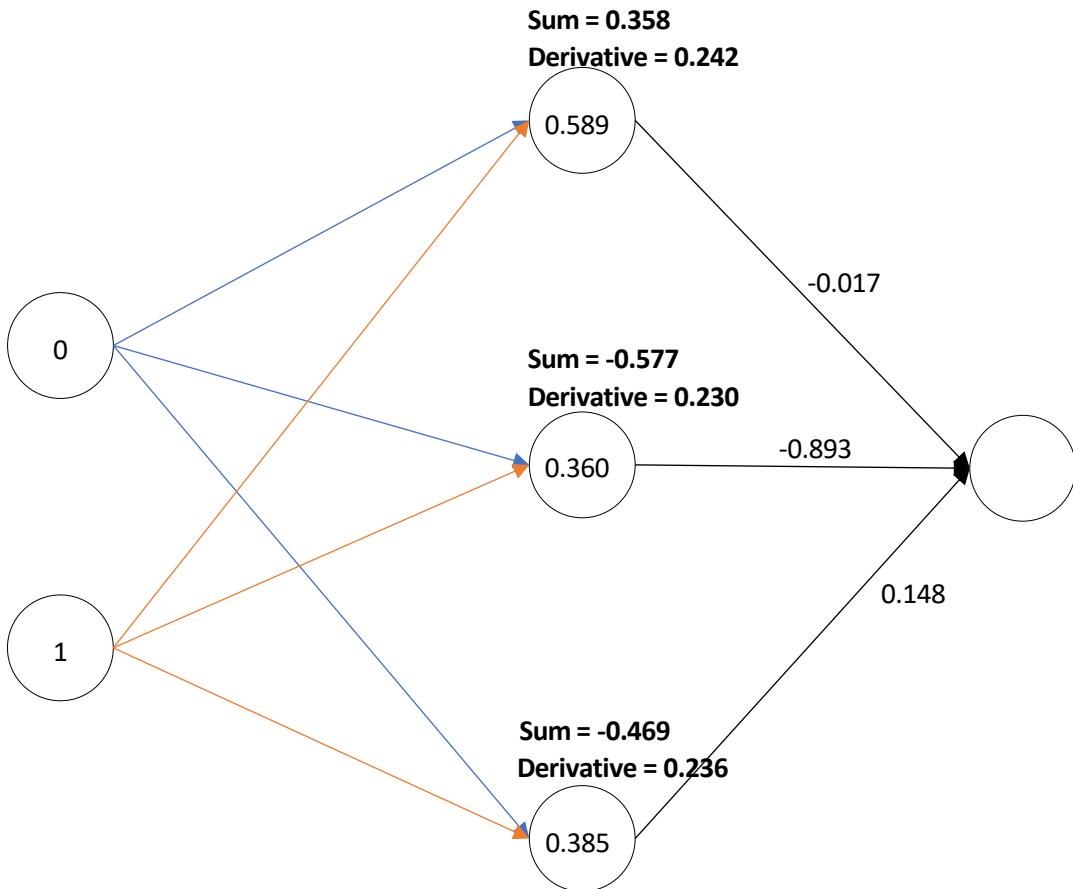
x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0

$$0.25 * (-0.017) * (-0.098) = 0.000$$

$$0.25 * (-0.893) * (-0.098) = 0.022$$

$$0.25 * 0.148 * (-0.098) = -0.004$$

$$\Delta_{\text{Hidden}} = \text{DerivativeSigmoid} * \text{weight} * \Delta_{\text{Output}}$$



x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0

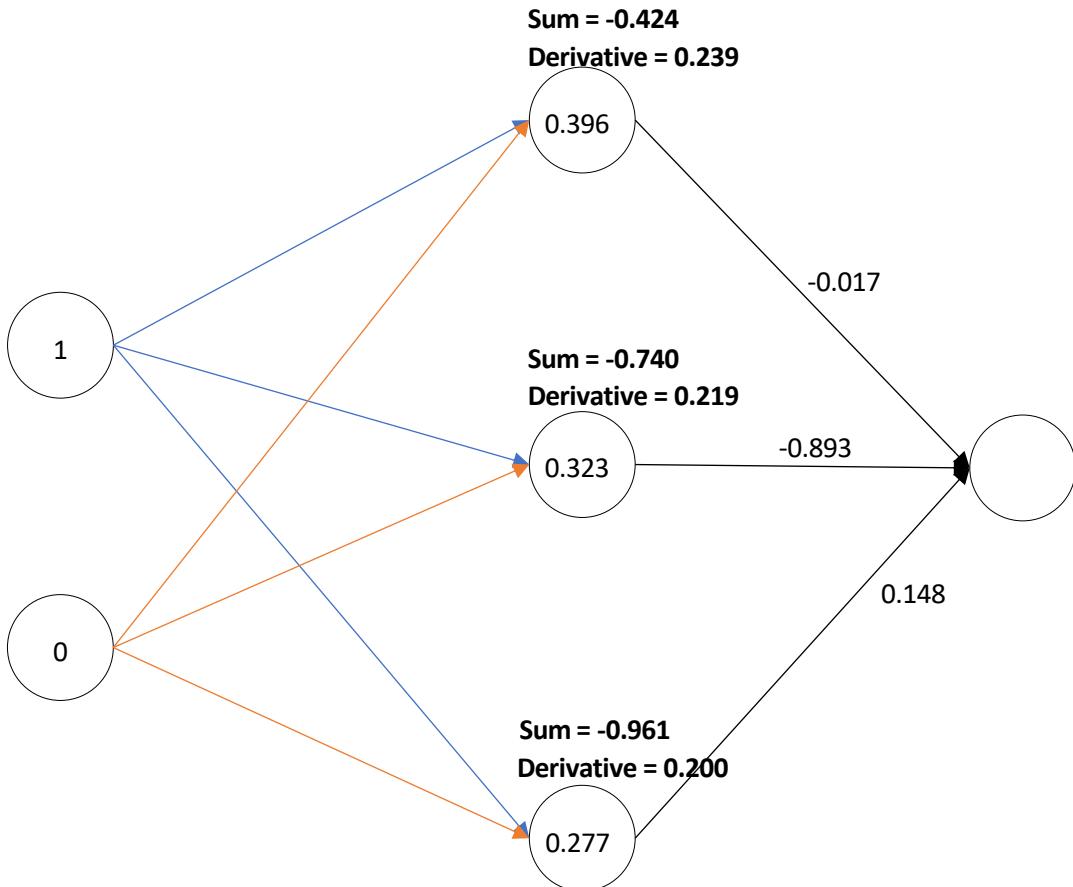
$$\Delta_{\text{Output}} = 0.139$$

$$0.242 * (-0.017) * 0.139 = -0.001$$

$$0.230 * (-0.893) * 0.139 = -0.029$$

$$0.236 * 0.148 * 0.139 = 0.005$$

$$\text{DeltaHidden} = \text{DerivativeSigmoid} * \text{weight} * \text{DeltaOutput}$$



x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0

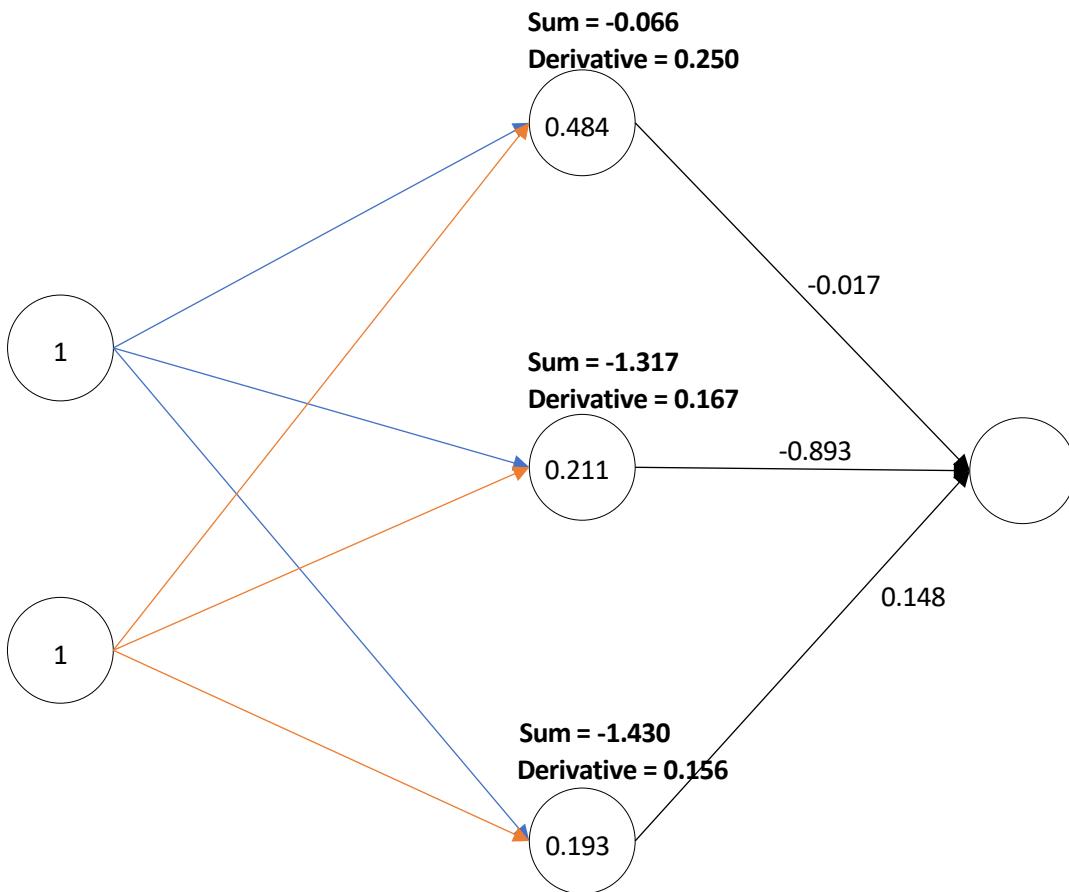
$$\text{DeltaOutput} = 0.139$$

$$0.239 * (-0.017) * 0.139 = -0.001$$

$$0.219 * (-0.893) * 0.139 = -0.027$$

$$0.200 * 0.148 * 0.139 = 0.004$$

$$\Delta_{\text{Hidden}} = \text{DerivativeSigmoid} * \text{weight} * \Delta_{\text{Output}}$$



x1	x2	Class
0	0	0
0	1	1
1	0	1
1	1	0

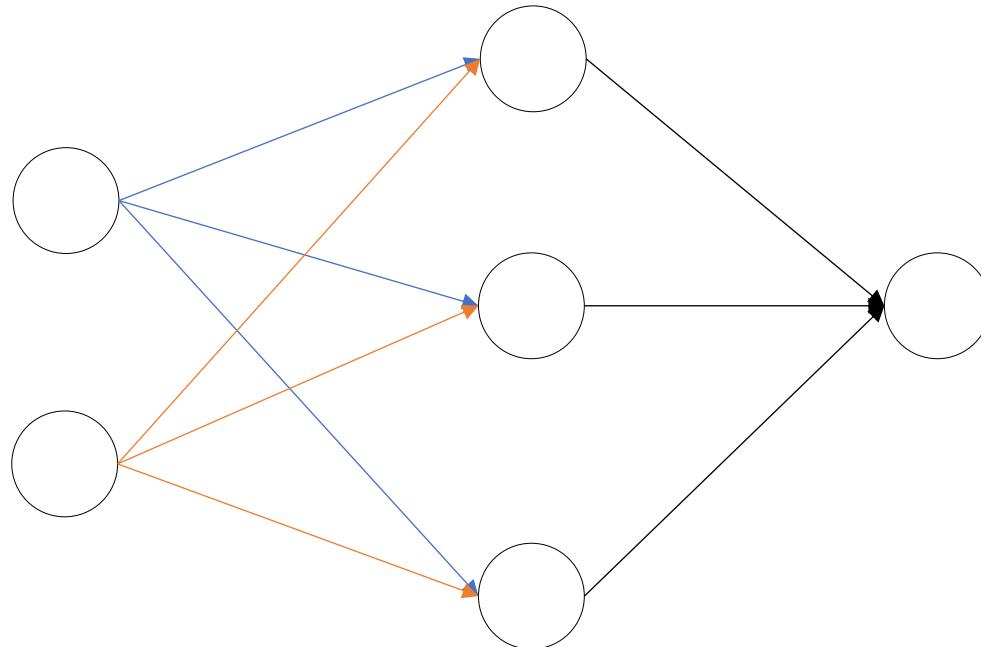
$$\Delta_{\text{Output}} = -0.114$$

$$0.250 * (-0.017) * (-0.114) = 0.000$$

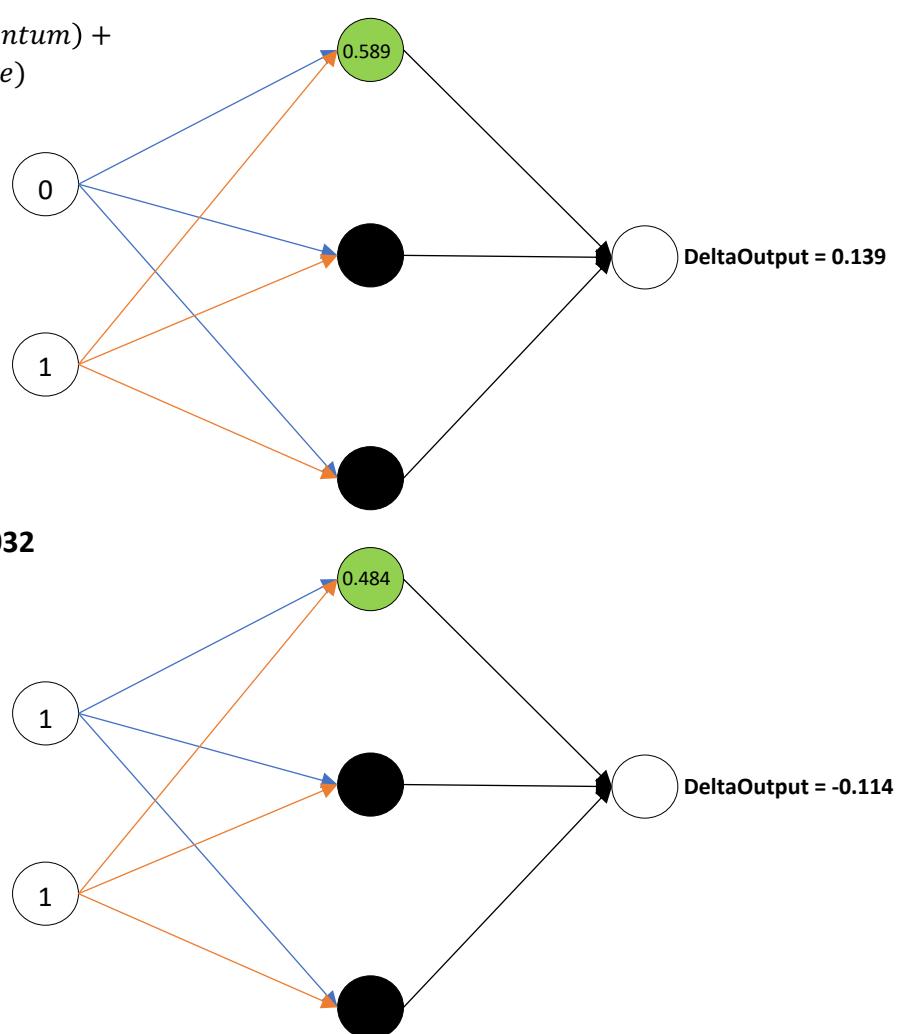
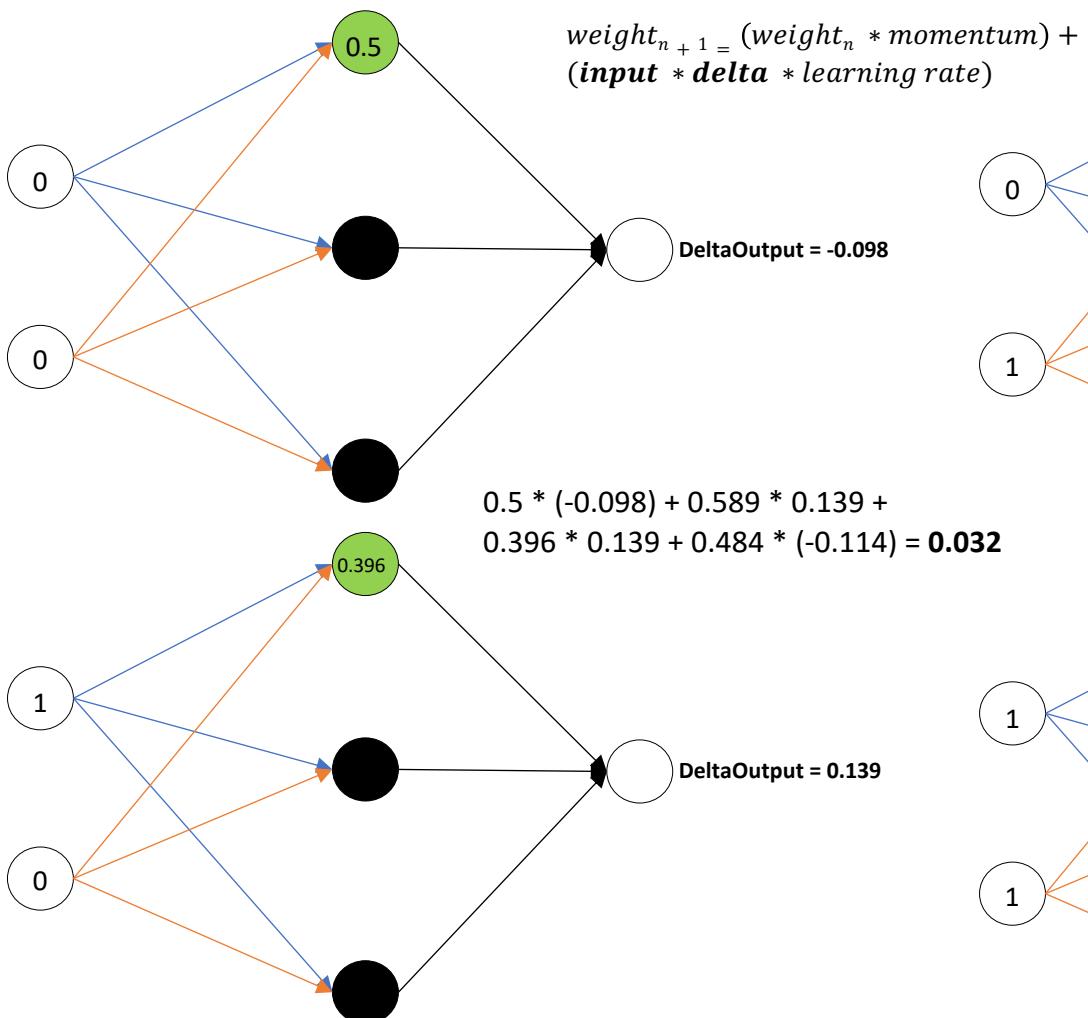
$$0.167 * (-0.893) * (-0.114) = 0.017$$

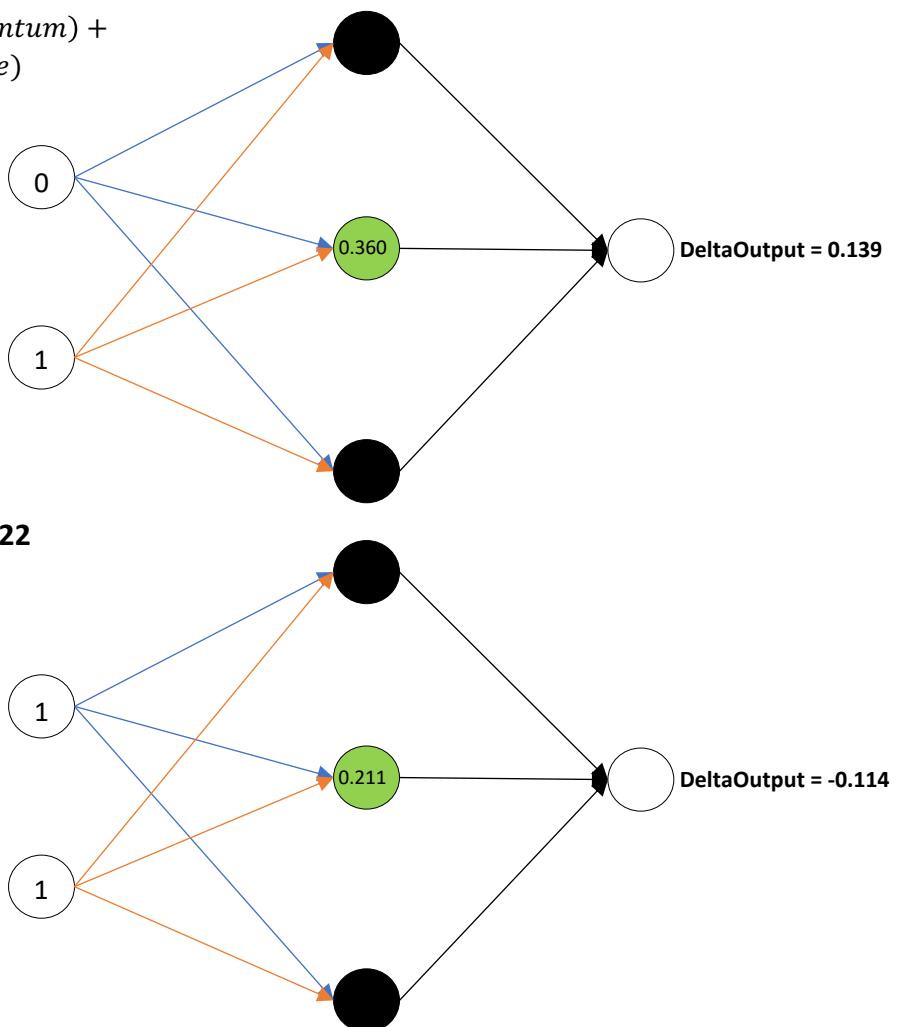
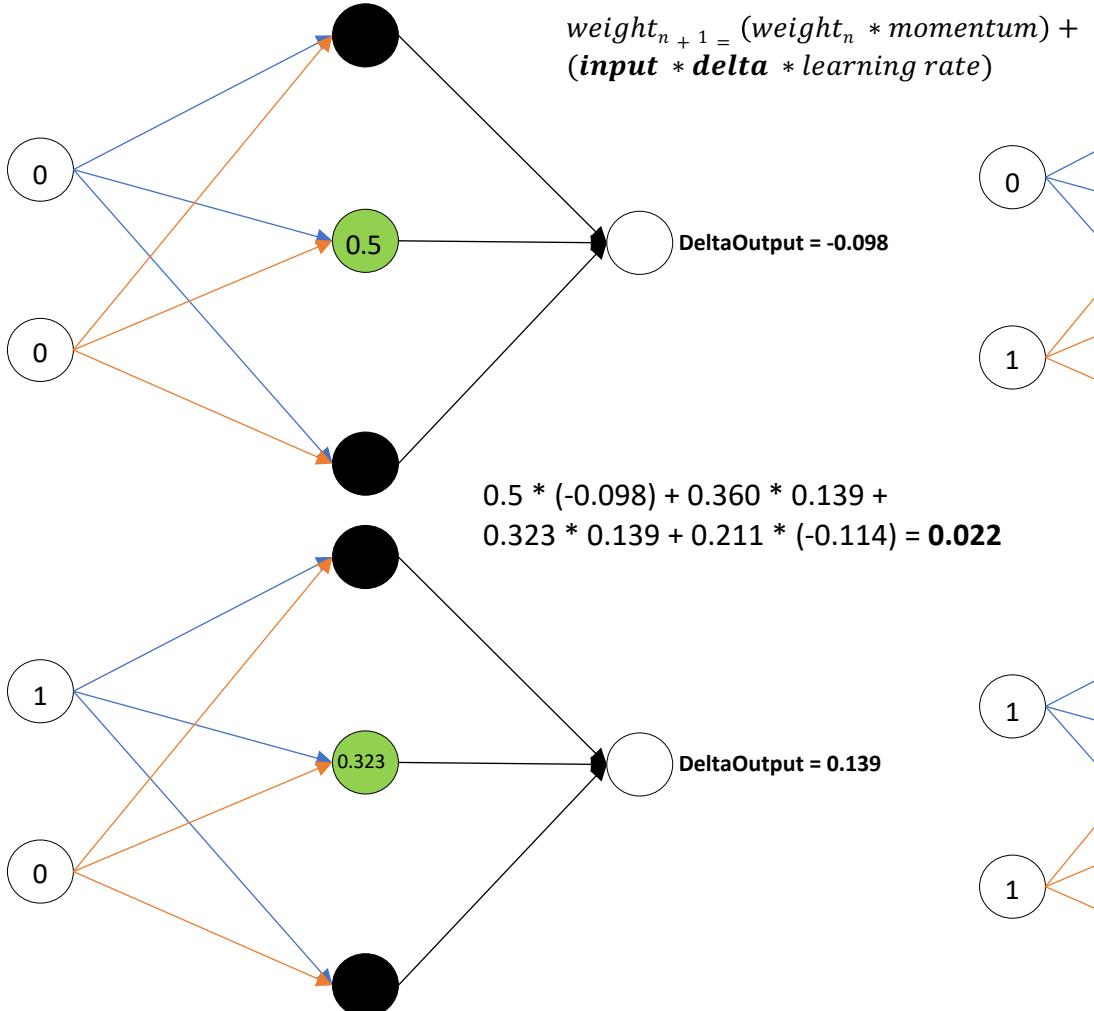
$$0.156 * 0.148 * (-0.114) = -0.003$$

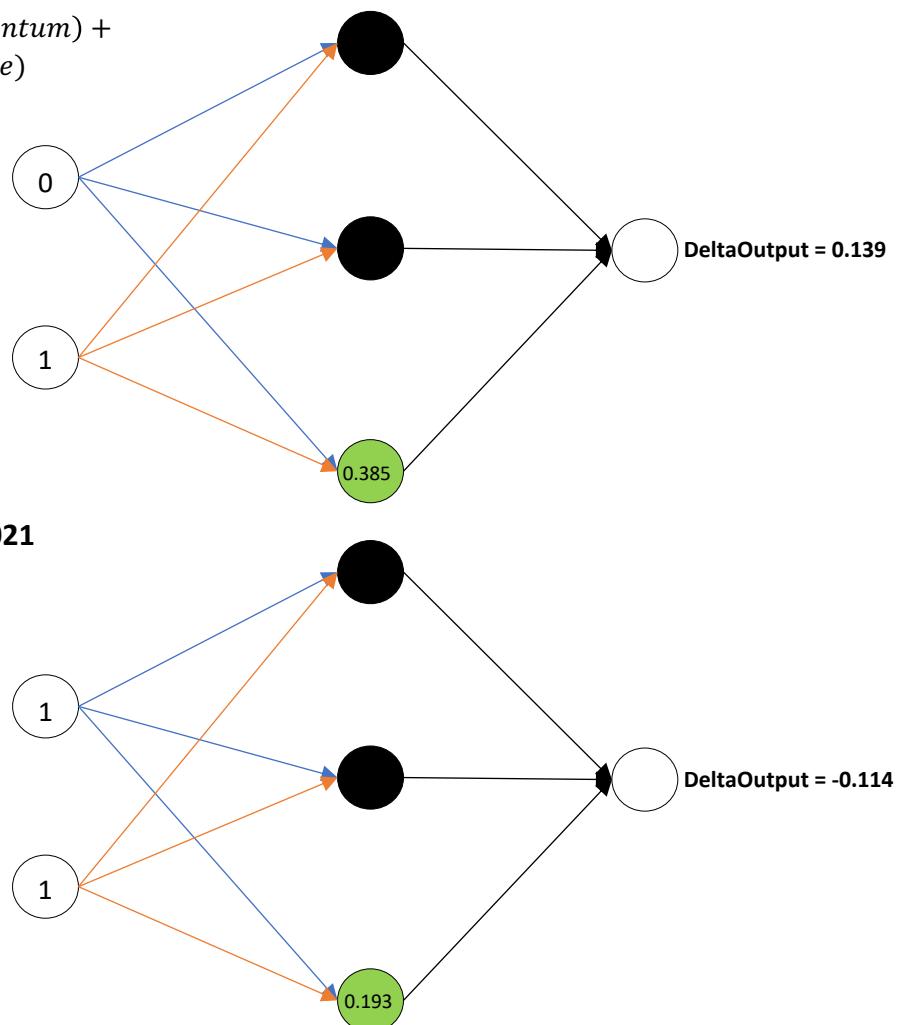
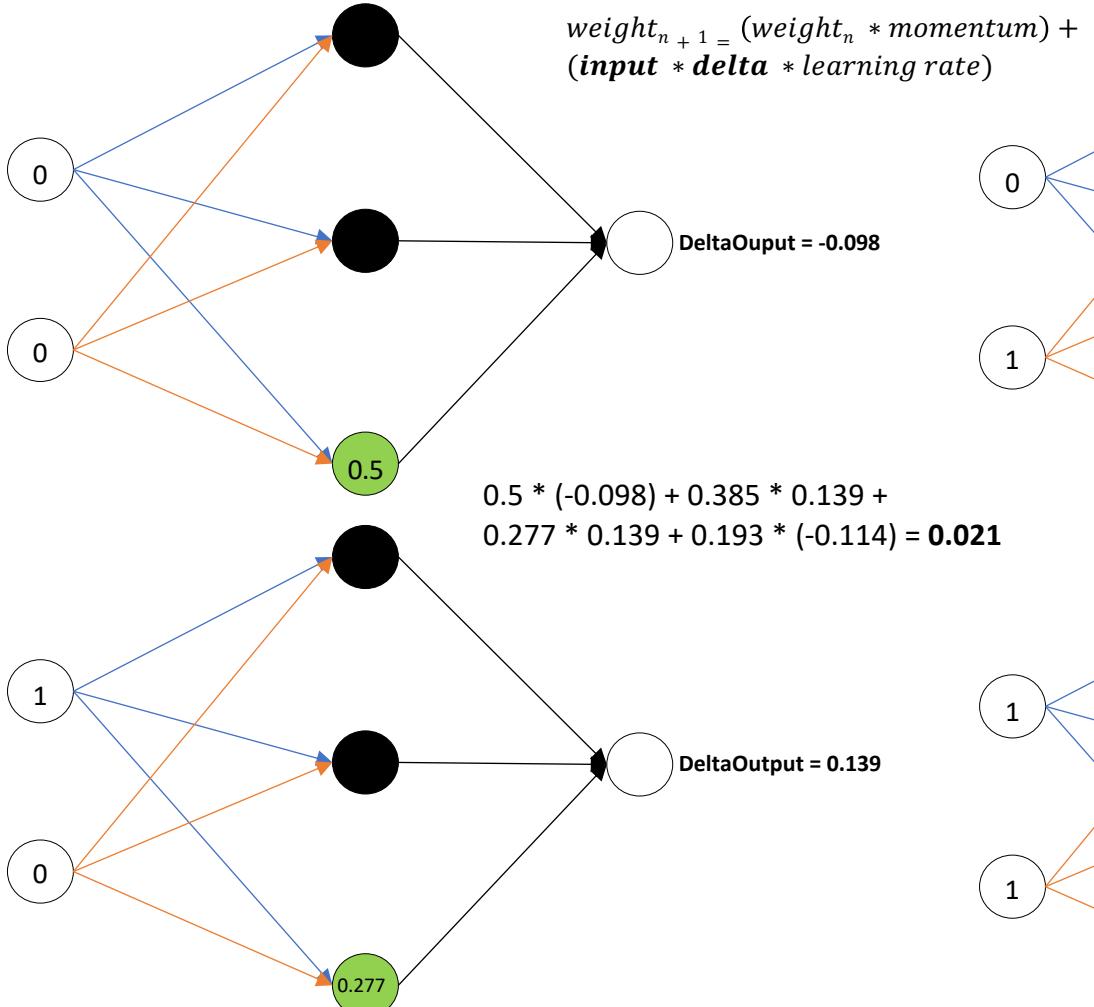
# BACKPROPAGATION



$$weight_{n+1} = (weight_n * momentum) + (input * delta * learning\ rate)$$







Learning rate = 0.3

Momentum = 1

Input x delta

0.032

0.022

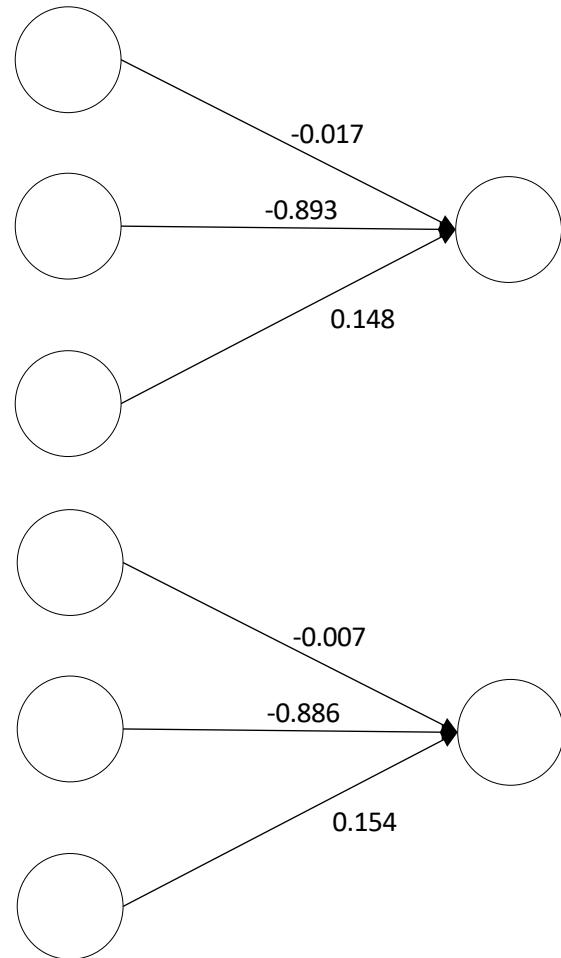
0.021

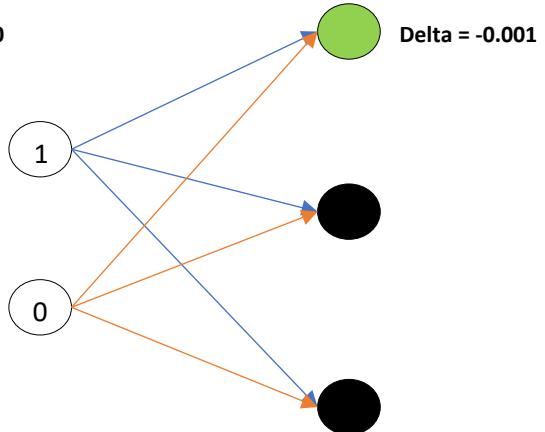
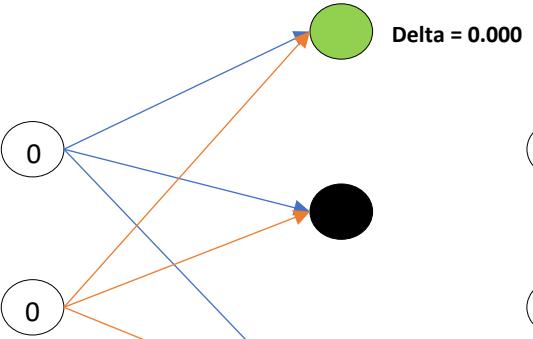
$$Weight_{n+1} = (weight_n * momentum) + (input * delta * learning rate)$$

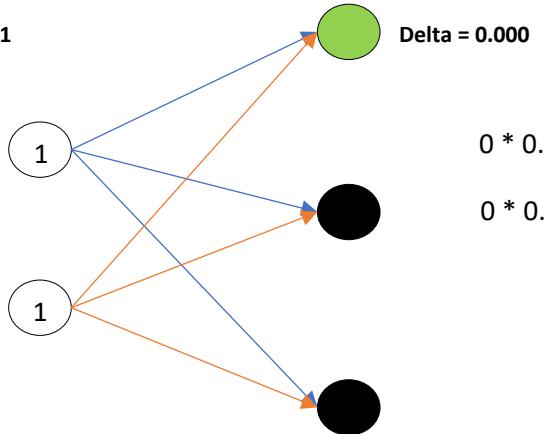
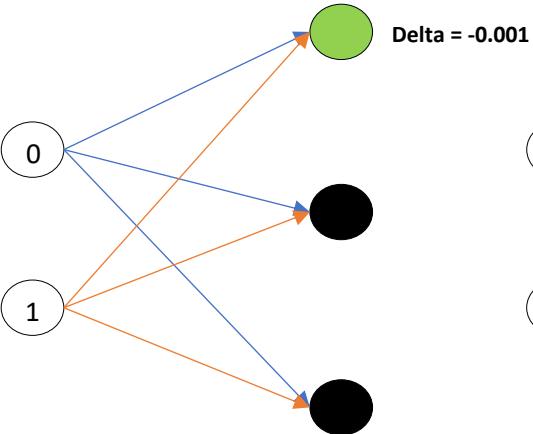
$$(-0.017 * 1) + 0.032 * 0.3 = \mathbf{-0.007}$$

$$(-0.893 * 1) + 0.022 * 0.3 = \mathbf{-0.886}$$

$$(0.148 * 1) + 0.021 * 0.3 = \mathbf{0.154}$$



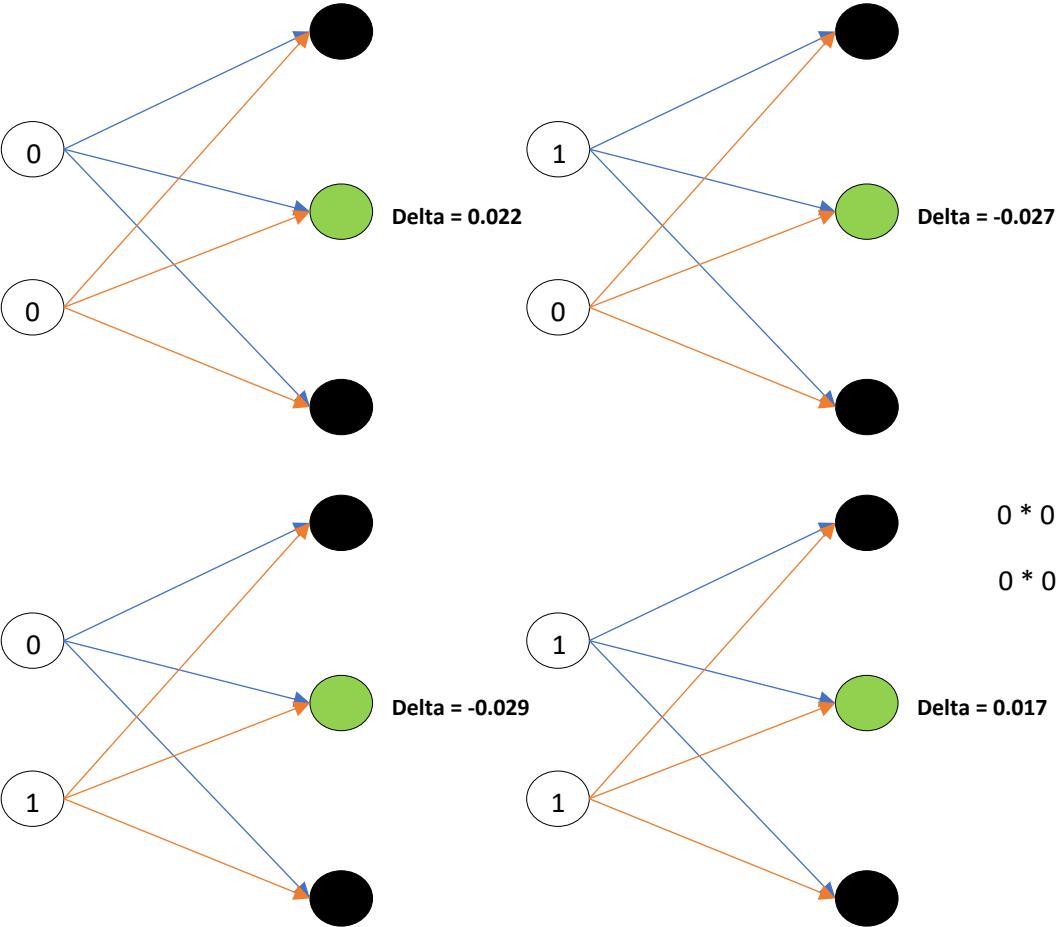


$$weight_{n+1} = (weight_n * momentum) + (input * delta * learning\ rate)$$


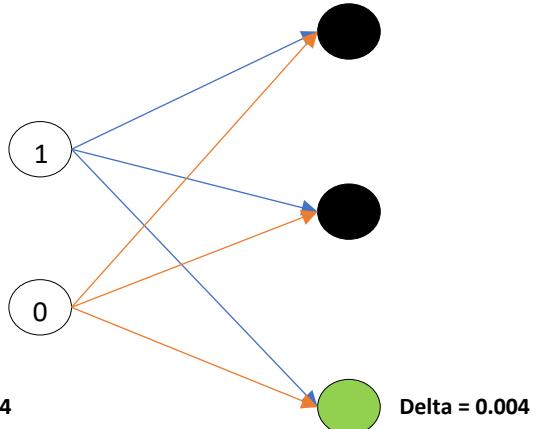
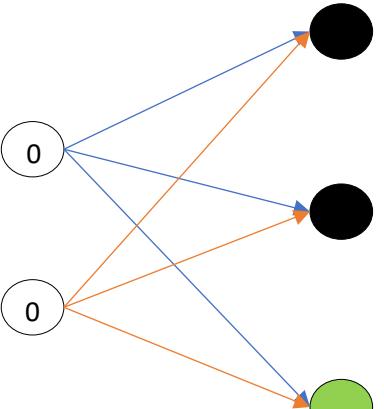
$$0 * 0.000 + 0 * (-0.001) + 1 * (-0.001) + 1 * 0.000 = -0.000$$

$$0 * 0.000 + 1 * (-0.001) + 0 * (-0.001) + 1 * 0.000 = -0.000$$

Rounded!



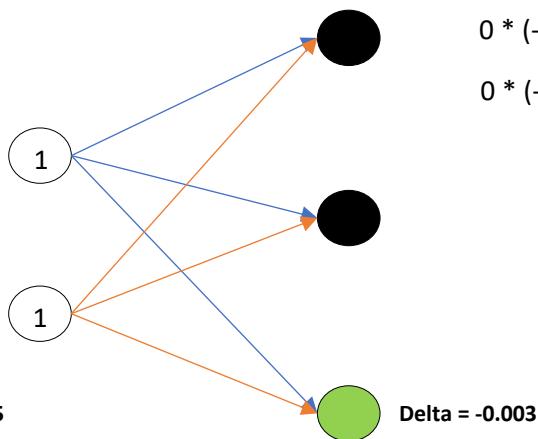
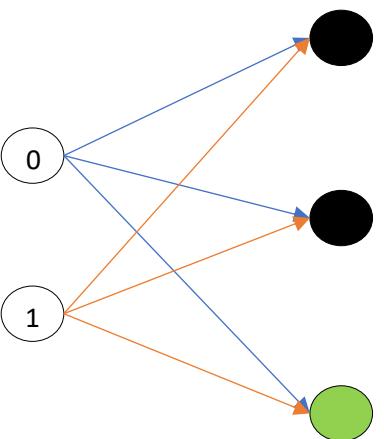
$$weight_{n+1} = (weight_n * momentum) + (input * delta * learning\ rate)$$



$$weight_{n+1} = (weight_n * momentum) + (input * delta * learning\ rate)$$

**Delta = -0.004**

**Delta = 0.004**



$$0 * (-0.004) + 0 * 0.005 + 1 * 0.004 + 1 * (-0.003) = 0.001$$

$$0 * (-0.004) + 1 * 0.005 + 0 * 0.004 + 1 * (-0.003) = 0.002$$

Learning rate = 0.3

Momentum = 1

Input x delta

-0.000 -0.010 0.001

-0.000 -0.012 0.002

$$Weight_{n+1} = (weight_n * momentum) + (input * delta * learning rate)$$

$$(-0.424 * 1) + (-0.000) * 0.3 = \mathbf{-0.424}$$

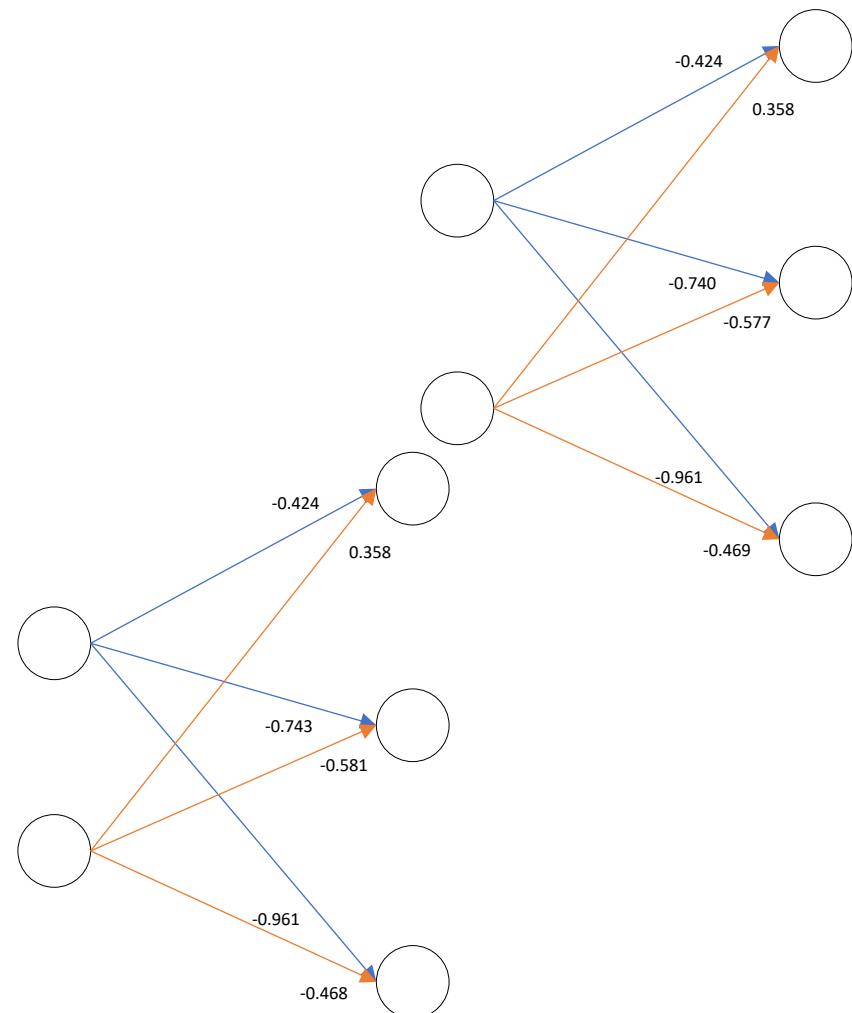
$$(0.358 * 1) + (-0.000) * 0.3 = \mathbf{0.358}$$

$$(-0.740 * 1) + (-0.010) * 0.3 = \mathbf{-0.743}$$

$$(-0.577 * 1) + (-0.012) * 0.3 = \mathbf{-0.581}$$

$$(-0.961 * 1) + 0.001 * 0.3 = \mathbf{-0.961}$$

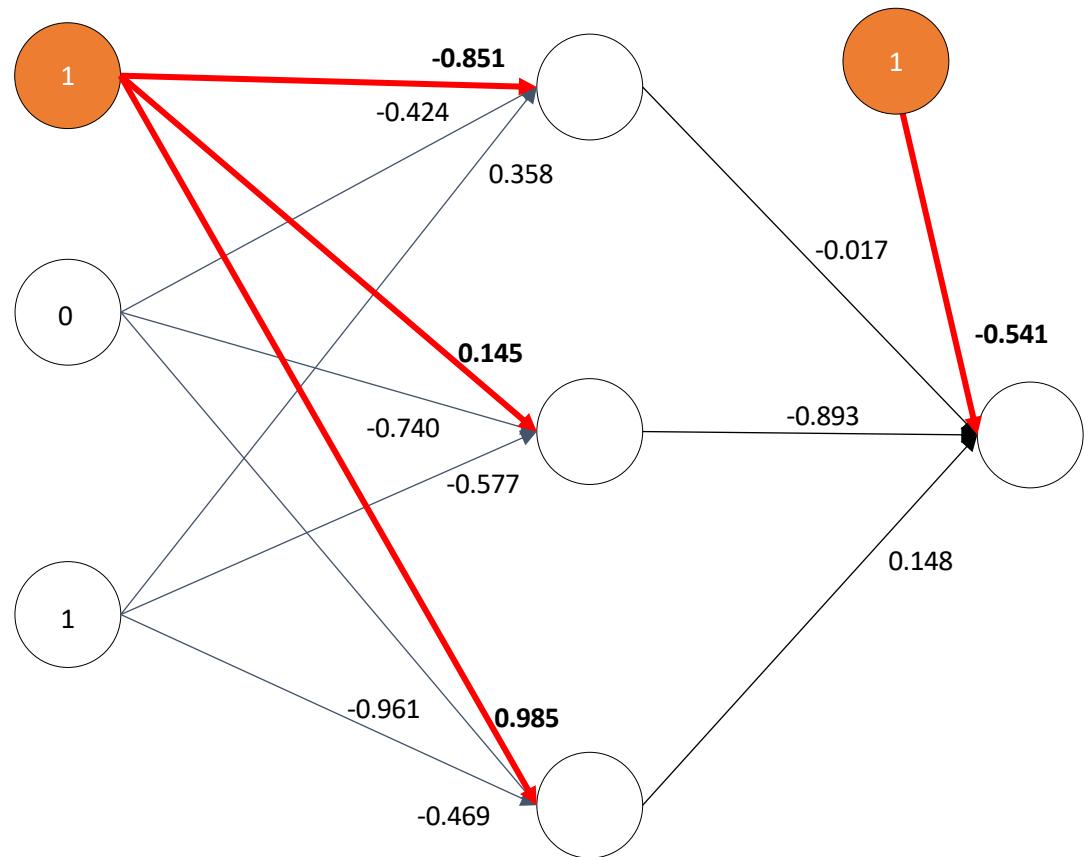
$$(-0.469 * 1) + 0.002 * 0.3 = \mathbf{-0.468}$$



# BIAS

Different values even if all entries are zero

Change the output



# ERROR (LOOS)

- Simplest algorithm
  - error = expected output – prediction

x1	x2	Class	Prediction	Error
0	0	0	0.406	-0.406
0	1	1	0.432	0.568
1	0	1	0.437	0.563
1	1	0	0.458	-0.458

Absolute average = 0.49

# MEAN SQUARED ERROR (MSE) AND ROOT MEAN SQUARED ERROR (RMSE)

$$MSE = \frac{1}{N} \sum_{i=1}^N (f_i - y_i)^2$$

$$RMSE = \sqrt{\frac{1}{N} \sum_{i=1}^N (f_i - o_i)^2}$$

x1	x2	Class	Prediction	Error
0	0	0	0.406	$(0 - 0.406)^2 = 0.164$
0	1	1	0.432	$(1 - 0.432)^2 = 0.322$
1	0	1	0.437	$(1 - 0.437)^2 = 0.316$
1	1	0	0.458	$(0 - 0.458)^2 = 0.209$

$$\text{Sum} = 1.011$$

$$MSE = 1.011 / 4 = 0.252$$

$$RMSE = 0.501$$

Credit history	Debts	Properties	Income	Risk
3	1	1	1	100
2	1	1	2	100
2	2	1	2	010
2	2	1	3	100
2	2	1	3	001
2	2	2	3	001
3	2	1	1	100
3	2	2	3	010
1	2	1	3	001
1	1	2	3	001
1	1	1	1	100
1	1	1	2	010
1	1	1	3	001
3	1	1	2	100

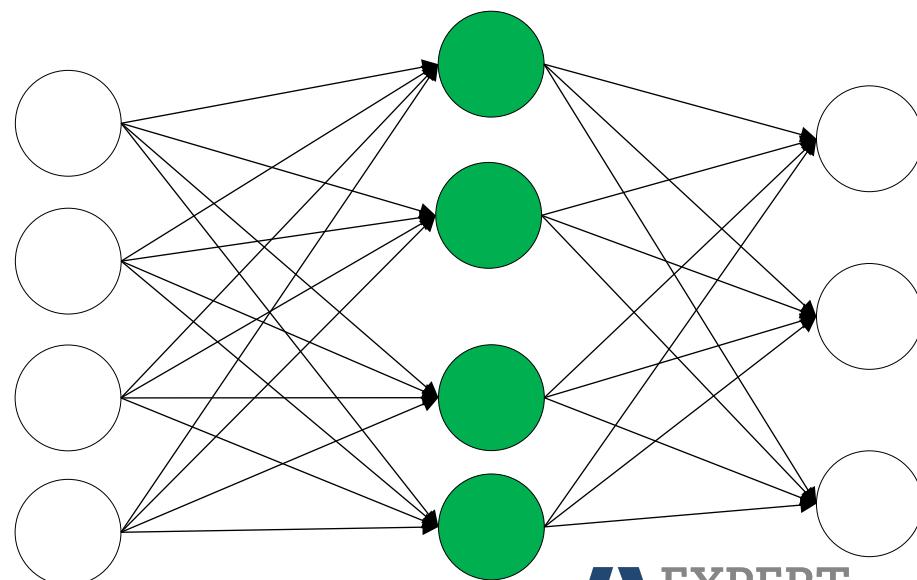
Calculates the error for all rows and updates the weights

## Batch gradient descent

Credit history	Debts	Properties	Income	Risk
3	1	1	1	100
2	1	1	2	100
2	2	1	2	010
2	2	1	3	100
2	2	1	3	001
2	2	2	3	001
3	2	1	1	100
3	2	2	3	010
1	2	1	3	001
1	1	2	3	001
1	1	1	1	100
1	1	1	2	010
1	1	1	3	001
3	1	1	2	100

Calculates the error for each row and updates the weights

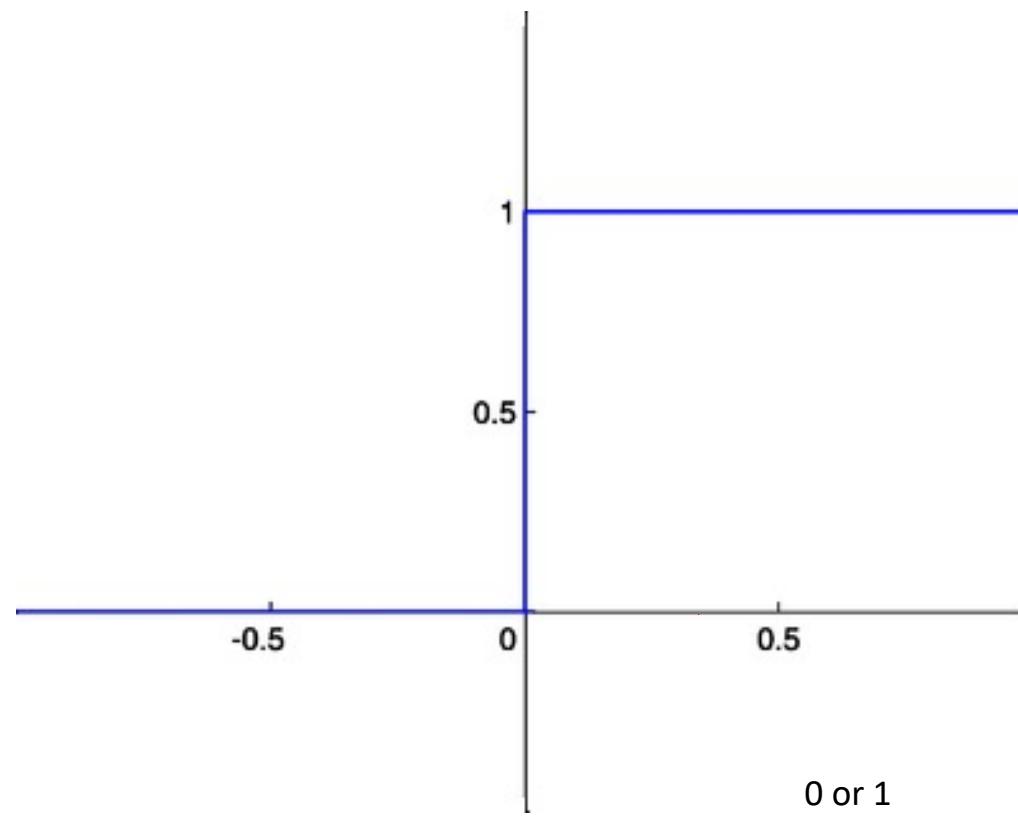
## Stochastic gradient descent



# GRADIENT DESCENT

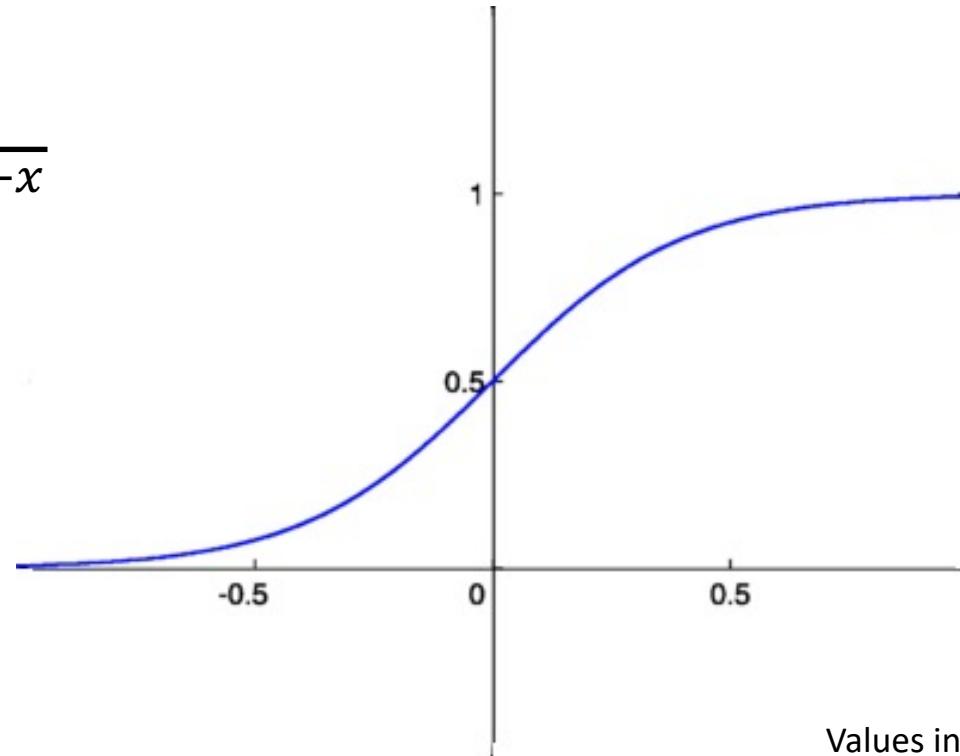
- Stochastic
  - Helps prevent local minima
  - Faster (no need to load all rows into memory)
- Mini batch gradient descent
  - Choose a number of rows to calculate the error and update the weights

# STEP FUNCTION



# SIGMOID FUNCTION

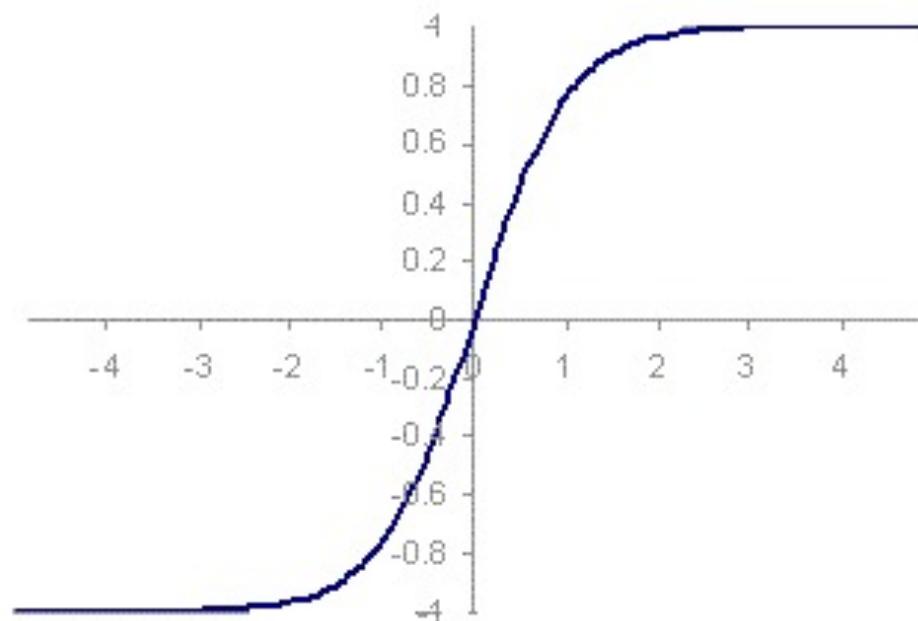
$$y = \frac{1}{1 + e^{-x}}$$



Values in the range from 0 to 1

# HYPERBOLIC TANGENT FUNCTION

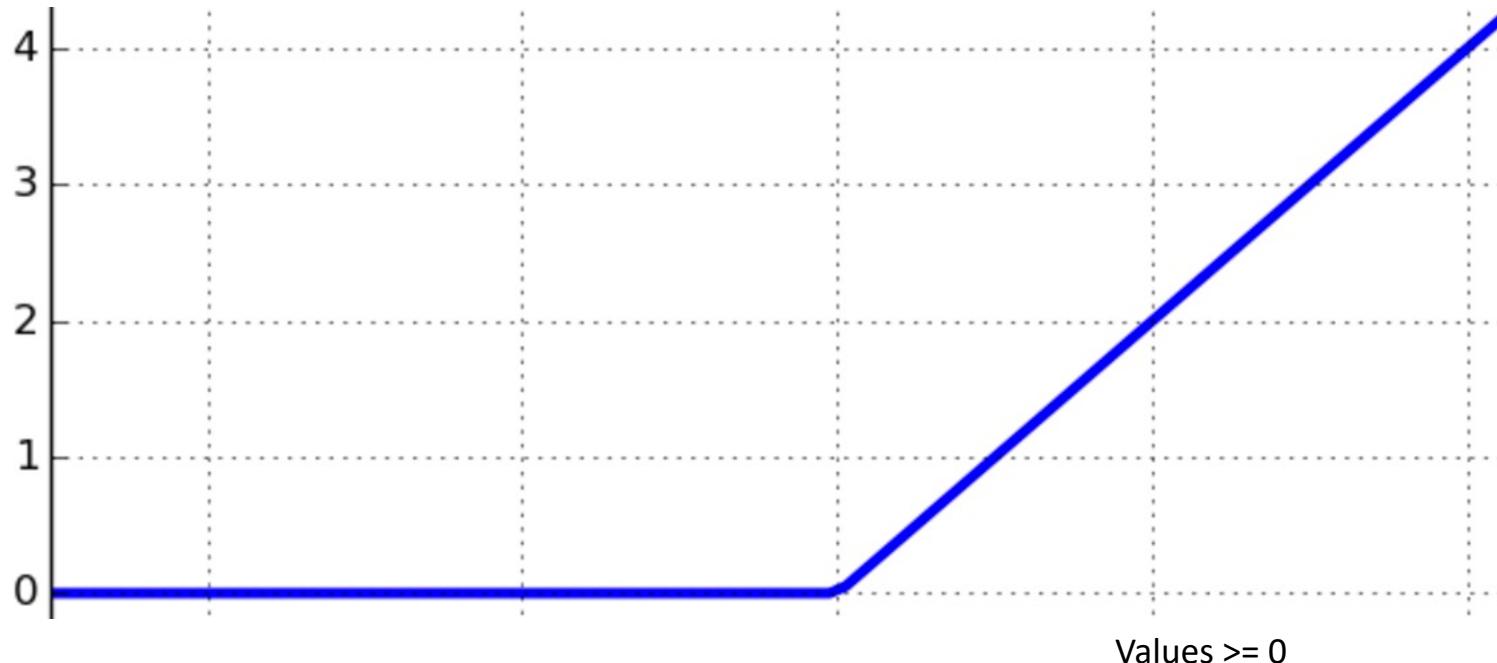
$$Y = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



Values in the range from -1 to 1

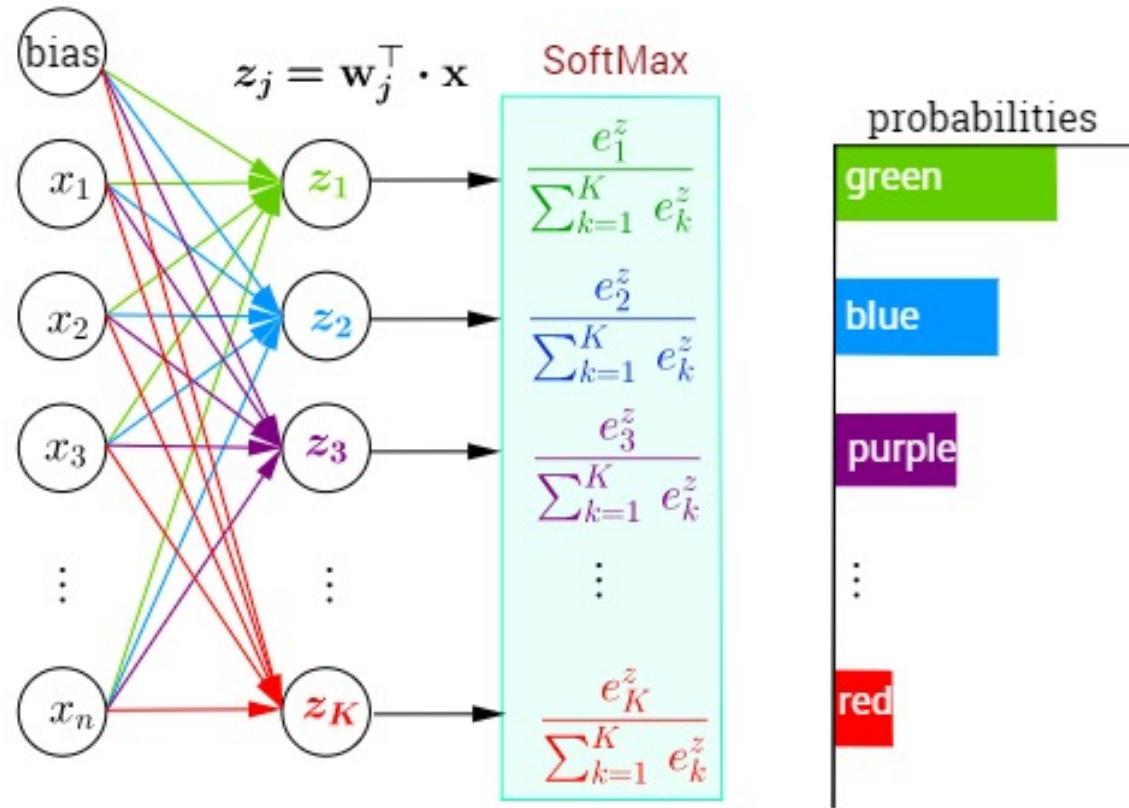
# RELU FUNCTION (RECTIFIED LINEAR UNITS)

$$Y = \max(0, x)$$



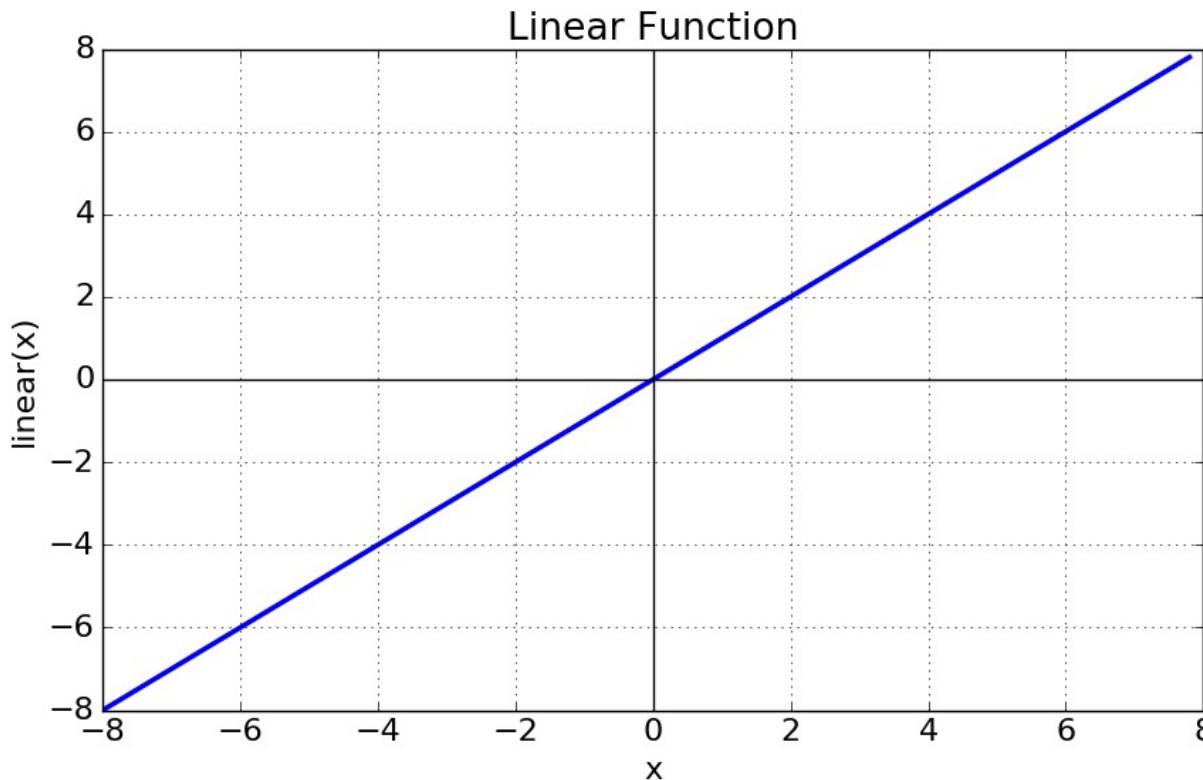
# SOFTMAX FUNCTION

$$Y = \frac{e(x)}{\sum e(x)}$$



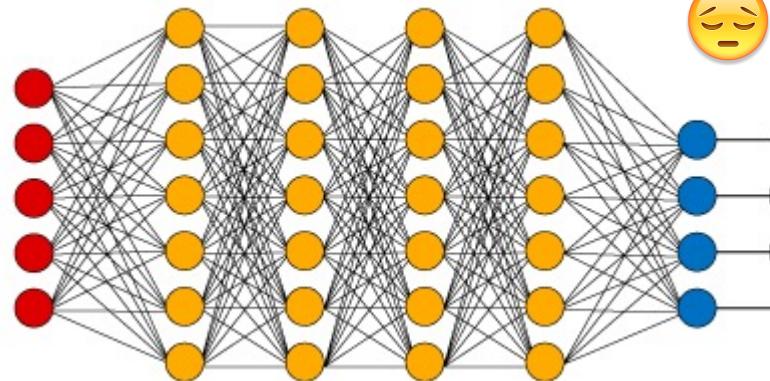
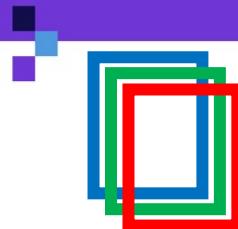
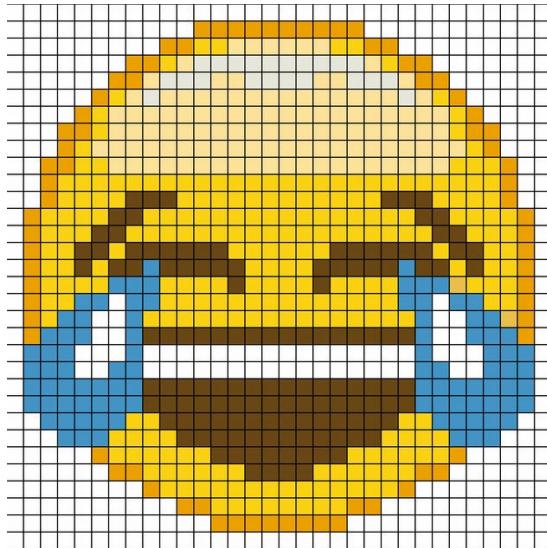
Source: <https://deeplearningsimplified.com/softmax-function/>

# LINEAR FUNCTION



# PIXELS

$$32 \times 32 = 1.024 \times 3 = \mathbf{3.072 \text{ inputs}}$$



- It does not use all pixels
- It applies a dense neural network, but at the beginning it transforms the data
- What are the most important features?

# CONVOLUTIONAL NEURAL NETWORK STEPS



1. Convolution operation
2. Pooling
3. Flattening
4. Dense neural network



# STEP 1: CONVOLUTION OPERATION

0	0	0	0	0	0	0	0
0	1	0	0	0	1	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	1	1	1
0	1	0	1	1	0	0	0
0	1	0	1	1	0	1	1
0	1	0	0	0	1	1	1

Image



x

1	0	0
1	0	1
0	1	1

Feature detector

=

0				

Feature map

$$0 * 1 + 0 * 0 + 0 * 0 + 0 * 1 + 1 * 0 + 0 * 1 + 0 * 0 + 0 * 1 + 0 * 1 = 0$$



# STEP 1: CONVOLUTION OPERATION



0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	1	1
0	1	0	1	1	0	0
0	1	0	1	1	0	1
0	1	0	0	0	1	1

Image

x



1	0	0
1	0	1
0	1	1

Feature detector

=

0	1			

Feature map

$$0 * 1 + 0 * 0 + 0 * 0 + 1 * 1 + 0 * 0 + 0 * 1 + 0 * 0 + 0 * 1 + 0 * 1 = 1$$



# STEP 1: CONVOLUTION OPERATION



0	0	0	0	0	0	0	0
0	1	0	0	0	1	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	1	1	1
0	1	0	1	1	0	0	0
0	1	0	1	1	0	1	1
0	1	0	0	0	1	1	1

Image

x



1	0	0
1	0	1
0	1	1

Feature detector

=

0	1	0		

Feature map

$$0 * 1 + 0 * 0 + 0 * 0 + 0 * 1 + 0 * 0 + 0 * 1 + 0 * 0 + 0 * 1 + 0 * 1 = 0$$



# STEP 1: CONVOLUTION OPERATION

0	0	0	0	0	0	0	0
0	1	0	0	0	1	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	1	1	1
0	1	0	1	1	0	0	0
0	1	0	1	1	0	1	1
0	1	0	0	0	1	1	1

Image



1	0	0
1	0	1
0	1	1

Feature detector

0	1	0	1	

Feature map

x

=

$$0 * 1 + 0 * 0 + 0 * 0 + 0 * 1 + 0 * 0 + 1 * 1 + 0 * 0 + 0 * 1 + 0 * 1 = 1$$

# STEP 1: CONVOLUTION OPERATION

0	0	0	0	0	0	0	0
0	1	0	0	0	1	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	1	1	0
0	1	0	1	1	0	0	0
0	1	0	1	1	0	1	1
0	1	0	0	0	0	1	1

Image



x

1	0	0
1	0	1
0	1	1

Feature detector

=

0	1	0	1	0
0	2	1	1	2
1	2	2	3	1
1	3	3	3	2
1	3	1	3	5

Feature map

$$1 * 1 + 0 * 0 + 0 * 0 + 1 * 1 + 0 * 0 + 1 * 1 + 0 * 0 + 1 * 1 + 1 * 1 = 5$$



# STEP 1: CONVOLUTION OPERATION – RELU



0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	1	1
0	1	0	1	1	0	0
0	1	0	1	1	0	1
0	1	0	0	0	1	1

Image


$$\begin{matrix} & \begin{matrix} 1 & 0 & 0 \\ 1 & 0 & 1 \\ 0 & 1 & 1 \end{matrix} \\ \times & \end{matrix} = \begin{matrix} & \begin{matrix} 0 & 1 & 0 & 1 & 0 \\ 0 & 2 & 1 & 1 & 2 \\ 1 & 2 & 2 & 3 & 1 \\ 1 & 3 & 3 & 3 & 2 \\ 1 & 3 & 1 & 3 & 5 \end{matrix} \\ \text{Feature detector} & \end{matrix}$$



0	1	0	1	0
0	2	1	1	2
1	2	2	3	1
1	3	3	3	2
1	3	1	3	5

Feature map



# STEP 1: CONVOLUTIONAL LAYER



0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	1	1
0	1	0	1	1	0	0
0	1	0	1	1	0	1
0	1	0	0	0	1	1

## Image



The diagram illustrates a sequence of binary strings and their corresponding numerical values. The strings are arranged in a grid, with each row representing a different string. The first four rows show the binary strings: 1010, 1001, 1011, and 1101. The last five rows show the numerical values: 0, 1, 0, 1, 0; 0, 2, 1, 1, 2; 1, 2, 2, 3, 1; 1, 3, 3, 3, 2; and 1, 3, 1, 3, 5.

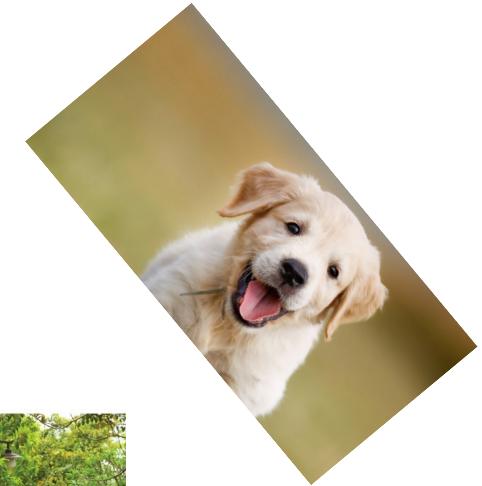
1	0	1	0	
1	0	0	1	
1	0	1	1	
1	1	0	1	
0	1	0	1	0
0	2	1	1	2
1	2	2	3	1
1	3	3	3	2
1	3	1	3	5

## Feature maps

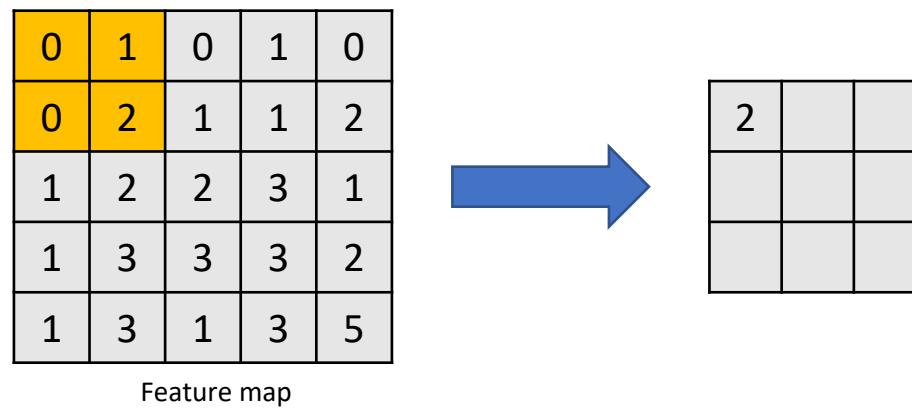
The network will decide which feature detector to use

The convolutional layer is the set of feature maps

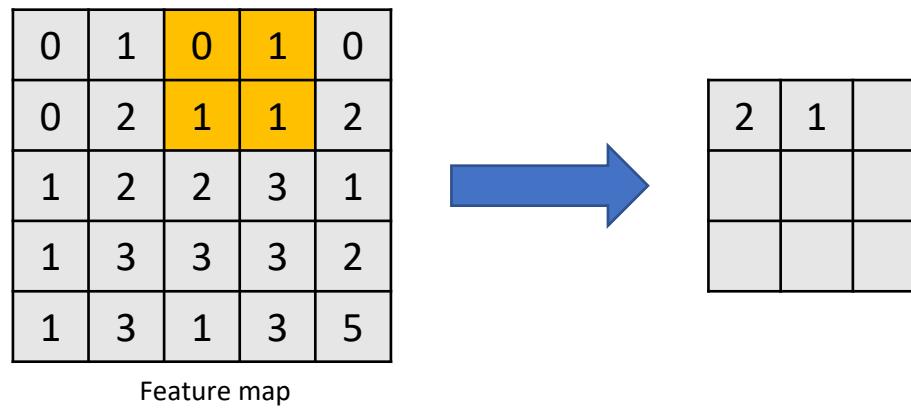
# STEP 2: POOLING



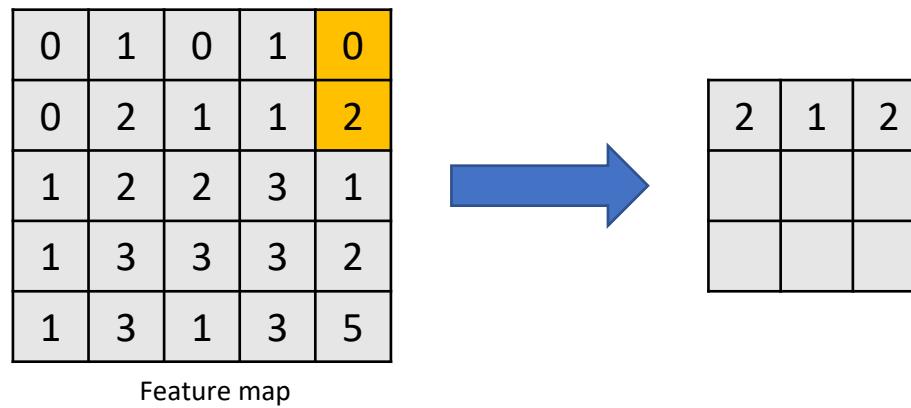
## STEP 2: POOLING



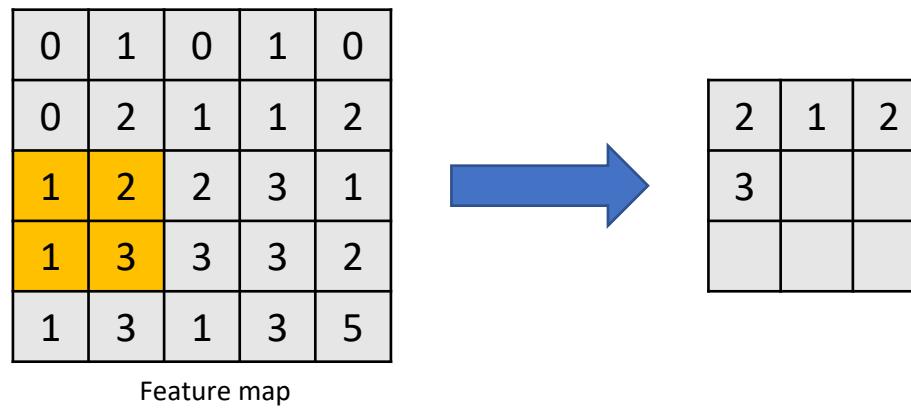
## STEP 2: POOLING



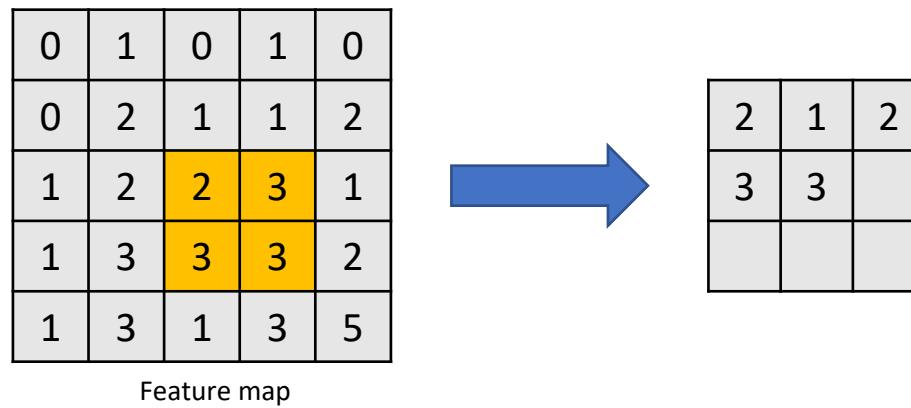
## STEP 2: POOLING



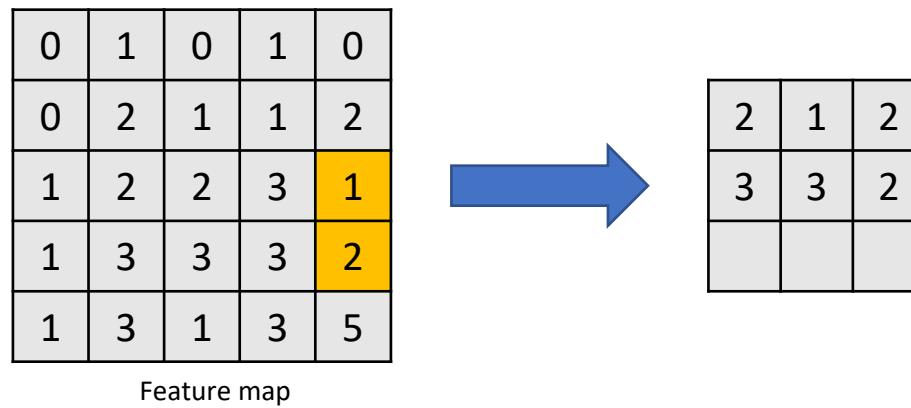
## STEP 2: POOLING



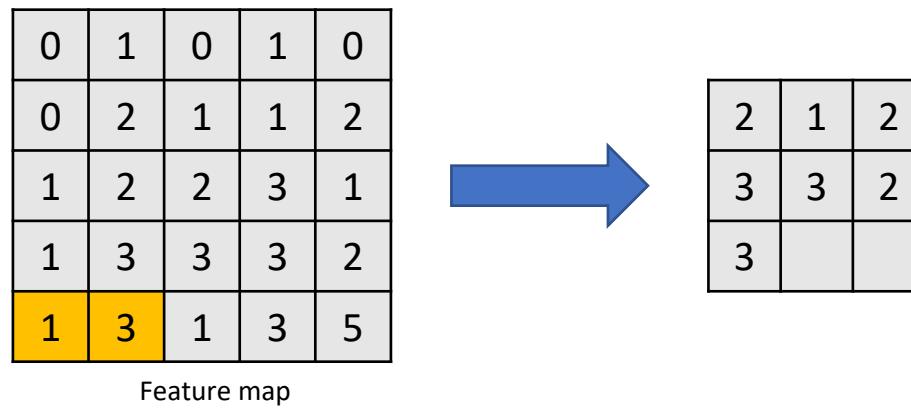
## STEP 2: POOLING



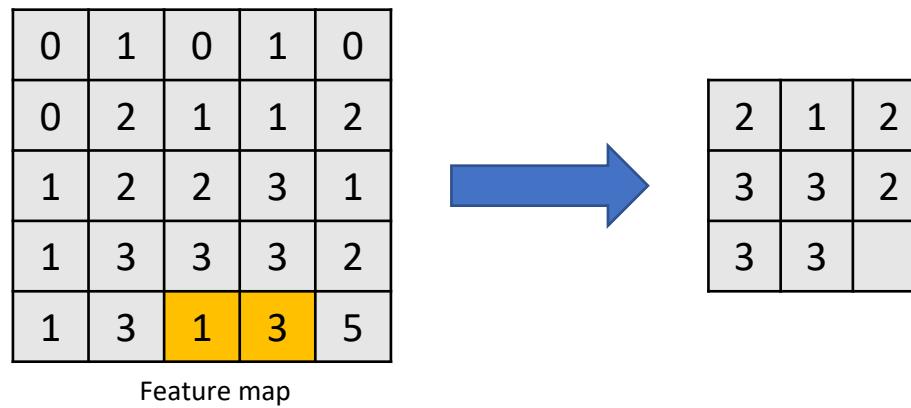
## STEP 2: POOLING



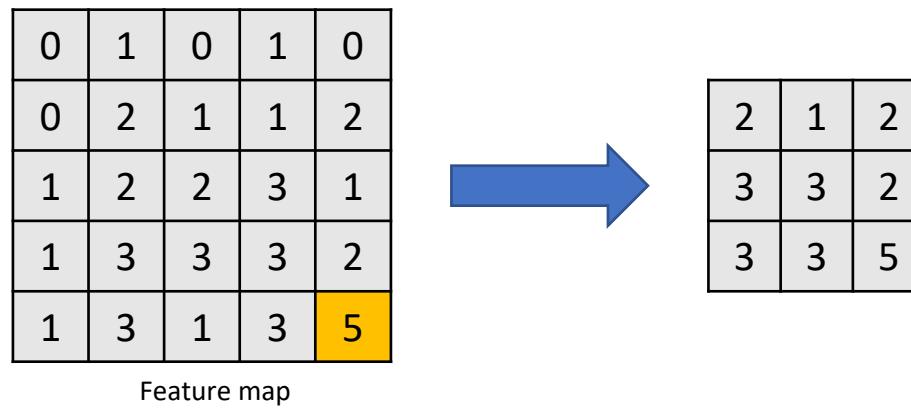
## STEP 2: POOLING



## STEP 2: POOLING



## STEP 2: POOLING

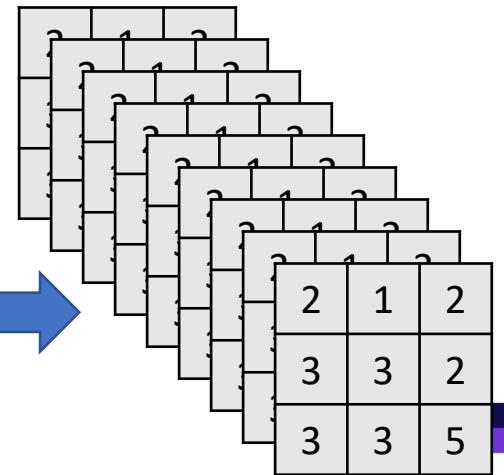
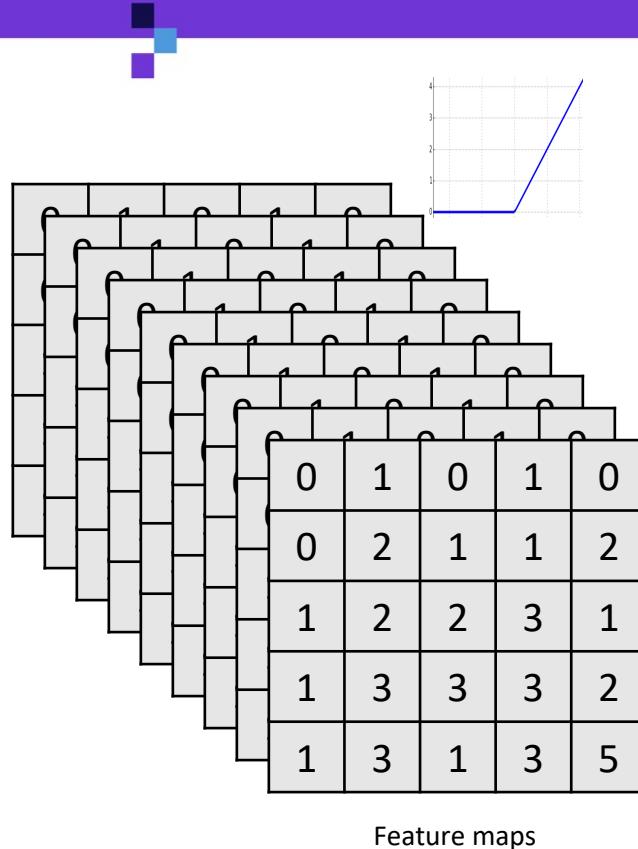
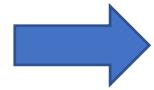


# CONVOLUTIONAL NEURAL NETWORK – POOLING

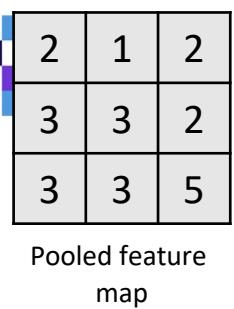


0	0	0	0	0	0	0	0
0	1	0	0	0	1	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	1	1	1
0	1	0	1	1	0	0	0
0	1	0	1	1	0	1	1
0	1	0	0	0	1	1	1

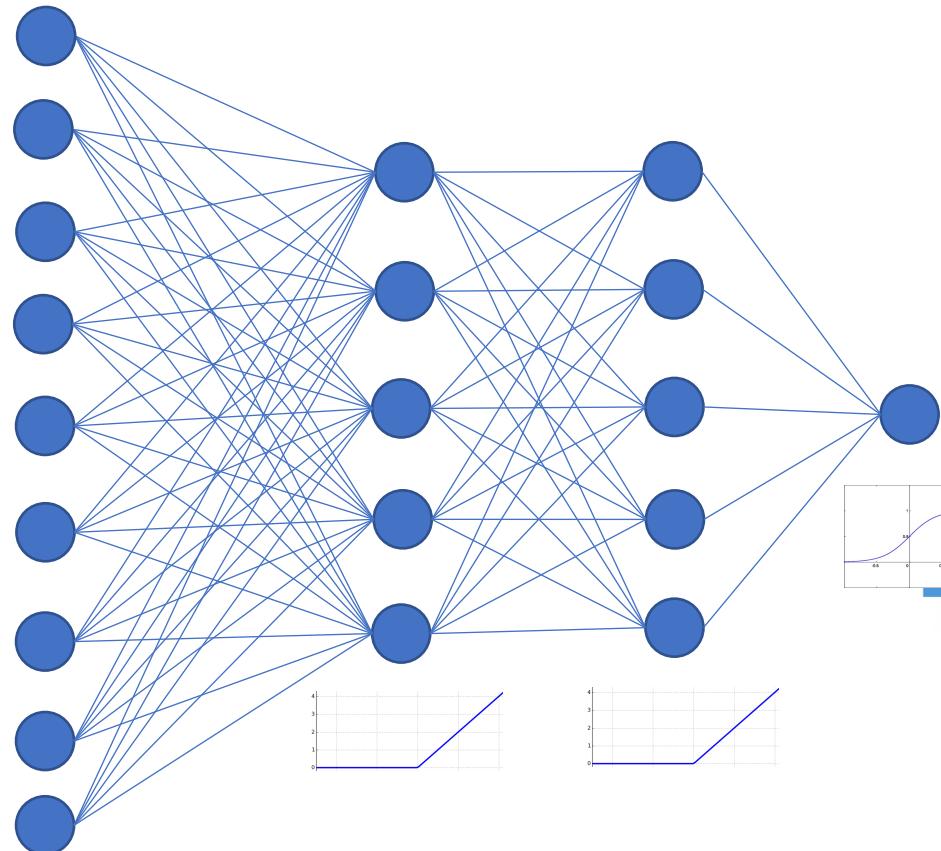
Image



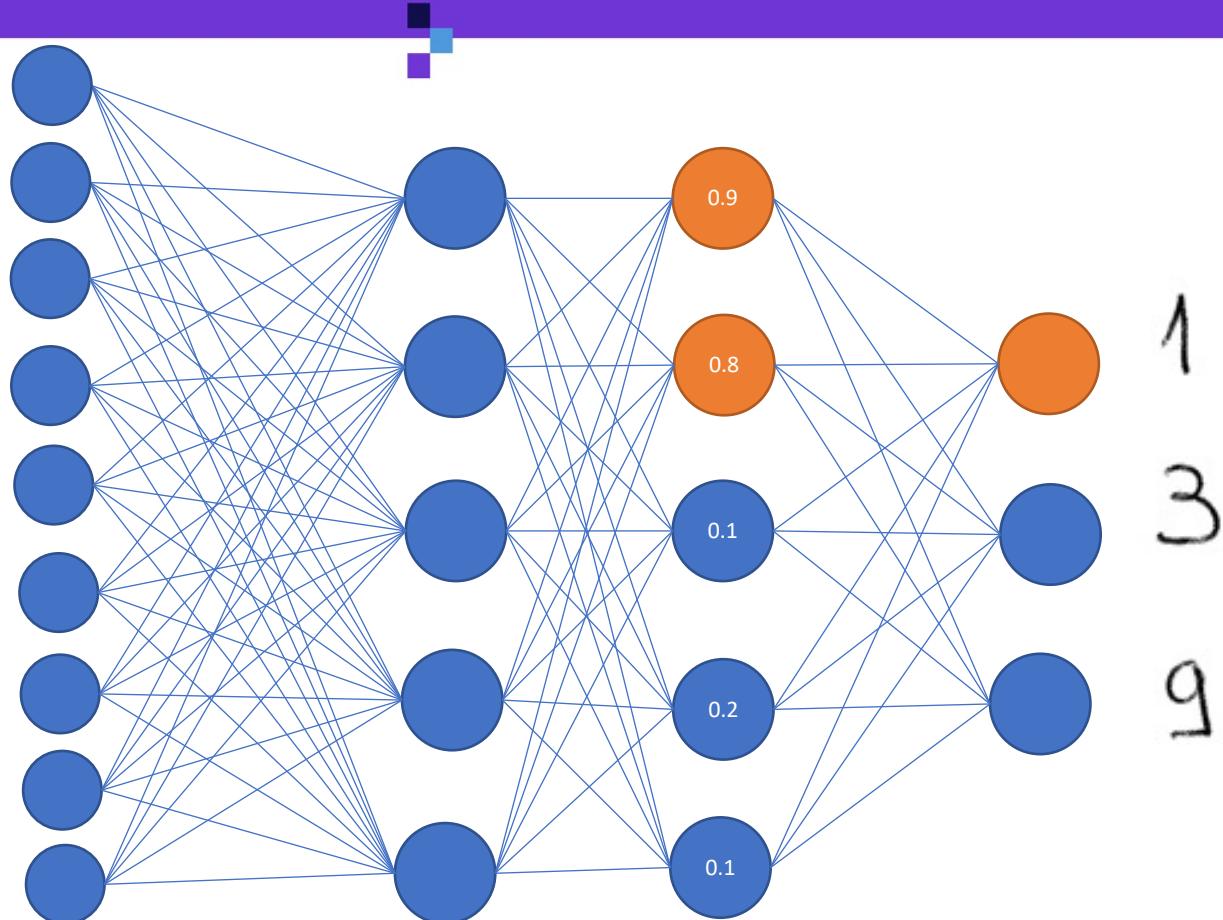
# STEP 3: FLATTENING



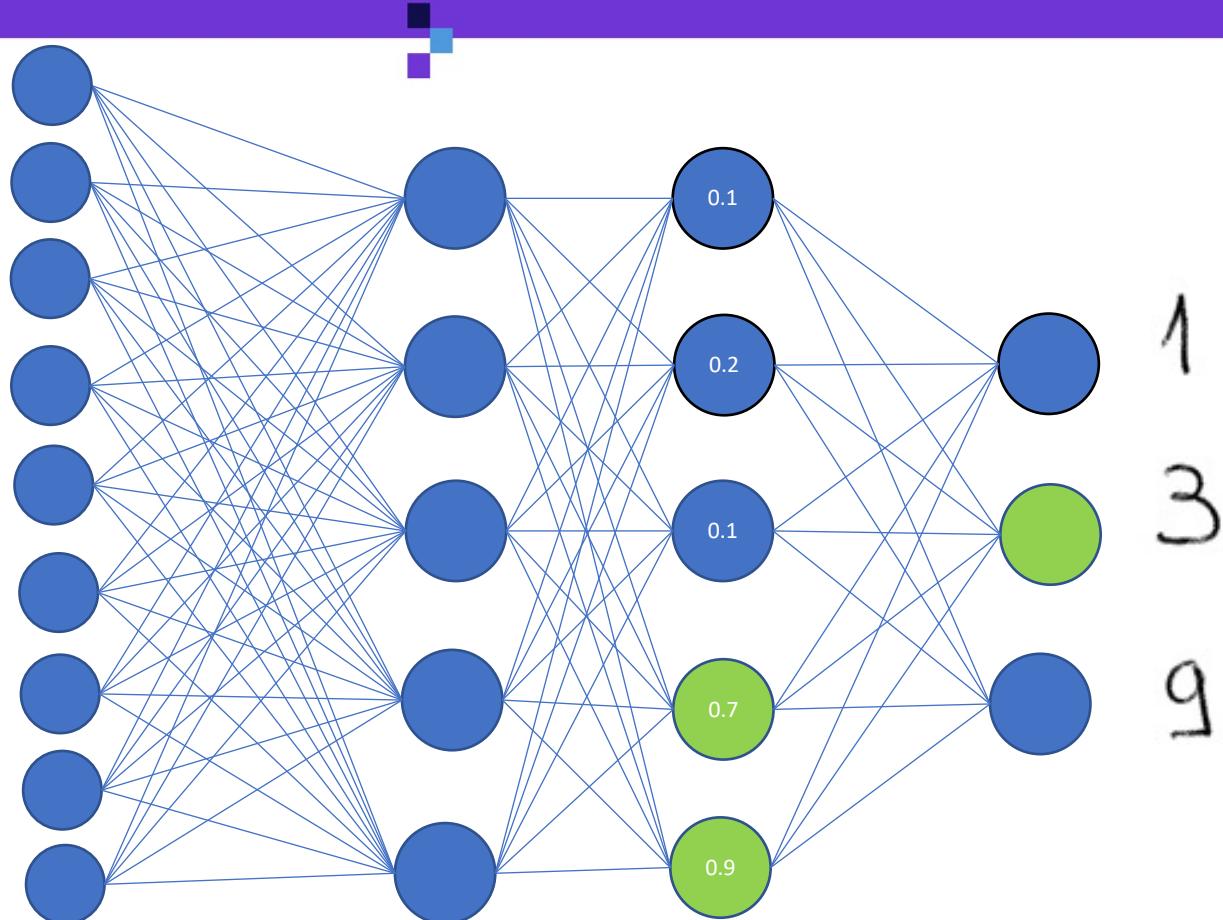
2
1
2
3
3
2
3
3
5



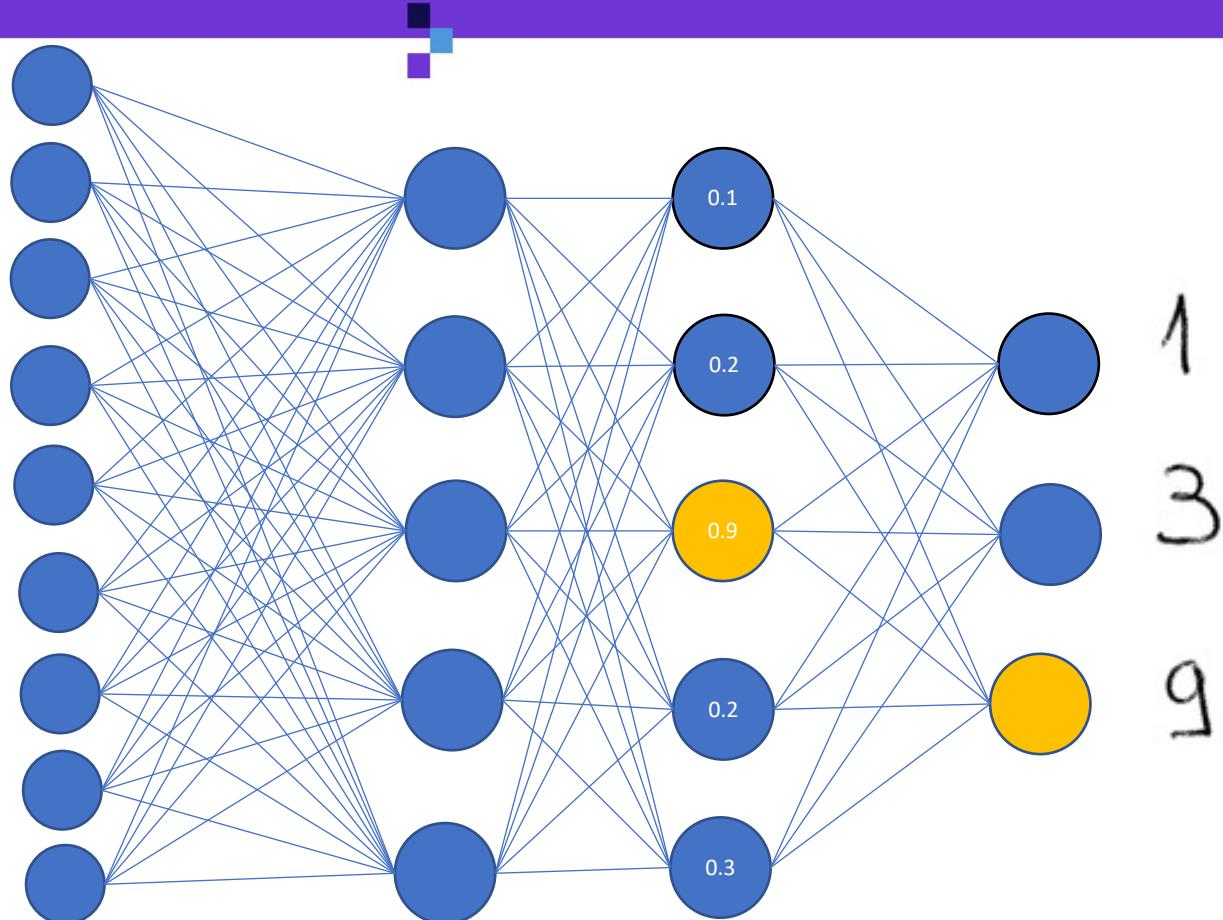
# STEP 3: DENSE NEURAL NETWORK



# STEP 3: DENSE NEURAL NETWORK



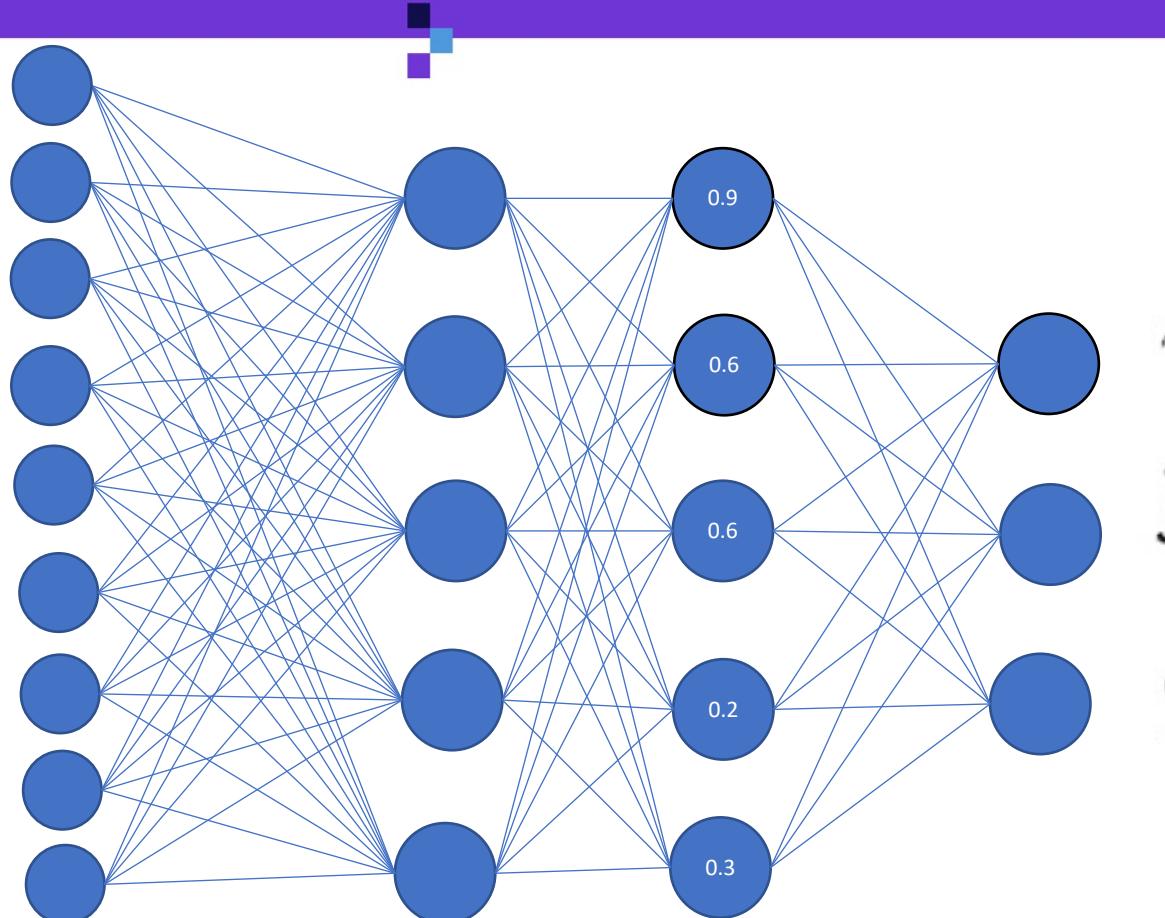
# STEP 3: DENSE NEURAL NETWORK



1  
3  
9



# STEP 3: DENSE NEURAL NETWORK



65%

05%

30%

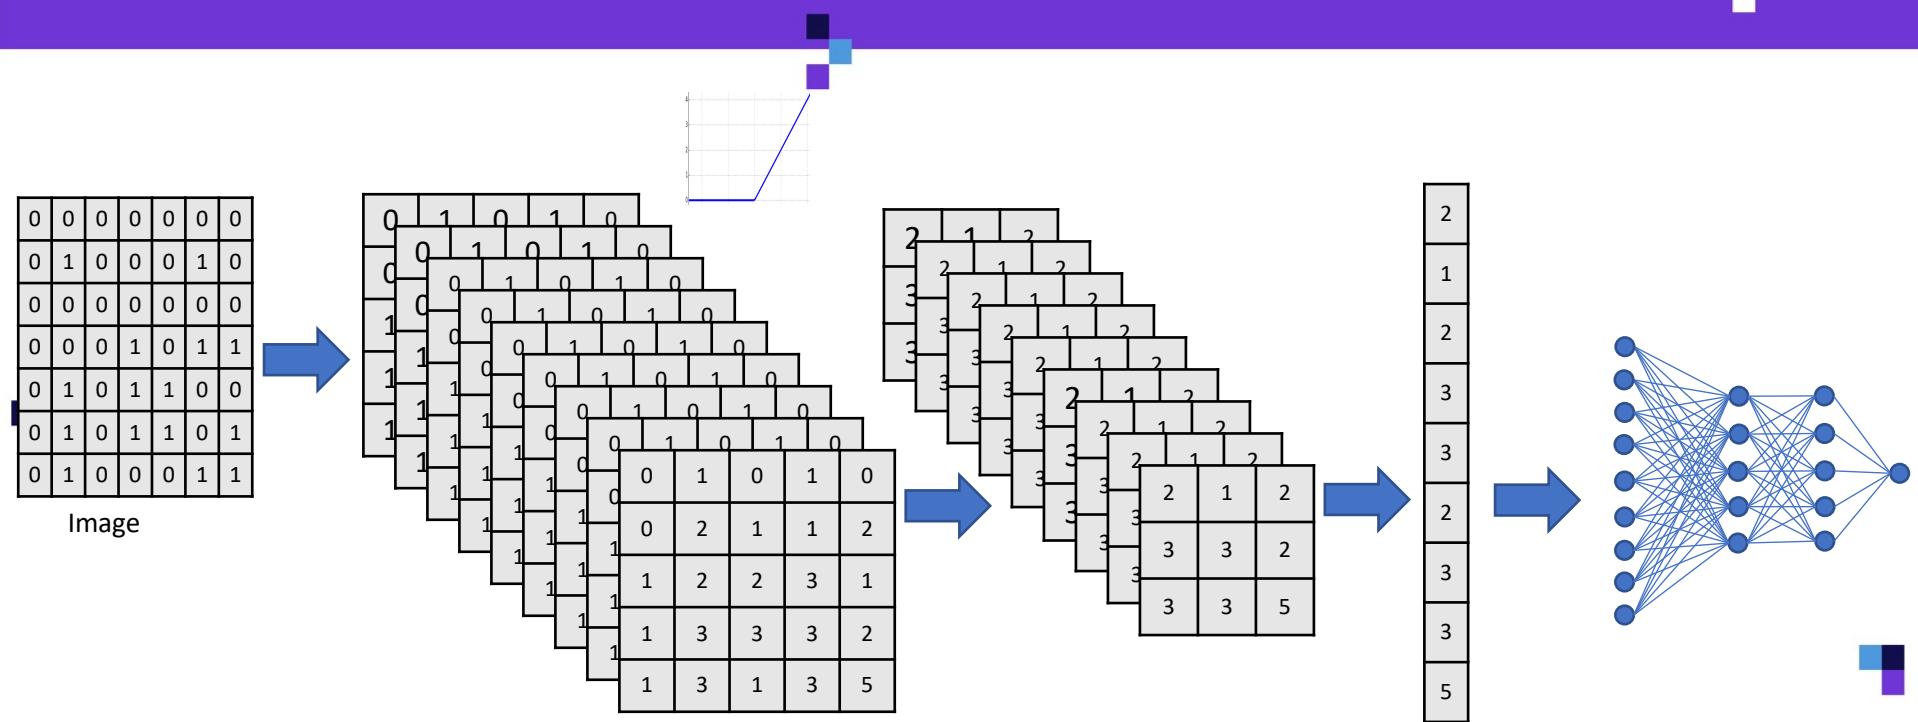
1

3

9



# CONVOLUTIONAL NEURAL NETWORK



## Training using gradient descent

In addition to adjusting the weights, the feature detector is also changed